

DSFBA: Data Wrangling

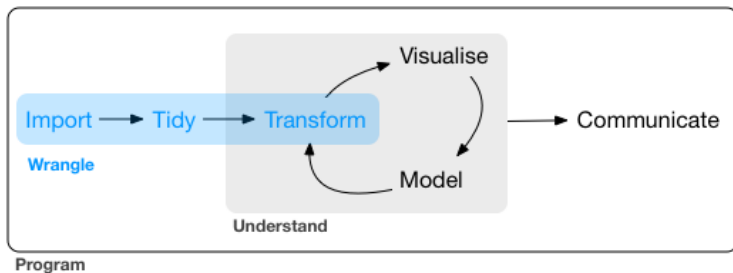
Data Science for Business Analytics

Thibault Vatter

Department of Statistics, Columbia University

10/13/2021

- 1 Tidy data
- 2 Filter
- 3 Arrange
- 4 Select
- 5 Mutate
- 6 Summarize
- 7 Relational data



Most of the material (e.g., the picture above) is borrowed from

R for data science

- When working with data you must:
 - ▶ Figure out what you want to do.
 - ▶ Describe those tasks as a computer program.
 - ▶ Execute the program.
- The `dplyr` package makes this fast and easy with 5 verbs!
 - ▶ `filter()`: select observations based on their values.
 - ▶ `arrange()`: reorder the observations.
 - ▶ `select()`: select variables based on their names.
 - ▶ `mutate()`: add variables as functions of existing variables.
 - ▶ `summarize()`: collapse many values down to a single summary.
- Two important features:
 - ▶ Verbs can be used with `group_by()` to operate groupwise.
 - ▶ Verbs work similarly. . .
 1. First argument: a data frame.
 2. Other arguments: what to do with it using variable names.
 3. Result: a new data frame.

All 336,776 flights that departed from NYC in 2013 (US BTS):

```
nycflights13::flights
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>      <dbl>   <int>
#> 1  2013     1     1     517           515         2     830
#> 2  2013     1     1     533           529         4     850
#> 3  2013     1     1     542           540         2     923
#> 4  2013     1     1     544           545        -1    1004
#> 5  2013     1     1     554           600        -6     812
#> 6  2013     1     1     554           558        -4     740
#> 7  2013     1     1     555           600        -5     913
#> 8  2013     1     1     557           600        -3     709
#> 9  2013     1     1     557           600        -3     838
#> 10 2013     1     1     558           600        -2     753
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```

What is this code doing?

```
a1 <- group_by(flights, year, month, day)
a2 <- select(a1, arr_delay, dep_delay)
a3 <- summarize(a2,
                arr = mean(arr_delay, na.rm = TRUE),
                dep = mean(dep_delay, na.rm = TRUE))
filter(a3, arr > 30 | dep > 30)
#> # A tibble: 49 x 5
#> # Groups:   year, month [11]
#>   year month   day   arr   dep
#>   <int> <int> <int> <dbl> <dbl>
#> 1  2013     1    16  34.2  24.6
#> 2  2013     1    31  32.6  28.7
#> 3  2013     2    11  36.3  39.1
#> 4  2013     2    27  31.3  37.8
#> 5  2013     3     8  85.9  83.5
#> 6  2013     3    18  41.3  30.1
#> 7  2013     4    10  38.4  33.0
#> 8  2013     4    12  36.0  34.8
#> 9  2013     4    18  36.0  34.9
#> 10 2013     4    19  47.9  46.1
#> # ... with 39 more rows
```

Same code (no unnecessary objects)

```
filter(summarize(select(group_by(flights, year, month, day),
  arr_delay, dep_delay),
  arr = mean(arr_delay, na.rm = TRUE),
  dep = mean(dep_delay, na.rm = TRUE)),
  arr > 30 | dep > 30)
#> # A tibble: 49 x 5
#> # Groups:   year, month [11]
#>   year month   day   arr   dep
#>   <int> <int> <int> <dbl> <dbl>
#> 1  2013     1    16  34.2  24.6
#> 2  2013     1    31  32.6  28.7
#> 3  2013     2    11  36.3  39.1
#> 4  2013     2    27  31.3  37.8
#> 5  2013     3     8  85.9  83.5
#> 6  2013     3    18  41.3  30.1
#> 7  2013     4    10  38.4  33.0
#> 8  2013     4    12  36.0  34.8
#> 9  2013     4    18  36.0  34.9
#> 10 2013     4    19  47.9  46.1
#> # ... with 39 more rows
```

... Or use %>%

```
flights %>%
  group_by(year, month, day) %>%
  select(arr_delay, dep_delay) %>%
  summarize(arr = mean(arr_delay, na.rm = TRUE),
            dep = mean(dep_delay, na.rm = TRUE)) %>%
  filter(arr > 30 | dep > 30)
#> # A tibble: 49 x 5
#> # Groups:   year, month [11]
#>   year month   day   arr   dep
#>   <int> <int> <int> <dbl> <dbl>
#> 1  2013     1    16  34.2  24.6
#> 2  2013     1    31  32.6  28.7
#> 3  2013     2    11  36.3  39.1
#> 4  2013     2    27  31.3  37.8
#> 5  2013     3     8  85.9  83.5
#> 6  2013     3    18  41.3  30.1
#> 7  2013     4    10  38.4  33.0
#> 8  2013     4    12  36.0  34.8
#> 9  2013     4    18  36.0  34.9
#> 10 2013     4    19  47.9  46.1
#> # ... with 39 more rows
```


- `x %>% f` is equivalent to `f(x)`
- `x %>% f(y)` is equivalent to `f(x, y)`
- `x %>% f(y) %>% g(z)` is equivalent to `g(f(x, y), z)`

```
x <- 1:10
y <- x + 1
z <- y + 1
f <- function(x, y) x + y

x %>% sum
#> [1] 55
x %>% f(y)
#> [1] 3 5 7 9 11 13 15 17 19 21
x %>% f(y) %>% f(z)
#> [1] 6 9 12 15 18 21 24 27 30 33
```

The argument (“dot”) placeholder

- `x %>% f(y, .)` is equivalent to `f(y, x)`
- `x %>% f(y, z = .)` is equivalent to `f(y, z = x)`

```
x <- 1:10
y <- 2 * x
f <- function(z, y) y / z

x %>% f(y, .)
#> [1] 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5
x %>% f(y, z = .)
#> [1] 2 2 2 2 2 2 2 2 2 2
```

- Each of the three options has its own strengths and weaknesses:
 - ▶ Nesting, $f(g(x))$:
 - Concise, and well suited for short sequences.
 - Longer sequences harder to read (inside out & right to left).
 - Arguments can get spread out over long distances (see [Dagwood sandwich](#)).
 - ▶ Intermediate objects, $y \leftarrow f(x); g(y)$:
 - Requires you to name intermediate objects.
 - A strength when objects are important, but a weakness when values are truly intermediate.
 - ▶ Piping, $x \%>\% f() \%>\% g()$:
 - Allows to read code in straightforward left-to-right fashion.
 - Doesn't require to name intermediate objects.
 - Only for linear sequences of transformations of a single object.
- Most code use a combination of all three styles, but...
- **Piping is more common in data analysis code!**

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

“Happy families are all alike; every unhappy family is unhappy in its own way.” — Leo Tolstoy

“Tidy datasets are all alike, but every messy dataset is messy in its own way.” — Hadley Wickham

To learn more about the underlying theory, see the [Tidy Data paper](#).

Which representation is “best”?

■ First representation?

table1

```
#> # A tibble: 6 x 4
#>   country    year cases population
#>   <chr>    <int> <int>      <int>
#> 1 Afghanistan 1999    745  19987071
#> 2 Afghanistan 2000   2666 20595360
#> 3 Brazil      1999  37737 172006362
#> 4 Brazil      2000   80488 174504898
#> 5 China       1999 212258 1272915272
#> 6 China       2000 213766 1280428583
```

■ Second representation?

table2

```
#> # A tibble: 12 x 4
#>   country    year type      count
#>   <chr>    <int> <chr>      <int>
#> 1 Afghanistan 1999 cases        745
#> 2 Afghanistan 1999 population 19987071
#> 3 Afghanistan 2000 cases        2666
#> 4 Afghanistan 2000 population 20595360
#> 5 Brazil      1999 cases        37737
#> 6 Brazil      1999 population 172006362
#> 7 Brazil      2000 cases        80488
#> 8 Brazil      2000 population 174504898
#> 9 China       1999 cases        212258
#> 10 China      1999 population 1272915272
#> 11 China      2000 cases        213766
#> 12 China      2000 population 1280428583
```

■ Third representation?

table3

```
#> # A tibble: 6 x 3
#>   country    year rate
#>   * <chr>    <int> <chr>
#> 1 Afghanistan 1999 745/19987071
#> 2 Afghanistan 2000 2666/20595360
#> 3 Brazil      1999 37737/172006362
#> 4 Brazil      2000 80488/174504898
#> 5 China       1999 212258/1272915272
#> 6 China       2000 213766/1280428583
```

■ Fourth representation?

table4a # cases

```
#> # A tibble: 3 x 3
#>   country    `1999` `2000`
#>   * <chr>    <int> <int>
#> 1 Afghanistan    745    2666
#> 2 Brazil        37737  80488
#> 3 China         212258 213766
```

table4b # population

```
#> # A tibble: 3 x 3
#>   country    `1999` `2000`
#>   * <chr>    <int> <int>
#> 1 Afghanistan 19987071 20595360
#> 2 Brazil      172006362 174504898
#> 3 China       1272915272 1280428583
```

What makes a dataset tidy?

Three interrelated rules:

- Each variable must have its own column.
- Each observation must have its own row.
- Each value must have its own cell.

country	year	cases	population
Afghanistan	1999	1845	12000071
Afghanistan	2000	1666	20095360
Brazil	1999	37737	172006362
Brazil	2000	80488	174004898
China	1999	210258	1270015272
China	2000	210766	1280008583

variables

country	year	cases	population
Afghanistan	1999	1845	12000071
Afghanistan	2000	1666	20095360
Brazil	1999	37737	172006362
Brazil	2000	80488	174004898
China	1999	210258	1270015272
China	2000	210766	1280008583

observations

country	year	cases	population
Afghanistan	1999	1845	12000071
Afghanistan	2000	1666	20095360
Brazil	1999	37737	172006362
Brazil	2000	80488	174004898
China	1999	210258	1270015272
China	2000	210766	1280008583

values

Because it's impossible to only satisfy two of the three:

- Put each dataset in a tibble.
- Put each variable in a column.

- Why?
 - ▶ With consistent data structure, it's easier to learn the tools that work with it because they have an underlying uniformity.
 - ▶ Placing variables in columns allows R's vectorized nature to shine.
- Tidy data principles seem obvious, BUT:
 - ▶ Most people aren't familiar with them.
 - ▶ Data often organized to facilitate something different than analysis.
- Hence, you'll most likely need to do some tidying.

- Figure out what the variables and observations are.
- Resolve one of two common problems:
 - ▶ One variable might be spread across multiple columns.
 - ▶ One observation might be scattered across multiple rows.

... To fix these problems, you'll need `pivot_longer()` and `pivot_wider()`.

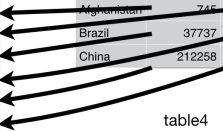
Longer with pivot_wider()

table4a

```
#> # A tibble: 3 x 3  
#>   country   `1999` `2000`  
#> * <chr>     <int> <int>  
#> 1 Afghanistan    745   2666  
#> 2 Brazil        37737  80488  
#> 3 China         212258 213766
```

table4a %>%

```
  pivot_longer(c(`1999`, `2000`),  
               names_to = "year",  
               values_to = "cases")  
#> # A tibble: 6 x 3  
#>   country    year  cases  
#>   <chr>     <chr> <int>  
#> 1 Afghanistan 1999     745  
#> 2 Afghanistan 2000    2666  
#> 3 Brazil      1999   37737  
#> 4 Brazil      2000   80488  
#> 5 China       1999  212258  
#> 6 China       2000  213766
```



country	year	cases
Afghanistan	1999	745
Afghanistan	2000	2666
Brazil	1999	37737
Brazil	2000	80488
China	1999	212258
China	2000	213766

table4

Wider with pivot_wider()

table2

```
#> # A tibble: 12 x 4
#>   country year type count
#>   <chr>    <int> <chr>   <int>
#> 1 Afghanistan 1999 cases    745
#> 2 Afghanistan 1999 population 19987071
#> 3 Afghanistan 2000 cases    2666
#> 4 Afghanistan 2000 population 20595360
#> 5 Brazil 1999 cases    37737
#> 6 Brazil 1999 population 172006362
#> 7 Brazil 2000 cases    80488
#> 8 Brazil 2000 population 174504898
#> 9 China 1999 cases    212258
#> 10 China 1999 population 1272915272
#> 11 China 2000 cases    213766
#> 12 China 2000 population 1280428583
```

table2 %>%

```
pivot_wider(names_from = type,
             values_from = count)
```

```
#> # A tibble: 6 x 4
#>   year country cases population
#>   <chr>    <int> <int>    <int>
#> 1 Afghanistan 1999    745 19987071
#> 2 Afghanistan 2000   2666 20595360
#> 3 Brazil 1999  37737 172006362
#> 4 Brazil 2000  80488 174504898
#> 5 China 1999 212258 1272915272
#> 6 China 2000 213766 1280428583
```

country	year	key	value
Afghanistan	1999	cases	745
Afghanistan	1999	population	19987071
Afghanistan	2000	cases	2666
Afghanistan	2000	population	20595360
Brazil	1999	cases	37737
Brazil	1999	population	172006362
Brazil	2000	cases	80488
Brazil	2000	population	174504898
China	1999	cases	212258
China	1999	population	1272915272
China	2000	cases	213766
China	2000	population	1280428583

table2

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

Separate a column with separate()

table3

```
#> # A tibble: 6 x 3  
#>   country    year rate  
#>   <chr>    <int> <chr>  
#> 1 Afghanistan 1999 745/19987071  
#> 2 Afghanistan 2000 2666/20595360  
#> 3 Brazil      1999 37737/172006362  
#> 4 Brazil      2000 80488/174504898  
#> 5 China       1999 212258/1272915272  
#> 6 China       2000 213766/1280428583
```

```
table3 %>% separate(rate,  
                     into = c("cases",  
                              "population"))  
#> # A tibble: 6 x 4  
#>   country    year cases population  
#>   <chr>    <int> <chr>    <chr>  
#> 1 Afghanistan 1999 745      19987071  
#> 2 Afghanistan 2000 2666     20595360  
#> 3 Brazil      1999 37737    172006362  
#> 4 Brazil      2000 80488    174504898  
#> 5 China       1999 212258   1272915272  
#> 6 China       2000 213766   1280428583
```



country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

table3

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

separate() using convert = TRUE

```
table3 %>%  
  separate(rate, into = c("cases", "population"), convert = TRUE)  
#> # A tibble: 6 x 4  
#>   country    year cases population  
#>   <chr>    <int> <int>      <int>  
#> 1 Afghanistan 1999    745   19987071  
#> 2 Afghanistan 2000   2666   20595360  
#> 3 Brazil      1999  37737   172006362  
#> 4 Brazil      2000  80488   174504898  
#> 5 China       1999 212258  1272915272  
#> 6 China       2000 213766  1280428583
```

Unite two columns with unite()

table5

```
#> # A tibble: 6 x 4
#>   country    century year  rate
#>   <chr>      <chr>  <chr> <chr>
#> 1 Afghanistan 19      99    745/19987071
#> 2 Afghanistan 20      00    2666/20595360
#> 3 Brazil      19      99    37737/172006362
#> 4 Brazil      20      00    80488/174504898
#> 5 China       19      99    212258/1272915272
#> 6 China       20      00    213766/1280428583
```

table5 %>%

```
  unite(new, century, year, sep = "")
#> # A tibble: 6 x 3
#>   country    new    rate
#>   <chr>      <chr> <chr>
#> 1 Afghanistan 1999    745/19987071
#> 2 Afghanistan 2000    2666/20595360
#> 3 Brazil      1999    37737/172006362
#> 4 Brazil      2000    80488/174504898
#> 5 China       1999    212258/1272915272
#> 6 China       2000    213766/1280428583
```



country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

country	century	year	rate
Afghanistan	19	99	745 / 19987071
Afghanistan	20	0	2666 / 20595360
Brazil	19	99	37737 / 172006362
Brazil	20	0	80488 / 174504898
China	19	99	212258 / 1272915272
China	20	0	213766 / 1280428583

table6

- A value can be missing in one of two possible ways:
 - ▶ **Explicitly**, i.e. flagged with NA.
 - ▶ **Implicitly**, i.e. simply not present in the data.

“An explicit missing value is the presence of an absence; an implicit missing value is the absence of a presence.” Hadley Wickham

- Are there missing values in this dataset?

```
stocks <- tibble(  
  year   = c(2015, 2015, 2015, 2015, 2016, 2016, 2016),  
  qtr    = c( 1,    2,    3,    4,    2,    3,    4),  
  return = c(1.88, 0.59, 0.35, NA, 0.92, 0.17, 2.66)  
)
```

- Implicit to explicit by pivoting:

```
stocks %>%  
  pivot_wider(names_from = year,  
              values_from = return)  
#> # A tibble: 4 x 3  
#>   qtr `2015` `2016`  
#>   <dbl> <dbl> <dbl>  
#> 1     1     1  1.88  NA  
#> 2     2     2  0.59  0.92  
#> 3     3     3  0.35  0.17  
#> 4     4     4  NA    2.66
```

- Implicit to explicit using complete:

```
stocks %>% complete(year, qtr)  
#> # A tibble: 8 x 3  
#>   year  qtr return  
#>   <dbl> <dbl> <dbl>  
#> 1  2015     1  1.88  
#> 2  2015     2  0.59  
#> 3  2015     3  0.35  
#> 4  2015     4  NA  
#> 5  2016     1  NA  
#> 6  2016     2  0.92  
#> 7  2016     3  0.17  
#> 8  2016     4  2.66
```

- Explicit to implicit via drop_na().

Fill in missing values with fill()

```
treatment <- tribble(
  ~ person,      ~ treatment, ~response,
  "Derrick Whitmore", 1,          7,
  NA,              2,          10,
  NA,              3,          9,
  "Katherine Burke", 1,          4
)
treatment %>%
  fill(person)
#> # A tibble: 4 x 3
#>   person      treatment response
#>   <chr>         <dbl>     <dbl>
#> 1 Derrick Whitmore      1         7
#> 2 Derrick Whitmore      2        10
#> 3 Derrick Whitmore      3         9
#> 4 Katherine Burke       1         4
```

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

Filter rows with filter()

```
filter(flights, month == 1, day == 1)
#> # A tibble: 842 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>       <dbl>   <int>
#> 1  2013     1     1     517           515         2     830
#> 2  2013     1     1     533           529         4     850
#> 3  2013     1     1     542           540         2     923
#> 4  2013     1     1     544           545        -1    1004
#> 5  2013     1     1     554           600        -6     812
#> 6  2013     1     1     554           558        -4     740
#> 7  2013     1     1     555           600        -5     913
#> 8  2013     1     1     557           600        -3     709
#> 9  2013     1     1     557           600        -3     838
#> 10 2013     1     1     558           600        -2     753
#> # ... with 832 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```

- The standard suite: $>$, $>=$, $<$, $<=$, $!=$, and $==$.
- Most common mistake:







```
filter(flights, month = 1)
```

- What happens in the following?

```
sqrt(2) ^ 2 == 2
#> [1] FALSE
1/49 * 49 == 1
#> [1] FALSE
near(sqrt(2) ^ 2, 2)
#> [1] TRUE
near(1 / 49 * 49, 1)
#> [1] TRUE
```

Multiple arguments to `filter()` are combined with:

- `&` for “and”
- `|` for “or”
- `!` for “not”



What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

- Literally “finds all flights that departed in November or December”.
- ... `filter(flights, month == 11 | 12)` ?

What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

- Literally “finds all flights that departed in November or December”.
- ... `filter(flights, month == 11 | 12)` ?
- No, but a solution:

```
filter(flights, month %in% c(11, 12))
```


- $!(x \ \& \ y)$ is the same as $!x \ | \ !y$
- $!(x \ | \ y)$ is the same as $!x \ \& \ !y$

```
all.equal(  
  filter(flights, !(arr_delay > 120 | dep_delay > 120)),  
  filter(flights, arr_delay <= 120, dep_delay <= 120)  
)  
#> [1] TRUE
```

Missing values and filter()

```
df <- tibble(x = c(1, NA, 3))
filter(df, x > 1)
#> # A tibble: 1 x 1
#>       x
#>   <dbl>
#> 1     3
filter(df, is.na(x) | x > 1)
#> # A tibble: 2 x 1
#>       x
#>   <dbl>
#> 1   NA
#> 2     3
```

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

Arrange rows with arrange()

```
arrange(flights, year, month, day)
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>       <dbl>   <int>
#> 1  2013     1     1     517           515         2     830
#> 2  2013     1     1     533           529         4     850
#> 3  2013     1     1     542           540         2     923
#> 4  2013     1     1     544           545        -1    1004
#> 5  2013     1     1     554           600        -6     812
#> 6  2013     1     1     554           558        -4     740
#> 7  2013     1     1     555           600        -5     913
#> 8  2013     1     1     557           600        -3     709
#> 9  2013     1     1     557           600        -3     838
#> 10 2013     1     1     558           600        -2     753
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```

arrange() and desc()

```
arrange(flights, desc(arr_delay))
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>      <dbl>   <int>
#> 1  2013     1     9     641           900        1301    1242
#> 2  2013     6    15    1432          1935        1137    1607
#> 3  2013     1    10    1121          1635        1126    1239
#> 4  2013     9    20    1139          1845        1014    1457
#> 5  2013     7    22     845          1600        1005    1044
#> 6  2013     4    10    1100          1900         960    1342
#> 7  2013     3    17    2321           810         911     135
#> 8  2013     7    22    2257           759         898     121
#> 9  2013    12     5     756          1700         896    1058
#> 10 2013     5     3    1133          2055         878    1250
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```

arrange() and missing values

```
df <- tibble(x = c(5, NA, 2))
arrange(df, x)
#> # A tibble: 3 x 1
#>       x
#>   <dbl>
#> 1     2
#> 2     5
#> 3    NA
arrange(df, desc(x))
#> # A tibble: 3 x 1
#>       x
#>   <dbl>
#> 1     5
#> 2     2
#> 3    NA
```

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

Select columns with `select()`

```
select(flights, year, month, day)
#> # A tibble: 336,776 x 3
#>   year month   day
#>   <int> <int> <int>
#> 1  2013     1     1
#> 2  2013     1     1
#> 3  2013     1     1
#> 4  2013     1     1
#> 5  2013     1     1
#> 6  2013     1     1
#> 7  2013     1     1
#> 8  2013     1     1
#> 9  2013     1     1
#> 10 2013     1     1
#> # ... with 336,766 more rows
```


All columns between year and day

```
select(flights, year:day)
#> # A tibble: 336,776 x 3
#>   year month   day
#>   <int> <int> <int>
#> 1  2013     1     1
#> 2  2013     1     1
#> 3  2013     1     1
#> 4  2013     1     1
#> 5  2013     1     1
#> 6  2013     1     1
#> 7  2013     1     1
#> 8  2013     1     1
#> 9  2013     1     1
#> 10 2013     1     1
#> # ... with 336,766 more rows
```

All columns except from year to day

```
select(flights, -(year:day))
#> # A tibble: 336,776 x 16
#>   dep_time sched_dep_time dep_delay arr_time sched_arr_time
#>   <int>         <int>         <dbl>    <int>         <int>
#> 1      517           515           2      830           819
#> 2      533           529           4      850           830
#> 3      542           540           2      923           850
#> 4      544           545          -1     1004          1022
#> 5      554           600          -6      812           837
#> 6      554           558          -4      740           728
#> 7      555           600          -5      913           854
#> 8      557           600          -3      709           723
#> 9      557           600          -3      838           846
#> 10     558           600          -2      753           745
#> # ... with 336,766 more rows, and 11 more variables:
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#> #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

select() and everything()

```
select(flights, time_hour, air_time, everything())
#> # A tibble: 336,776 x 19
#>   time_hour          air_time year month  day dep_time
#>   <dtm>          <dbl> <int> <int> <int>   <int>
#> 1 2013-01-01 05:00:00    227  2013     1     1    517
#> 2 2013-01-01 05:00:00    227  2013     1     1    533
#> 3 2013-01-01 05:00:00    160  2013     1     1    542
#> 4 2013-01-01 05:00:00    183  2013     1     1    544
#> 5 2013-01-01 06:00:00    116  2013     1     1    554
#> 6 2013-01-01 05:00:00    150  2013     1     1    554
#> 7 2013-01-01 06:00:00    158  2013     1     1    555
#> 8 2013-01-01 06:00:00     53  2013     1     1    557
#> 9 2013-01-01 06:00:00    140  2013     1     1    557
#> 10 2013-01-01 06:00:00    138  2013     1     1    558
#> # ... with 336,766 more rows, and 13 more variables:
#> #   sched_dep_time <int>, dep_delay <dbl>, arr_time <int>,
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>
```

- Helper functions you can use within `select()`:
 - ▶ `starts_with("abc")`: matches names that begin with "abc".
 - ▶ `ends_with("xyz")`: matches names that end with "xyz".
 - ▶ `contains("ijk")`: matches names that contain "ijk".
 - ▶ `matches("(.)\\1")`: selects variables that match a regular expression (this one matches any variables that contain repeated characters).
 - ▶ `num_range("x", 1:3)` matches `x1`, `x2` and `x3`.
- `select()` can be used to rename variables, but it drops all of the variables not explicitly mentioned. Instead, use `rename()`
- See `?select` for more details.

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

Create a narrower dataset

```
(flights_sml <- select(flights,
  ends_with("delay"),
  distance,
  air_time))
```

#> # A tibble: 336,776 x 4

#>	dep_delay	arr_delay	distance	air_time
#>	<dbl>	<dbl>	<dbl>	<dbl>
#> 1	2	11	1400	227
#> 2	4	20	1416	227
#> 3	2	33	1089	160
#> 4	-1	-18	1576	183
#> 5	-6	-25	762	116
#> 6	-4	12	719	150
#> 7	-5	19	1065	158
#> 8	-3	-14	229	53
#> 9	-3	-8	944	140
#> 10	-2	8	733	138

#> # ... with 336,766 more rows

Add new variables with mutate()

```
mutate(flights_sml,  
  gain = arr_delay - dep_delay,  
  speed = distance / air_time * 60)  
#> # A tibble: 336,776 x 6  
#>   dep_delay arr_delay distance air_time gain speed  
#>   <dbl>     <dbl>     <dbl>   <dbl> <dbl> <dbl>  
#> 1         2         11    1400    227     9  370.  
#> 2         4         20    1416    227    16  374.  
#> 3         2         33    1089    160    31  408.  
#> 4        -1        -18    1576    183   -17  517.  
#> 5        -6        -25     762    116   -19  394.  
#> 6        -4         12     719    150    16  288.  
#> 7        -5         19    1065    158    24  404.  
#> 8        -3        -14     229     53   -11  259.  
#> 9        -3         -8     944    140    -5  405.  
#> 10       -2          8     733    138    10  319.  
#> # ... with 336,766 more rows
```

Refer to columns just created

```
mutate(flights_sml,  
  gain = arr_delay - dep_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours)  
#> # A tibble: 336,776 x 7  
#>   dep_delay arr_delay distance air_time gain hours gain_per_hour  
#>   <dbl>     <dbl>     <dbl>   <dbl> <dbl> <dbl>      <dbl>  
#> 1         2         11    1400    227     9 3.78        2.38  
#> 2         4         20    1416    227    16 3.78        4.23  
#> 3         2         33    1089    160    31 2.67       11.6  
#> 4        -1        -18    1576    183   -17 3.05       -5.57  
#> 5        -6        -25     762    116   -19 1.93       -9.83  
#> 6        -4         12     719    150    16 2.5         6.4  
#> 7        -5         19    1065    158    24 2.63         9.11  
#> 8        -3        -14     229     53   -11 0.883      -12.5  
#> 9        -3         -8     944    140    -5 2.33        -2.14  
#> 10       -2         8      733    138    10 2.3         4.35  
#> # ... with 336,766 more rows
```



```
transmute(flights,  
  gain = arr_delay - dep_delay,  
  hours = air_time / 60,  
  gain_per_hour = gain / hours)  
#> # A tibble: 336,776 x 3  
#>   gain hours gain_per_hour  
#>   <dbl> <dbl>         <dbl>  
#> 1     9  3.78           2.38  
#> 2    16  3.78           4.23  
#> 3    31  2.67          11.6  
#> 4   -17  3.05          -5.57  
#> 5   -19  1.93          -9.83  
#> 6    16  2.5           6.4  
#> 7    24  2.63          9.11  
#> 8   -11  0.883         -12.5  
#> 9    -5  2.33          -2.14  
#> 10   10  2.3           4.35  
#> # ... with 336,766 more rows
```

Any vectorized function would work, but frequently useful are:

- Arithmetic operators: `+`, `-`, `*`, `/`, `^`.
 - ▶ Vectorized with “recycling rules” (e.g., `air_time / 60`).
 - ▶ Useful in conjunction with aggregate functions (e.g., `x / sum(x)` or `y - mean(y)`).
- Modular arithmetic: `%/%` (integer division) and `%%` (remainder), where `x == y * (x %/% y) + (x %% y)`.
 - ▶ Allows you to break integers up into pieces (e.g., `hour = dep_time %/% 100` and `minute = dep_time %% 100`)
- Logs: `log()`, `log2()`, `log10()`.
 - ▶ Useful for data ranging across multiple orders of magnitude.
 - ▶ Convert multiplicative relationships to additive.

- Offsets: `lead()` and `lag()`:
 - ▶ Refer to lead-/lagging values (e.g., compute running differences $x - \text{lag}(x)$ or find values change $x \neq \text{lag}(x)$).

```
x <- 1:10
lag(x)
#> [1] NA 1 2 3 4 5 6 7 8 9
lead(x)
#> [1] 2 3 4 5 6 7 8 9 10 NA
```

- Cumulative aggregates: `cumsum()`, `cumprod()`, `cummin()`, `cummax()`, `cummean()`.

```
cumsum(x)
#> [1] 1 3 6 10 15 21 28 36 45 55
cummean(x)
#> [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5
```

- Logical comparisons, <, <=, >, >=, !=
- Ranking functions: min_rank(), row_number(), dense_rank(), percent_rank(), cume_dist(), ntile()

```
y <- c(1, 2, 2, NA, 3, 4)
min_rank(y)
#> [1] 1 2 2 NA 4 5
min_rank(desc(y))
#> [1] 5 3 3 NA 2 1
row_number(y)
#> [1] 1 2 3 NA 4 5
dense_rank(y)
#> [1] 1 2 2 NA 3 4
percent_rank(y)
#> [1] 0.00 0.25 0.25 NA 0.75 1.00
cume_dist(y)
#> [1] 0.2 0.6 0.6 NA 0.8 1.0
```

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

Collapse values with `summarize()`

```
summarize(flights, delay = mean(dep_delay, na.rm = TRUE))  
#> # A tibble: 1 x 1  
#>   delay  
#>   <dbl>  
#> 1  12.6
```

summarize() paired with group_by()

```
by_day <- group_by(flights, year, month, day)
summarize(by_day, delay = mean(dep_delay, na.rm = TRUE))
#> `summarise()` has grouped output by 'year', 'month'. You can override using the `.groups` argument.
#> # A tibble: 365 x 4
#> # Groups:   year, month [12]
#>   year month   day delay
#>   <int> <int> <int> <dbl>
#> 1  2013     1     1  11.5
#> 2  2013     1     2  13.9
#> 3  2013     1     3  11.0
#> 4  2013     1     4   8.95
#> 5  2013     1     5   5.73
#> 6  2013     1     6   7.15
#> 7  2013     1     7   5.42
#> 8  2013     1     8   2.55
#> 9  2013     1     9   2.28
#> 10 2013     1    10   2.84
#> # ... with 355 more rows
```

- To suppress the summarize info

```
options(dplyr.summarise.inform = FALSE)
```

An alternative to na.rm: pre-filter

```
not_cancelled <- flights %>% filter(!is.na(dep_delay))
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(mean = mean(dep_delay))
#> # A tibble: 365 x 4
#> # Groups:   year, month [12]
#>   year month   day mean
#>   <int> <int> <int> <dbl>
#> 1  2013     1     1  11.5
#> 2  2013     1     2  13.9
#> 3  2013     1     3  11.0
#> 4  2013     1     4   8.95
#> 5  2013     1     5   5.73
#> 6  2013     1     6   7.15
#> 7  2013     1     7   5.42
#> 8  2013     1     8   2.55
#> 9  2013     1     9   2.28
#> 10 2013     1    10   2.84
#> # ... with 355 more rows
```


Useful summary functions I

- Measures of location: `mean()`, `median()`.
- Measures of spread: `sd()`, `IQR()`, `mad()`.
- Measures of rank: `min(x)`, `quantile(x, 0.25)`, `max(x)`.

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(first = min(dep_time), last = max(dep_time))
#> # A tibble: 365 x 5
#> # Groups:   year, month [12]
#>   year month   day first last
#>   <int> <int> <int> <int> <int>
#> 1  2013     1     1   517 2356
#> 2  2013     1     2    42 2354
#> 3  2013     1     3    32 2349
#> 4  2013     1     4    25 2358
#> 5  2013     1     5    14 2357
#> 6  2013     1     6    16 2355
#> 7  2013     1     7    49 2359
#> 8  2013     1     8   454 2351
#> 9  2013     1     9     2 2252
#> 10 2013     1    10     3 2320
#> # ... with 355 more rows
```

- Measures of position: `first(x)`, `nth(x, 2)`, `last(x)`.

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(first_dep = first(dep_time), last_dep = last(dep_time))
#> # A tibble: 365 x 5
#> # Groups:   year, month [12]
#>   year month   day first_dep last_dep
#>   <int> <int> <int>     <int>     <int>
#> 1  2013     1     1       517      2356
#> 2  2013     1     2        42      2354
#> 3  2013     1     3        32      2349
#> 4  2013     1     4        25      2358
#> 5  2013     1     5        14      2357
#> 6  2013     1     6        16      2355
#> 7  2013     1     7         49      2359
#> 8  2013     1     8       454      2351
#> 9  2013     1     9         2      2252
#> 10 2013     1    10         3      2320
#> # ... with 355 more rows
```

- Counts: `n(x)`, `sum(!is.na(x))`, `n_distinct(x)`.

```
not_cancelled %>%
  group_by(dest) %>%
  summarize(carriers = n_distinct(carrier)) %>%
  arrange(desc(carriers))
#> # A tibble: 104 x 2
#>   dest carriers
#>   <chr>     <int>
#> 1 ATL         7
#> 2 BOS         7
#> 3 CLT         7
#> 4 ORD         7
#> 5 TPA         7
#> 6 AUS         6
#> 7 DCA         6
#> 8 DTW         6
#> 9 IAD         6
#> 10 MSP        6
#> # ... with 94 more rows
```

- A simple helper function for counts:

```
not_cancelled %>% count(dest)
```

```
#> # A tibble: 104 x 2
```

```
#>   dest      n
```

```
#>   <chr> <int>
```

```
#> 1 ABQ    254
```

```
#> 2 ACK    265
```

```
#> 3 ALB    419
```

```
#> 4 ANC      8
```

```
#> 5 ATL  16898
```

```
#> 6 AUS   2418
```

```
#> 7 AVL    263
```

```
#> 8 BDL    412
```

```
#> 9 BGR    360
```

```
#> 10 BHM   272
```

```
#> # ... with 94 more rows
```

■ Counts with an optional weight variable:

```
not_cancelled %>% count(tailnum, wt = distance)
```

```
#> # A tibble: 4,037 x 2
```

```
#>   tailnum      n
```

```
#>   <chr>    <dbl>
```

```
#> 1 D942DN    3418
```

```
#> 2 NOEGMQ 240626
```

```
#> 3 N10156 110389
```

```
#> 4 N102UW 25722
```

```
#> 5 N103US 24619
```

```
#> 6 N104UW 25157
```

```
#> 7 N10575 141475
```

```
#> 8 N105UW 23618
```

```
#> 9 N107US 21677
```

```
#> 10 N108UW 32070
```

```
#> # ... with 4,027 more rows
```

- Counts of logical values: e.g., `sum(x > 10)`.

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(n_early = sum(dep_time < 500))
#> # A tibble: 365 x 4
#> # Groups:   year, month [12]
#>   year month   day n_early
#>   <int> <int> <int>   <int>
#> 1  2013     1     1       0
#> 2  2013     1     2       3
#> 3  2013     1     3       4
#> 4  2013     1     4       3
#> 5  2013     1     5       3
#> 6  2013     1     6       2
#> 7  2013     1     7       2
#> 8  2013     1     8       1
#> 9  2013     1     9       3
#> 10 2013     1    10       3
#> # ... with 355 more rows
```

- Proportions of logical values: e.g., `mean(y == 0)`.

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(hour_perc = mean(arr_delay > 60))
#> # A tibble: 365 x 4
#> # Groups:   year, month [12]
#>   year month   day hour_perc
#>   <int> <int> <int>     <dbl>
#> 1  2013     1     1      NA
#> 2  2013     1     2      NA
#> 3  2013     1     3      NA
#> 4  2013     1     4      NA
#> 5  2013     1     5    0.0349
#> 6  2013     1     6      NA
#> 7  2013     1     7    0.0333
#> 8  2013     1     8      NA
#> 9  2013     1     9      NA
#> 10 2013     1    10    0.0183
#> # ... with 355 more rows
```

Grouping by multiple variables I

```
daily <- group_by(flights, year, month, day)
(per_day <- summarize(daily, flights = n()))
#> # A tibble: 365 x 4
#> # Groups:   year, month [12]
#>   year month   day flights
#>   <int> <int> <int>   <int>
#> 1  2013     1     1     842
#> 2  2013     1     2     943
#> 3  2013     1     3     914
#> 4  2013     1     4     915
#> 5  2013     1     5     720
#> 6  2013     1     6     832
#> 7  2013     1     7     933
#> 8  2013     1     8     899
#> 9  2013     1     9     902
#> 10 2013     1    10     932
#> # ... with 355 more rows
```


Grouping by multiple variables II

```
(per_month <- summarize(per_day, flights = sum(flights)))  
#> # A tibble: 12 x 3  
#> # Groups:   year [1]  
#>   year month flights  
#>   <int> <int>   <int>  
#> 1  2013     1  27004  
#> 2  2013     2  24951  
#> 3  2013     3  28834  
#> 4  2013     4  28330  
#> 5  2013     5  28796  
#> 6  2013     6  28243  
#> 7  2013     7  29425  
#> 8  2013     8  29327  
#> 9  2013     9  27574  
#> 10 2013    10  28889  
#> 11 2013    11  27268  
#> 12 2013    12  28135  
(per_year <- summarize(per_month, flights = sum(flights)))  
#> # A tibble: 1 x 2  
#>   year flights  
#>   <int>   <int>  
#> 1  2013  336776
```

```
daily %>%  
  ungroup() %>%           # no longer grouped by date  
  summarize(flights = n()) # all flights  
#> # A tibble: 1 x 1  
#>   flights  
#>   <int>  
#> 1  336776
```

Grouped filters

```
(popular_dests <- flights %>%  
  group_by(dest) %>%  
  filter(n() > 365))  
#> # A tibble: 332,577 x 19  
#> # Groups:   dest [77]  
#>   year month   day dep_time sched_dep_time dep_delay arr_time  
#>   <int> <int> <int>   <int>         <int>      <dbl>   <int>  
#> 1  2013     1     1     517           515         2     830  
#> 2  2013     1     1     533           529         4     850  
#> 3  2013     1     1     542           540         2     923  
#> 4  2013     1     1     544           545        -1    1004  
#> 5  2013     1     1     554           600        -6     812  
#> 6  2013     1     1     554           558        -4     740  
#> 7  2013     1     1     555           600        -5     913  
#> 8  2013     1     1     557           600        -3     709  
#> 9  2013     1     1     557           600        -3     838  
#> 10 2013     1     1     558           600        -2     753  
#> # ... with 332,567 more rows, and 12 more variables:  
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,  
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#> #   time_hour <dtm>
```

```
popular_dests %>%
  filter(arr_delay > 0) %>%
  mutate(prop_delay = arr_delay / sum(arr_delay)) %>%
  select(year:day, dest, arr_delay, prop_delay)

#> # A tibble: 131,106 x 6
#> # Groups:   dest [77]
#>   year month   day dest  arr_delay prop_delay
#>   <int> <int> <int> <chr>    <dbl>    <dbl>
#> 1  2013     1     1 IAH      11  0.000111
#> 2  2013     1     1 IAH      20  0.000201
#> 3  2013     1     1 MIA      33  0.000235
#> 4  2013     1     1 ORD      12  0.0000424
#> 5  2013     1     1 FLL      19  0.0000938
#> 6  2013     1     1 ORD       8  0.0000283
#> 7  2013     1     1 LAX       7  0.0000344
#> 8  2013     1     1 DFW      31  0.000282
#> 9  2013     1     1 ATL      12  0.0000400
#> 10 2013     1     1 DTW      16  0.000116
#> # ... with 131,096 more rows
```

1 Tidy data

2 Filter

3 Arrange

4 Select


5 Mutate

6 Summarize

7 Relational data

- Until now: analysis of a single table of data.
- Typically: multiple tables of data to be combined.
 - ▶ Called **relational data**:
 - Because relations, not just the individual datasets, are important.
 - Relations are always defined for a pair of tables.
 - Relations of three or more tables are built from the relations between pairs.

- Common place to find relational data.
- Oracle, MySQL, Microsoft SQL Server, PostgreSQL, IBM DB2, Microsoft Access, SQLite, and others.



- All 336,776 flights that departed from NYC in 2013 (US BTS):

```
flights
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>      <dbl>   <int>
#> 1  2013     1     1     517           515         2     830
#> 2  2013     1     1     533           529         4     850
#> 3  2013     1     1     542           540         2     923
#> 4  2013     1     1     544           545        -1    1004
#> 5  2013     1     1     554           600        -6     812
#> 6  2013     1     1     554           558        -4     740
#> 7  2013     1     1     555           600        -5     913
#> 8  2013     1     1     557           600        -3     709
#> 9  2013     1     1     557           600        -3     838
#> 10 2013     1     1     558           600        -2     753
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```


airlines

```
#> # A tibble: 16 x 2
#>   carrier name
#>   <chr>     <chr>
#> 1 9E        Endeavor Air Inc.
#> 2 AA        American Airlines Inc.
#> 3 AS        Alaska Airlines Inc.
#> 4 B6        JetBlue Airways
#> 5 DL        Delta Air Lines Inc.
#> 6 EV        ExpressJet Airlines Inc.
#> 7 F9        Frontier Airlines Inc.
#> 8 FL        AirTran Airways Corporation
#> 9 HA        Hawaiian Airlines Inc.
#> 10 MQ       Envoy Air
#> 11 OO       SkyWest Airlines Inc.
#> 12 UA       United Air Lines Inc.
#> 13 US       US Airways Inc.
#> 14 VX       Virgin America
#> 15 WN       Southwest Airlines Co.
#> 16 YV       Mesa Airlines Inc.
```

airports

```
#> # A tibble: 1,458 x 8
```

#>	faa	name	lat	lon	alt	tz	dst	tzone
#>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<chr>	<chr>
#> 1	04G	Lansdowne Airport	41.1	-80.6	1044	-5	A	America/Ne~
#> 2	06A	Moton Field Muni~	32.5	-85.7	264	-6	A	America/Ch~
#> 3	06C	Schaumburg Regio~	42.0	-88.1	801	-6	A	America/Ch~
#> 4	06N	Randall Airport	41.4	-74.4	523	-5	A	America/Ne~
#> 5	09J	Jekyll Island Ai~	31.1	-81.4	11	-5	A	America/Ne~
#> 6	0A9	Elizabethton Mun~	36.4	-82.2	1593	-5	A	America/Ne~
#> 7	0G6	Williams County ~	41.5	-84.5	730	-5	A	America/Ne~
#> 8	0G7	Finger Lakes Reg~	42.9	-76.8	492	-5	A	America/Ne~
#> 9	0P2	Shoestring Aviat~	39.8	-76.6	1000	-5	U	America/Ne~
#> 10	0S9	Jefferson County~	48.1	-123.	108	-8	A	America/Lo~

#> # ... with 1,448 more rows

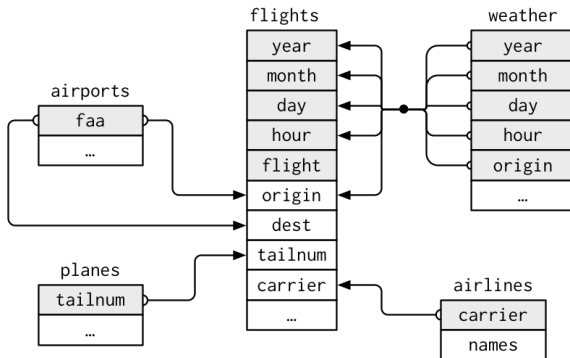
planes

```
#> # A tibble: 3,322 x 9
```


```
#>   tailnum year type manufacturer model engines seats speed engine
#>   <chr>   <int> <chr>   <chr>           <chr>   <int> <int> <int> <chr>
#> 1 N10156  2004 Fixed~ EMBRAER      EMB~      2    55    NA Turbo~
#> 2 N102UW  1998 Fixed~ AIRBUS      INDU~ A320~    2   182    NA Turbo~
#> 3 N103US  1999 Fixed~ AIRBUS      INDU~ A320~    2   182    NA Turbo~
#> 4 N104UW  1999 Fixed~ AIRBUS      INDU~ A320~    2   182    NA Turbo~
#> 5 N10575  2002 Fixed~ EMBRAER      EMB~      2    55    NA Turbo~
#> 6 N105UW  1999 Fixed~ AIRBUS      INDU~ A320~    2   182    NA Turbo~
#> 7 N107US  1999 Fixed~ AIRBUS      INDU~ A320~    2   182    NA Turbo~
#> 8 N108UW  1999 Fixed~ AIRBUS      INDU~ A320~    2   182    NA Turbo~
#> 9 N109UW  1999 Fixed~ AIRBUS      INDU~ A320~    2   182    NA Turbo~
#> 10 N110UW 1999 Fixed~ AIRBUS      INDU~ A320~    2   182    NA Turbo~
#> # ... with 3,312 more rows
```

weather

```
#> # A tibble: 26,115 x 15
#>   origin year month   day hour  temp  dewp humid wind_dir
#>   <chr>   <int> <int> <int> <int> <dbl> <dbl> <dbl>   <dbl>
#> 1 EWR     2013     1     1     1  39.0  26.1  59.4     270
#> 2 EWR     2013     1     1     2  39.0  27.0  61.6     250
#> 3 EWR     2013     1     1     3  39.0  28.0  64.4     240
#> 4 EWR     2013     1     1     4  39.9  28.0  62.2     250
#> 5 EWR     2013     1     1     5  39.0  28.0  64.4     260
#> 6 EWR     2013     1     1     6  37.9  28.0  67.2     240
#> 7 EWR     2013     1     1     7  39.0  28.0  64.4     240
#> 8 EWR     2013     1     1     8  39.9  28.0  62.2     250
#> 9 EWR     2013     1     1     9  39.9  28.0  62.2     260
#> 10 EWR    2013     1     1    10  41    28.0  59.6     260
#> # ... with 26,105 more rows, and 6 more variables: wind_speed <dbl>,
#> #   wind_gust <dbl>, precip <dbl>, pressure <dbl>, visib <dbl>,
#> #   time_hour <dtm>
```

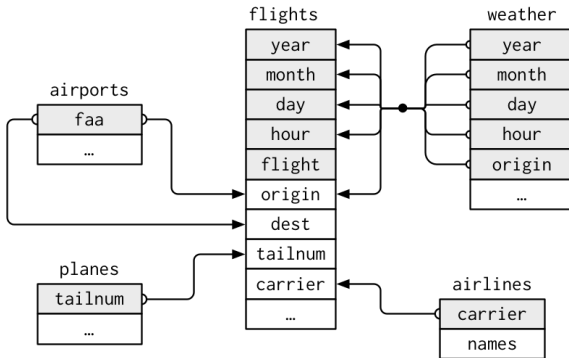


Exercise 1



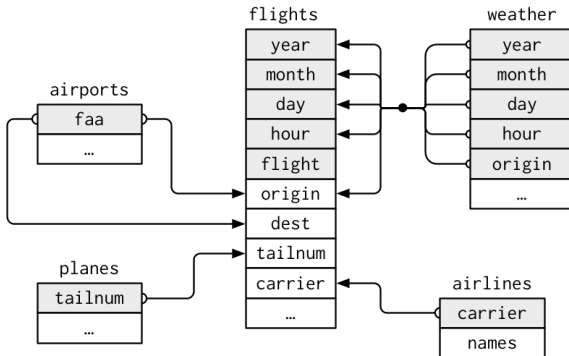
- Imagine you wanted to draw (approximately) the route each plane flies from its origin to its destination.
 - ▶ What variables would you need?
 - ▶ What tables would you need to combine?

Exercise 2



- I forgot to draw the relationship between weather and airports.
 - ▶ What is the relationship and how should it appear in the diagram?

Exercise 3



- weather only contains information for the origin (NYC) airports.
 - ▶ If it contained weather records for all airports in the USA, what additional relation would it define with flights?

- Keys:
 - ▶ Variables used to connect pair of tables.
 - ▶ Uniquely identifies an observation.
 - ▶ Can be:
 - A single variable (e.g., tailnum for planes).
 - Multiple variables (e.g., year, month, day, hour, and origin for weather).
- Two types of **keys**:
 - ▶ **Primary**: uniquely identifies an observation **in its own table**.
 - E.g., planes\$tailnum.
 - ▶ **Foreign**: uniquely identifies an observation **in another table**.
 - E.g., flights\$tailnum.
- Note that:
 - ▶ A variable can be both a primary key *and* a foreign key.
 - ▶ A primary key and the corresponding foreign key in another table form a **relation**.
 - ▶ Relations are typically one-to-many (e.g., flights and planes).

Is a given key primary?

```
planes %>%
  count(tailnum) %>%
  filter(n > 1)
#> # A tibble: 0 x 2
#> # ... with 2 variables: tailnum <chr>, n <int>

weather %>%
  count(year, month, day, hour, origin) %>%
  filter(n > 1)
#> # A tibble: 3 x 6
#>   year month   day hour origin      n
#>   <int> <int> <int> <int> <chr>  <int>
#> 1  2013    11     3     1  EWR      2
#> 2  2013    11     3     1  JFK      2
#> 3  2013    11     3     1  LGA      2
```

No explicit primary key?

```
flights %>%  
  count(year, month, day, flight) %>%  
  filter(n > 1)  
  
#> # A tibble: 29,768 x 5  
#>   year month   day flight     n  
#>   <int> <int> <int> <int> <int>  
#> 1  2013     1     1     1     2  
#> 2  2013     1     1     3     2  
#> 3  2013     1     1     4     2  
#> 4  2013     1     1    11     3  
#> 5  2013     1     1    15     2  
#> 6  2013     1     1    21     2  
#> 7  2013     1     1    27     4  
#> 8  2013     1     1    31     2  
#> 9  2013     1     1    32     2  
#> 10 2013     1     1    35     2  
#> # ... with 29,758 more rows
```

- Solution: add one with `mutate()` and `row_number()`.
- This is called a **surrogate key**.

- Two families of verbs to work with relational data:
 - ▶ **Mutating joins**
 - Add new variables to one data frame from matching observations in another.
 - ▶ **Filtering joins**
 - Filter observations from one data frame based on whether or not they match an observation in the other table.

Create a narrower dataset

```
flights2 <- flights %>%  
  select(year:day, hour, origin, dest, tailnum, carrier)
```

```
flights2
```

```
#> # A tibble: 336,776 x 8
```

```
#>   year month   day hour origin dest  tailnum carrier
```

```
#>   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr>
```

```
#> 1  2013     1     1     5 EWR   IAH   N14228 UA
```

```
#> 2  2013     1     1     5 LGA   IAH   N24211 UA
```

```
#> 3  2013     1     1     5 JFK   MIA   N619AA AA
```

```
#> 4  2013     1     1     5 JFK   BQN   N804JB B6
```

```
#> 5  2013     1     1     6 LGA   ATL   N668DN DL
```

```
#> 6  2013     1     1     5 EWR   ORD   N39463 UA
```

```
#> 7  2013     1     1     6 EWR   FLL   N516JB B6
```

```
#> 8  2013     1     1     6 LGA   IAD   N829AS EV
```

```
#> 9  2013     1     1     6 JFK   MCO   N593JB B6
```

```
#> 10 2013     1     1     6 LGA   ORD   N3ALAA AA
```

```
#> # ... with 336,766 more rows
```

A simple example


```
flights2 %>%
  select(-origin, -dest) %>%
  left_join(airlines, by = "carrier")
#> # A tibble: 336,776 x 7
#>   year month   day hour tailnum carrier name
#>   <int> <int> <int> <dbl> <chr>   <chr>   <chr>
#> 1  2013     1     1     5 N14228  UA      United Air Lines Inc.
#> 2  2013     1     1     5 N24211  UA      United Air Lines Inc.
#> 3  2013     1     1     5 N619AA  AA      American Airlines Inc.
#> 4  2013     1     1     5 N804JB  B6      JetBlue Airways
#> 5  2013     1     1     6 N668DN  DL      Delta Air Lines Inc.
#> 6  2013     1     1     5 N39463  UA      United Air Lines Inc.
#> 7  2013     1     1     6 N516JB  B6      JetBlue Airways
#> 8  2013     1     1     6 N829AS  EV      ExpressJet Airlines Inc.
#> 9  2013     1     1     6 N593JB  B6      JetBlue Airways
#> 10 2013     1     1     6 N3ALAA  AA      American Airlines Inc.
#> # ... with 336,766 more rows
```

Why mutating join?


```
flights2 %>%
  select(-origin, -dest) %>%
  mutate(name = airlines$name[match(carrier, airlines$carrier)])
#> # A tibble: 336,776 x 7
#>   year month   day hour tailnum carrier name
#>   <int> <int> <int> <dbl> <chr>   <chr>   <chr>
#> 1  2013     1     1     5 N14228 UA      United Air Lines Inc.
#> 2  2013     1     1     5 N24211 UA      United Air Lines Inc.
#> 3  2013     1     1     5 N619AA AA      American Airlines Inc.
#> 4  2013     1     1     5 N804JB B6      JetBlue Airways
#> 5  2013     1     1     6 N668DN DL      Delta Air Lines Inc.
#> 6  2013     1     1     5 N39463 UA      United Air Lines Inc.
#> 7  2013     1     1     6 N516JB B6      JetBlue Airways
#> 8  2013     1     1     6 N829AS EV      ExpressJet Airlines Inc.
#> 9  2013     1     1     6 N593JB B6      JetBlue Airways
#> 10 2013     1     1     6 N3ALAA AA      American Airlines Inc.
#> # ... with 336,766 more rows
```

Understanding mutating joins

```
x <- tribble(~key, ~val_x,  
            1, "x1",  
            2, "x2",  
            3, "x3")  
y <- tribble(~key, ~val_y,  
            1, "y1",  
            2, "y2",  
            4, "y3")
```



Inner join

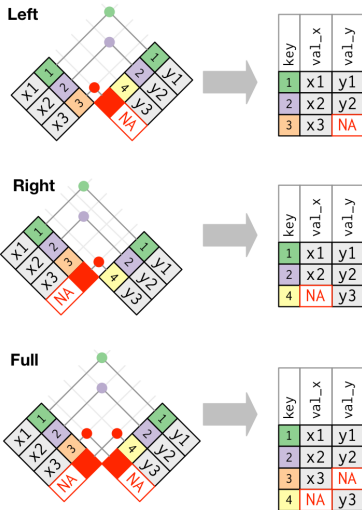


key	val_x	val_y
1	x1	y1
2	x2	y2





```
x %>%  
  inner_join(y, by = "key")  
#> # A tibble: 2 x 3  
#>   key val_x val_y  
#>   <dbl> <chr> <chr>  
#> 1     1   x1    y1  
#> 2     2   x2    y2
```

- **Outer joins** keep observations that appear in at least one of the tables:
 - ▶ **Left join:** keeps all observations in x .
 - ▶ **Right join:** keeps all observations in y .
 - ▶ **Full join:** keeps all observations in x and y
- They work by adding to each table an additional “virtual” observation which
 - ▶ has a key that always matches (if no other key matches),
 - ▶ and a value filled with NA.

Outer joins II



A Venn diagram for joins



- Two possibilities:
 - ▶ One table has duplicate keys.
 - Useful to add in additional information as there is typically a one-to-many relationship.
 - ▶ Both tables have duplicate keys.
 - Usually an error because in neither table do the keys uniquely identify an observation.
 - When you join duplicated keys, you get all possible combinations (i.e., the Cartesian product).

One table has duplicate keys


- Only x has duplicated keys:

```
x <- tribble(~key, ~val_x,  
            1, "x1",  
            2, "x2",  
            2, "x3",  
            1, "x4")
```

```
y <- tribble(~key, ~val_y,  
            1, "y1",  
            2, "y2")
```

- The join adds val_y to the matching rows:

```
left_join(x, y, by = "key")  
#> # A tibble: 4 x 3  
#>   key val_x val_y  
#>   <dbl> <chr> <chr>  
#> 1     1 x1    y1  
#> 2     2 x2    y2  
#> 3     2 x3    y2  
#> 4     1 x4    y1
```



val_x	key	val_y
x1	1	y1
x2	2	y2
x3	2	y2
x4	1	y1

Both tables have duplicate keys


- Both x and y have duplicated keys:

```
x <- tribble(~key, ~val_x,  
            1, "x1",  
            2, "x2",  
            2, "x3",  
            3, "x4")
```

```
y <- tribble(~key, ~val_y,  
            1, "y1",  
            2, "y2",  
            2, "y3",  
            3, "y4")
```

- The joint creates all combinations:

```
left_join(x, y, by = "key")  
#> # A tibble: 6 x 3  
#>   key val_x val_y  
#>   <dbl> <chr> <chr>  
#> 1     1 x1    y1  
#> 2     2 x2    y2  
#> 3     2 x2    y3  
#> 4     2 x3    y2  
#> 5     2 x3    y3  
#> 6     3 x4    y4
```



key	val_x	val_y
1	x1	y1
2	x2	y2
2	x2	y3
2	x3	y2
2	x3	y3
3	x4	y4

Defining the key columns

- Default uses all variables that appear in both tables.
- Called a **natural join**.

```
flights2 %>%  
  left_join(weather)  
#> Joining, by = c("year", "month", "day", "hour", "origin")  
#> # A tibble: 336,776 x 18  
#>   year month   day hour origin dest tailnum carrier temp dewp  
#>   <int> <int> <int> <dbl> <chr>  <chr> <chr>    <chr>   <dbl> <dbl>  
#> 1  2013     1     1     5 EWR   IAH   N14228 UA      39.0  28.0  
#> 2  2013     1     1     5 LGA   IAH   N24211 UA      39.9  25.0  
#> 3  2013     1     1     5 JFK   MIA   N619AA AA      39.0  27.0  
#> 4  2013     1     1     5 JFK   BQN   N804JB B6      39.0  27.0  
#> 5  2013     1     1     6 LGA   ATL   N668DN DL      39.9  25.0  
#> 6  2013     1     1     5 EWR   ORD   N39463 UA      39.0  28.0  
#> 7  2013     1     1     6 EWR   FLL   N516JB B6      37.9  28.0  
#> 8  2013     1     1     6 LGA   IAD   N829AS EV      39.9  25.0  
#> 9  2013     1     1     6 JFK   MCO   N593JB B6      37.9  27.0  
#> 10 2013     1     1     6 LGA   ORD   N3ALAA AA      39.9  25.0  
#> # ... with 336,766 more rows, and 8 more variables: humid <dbl>,  
#> #   wind_dir <dbl>, wind_speed <dbl>, wind_gust <dbl>, precip <dbl>,  
#> #   pressure <dbl>, visib <dbl>, time_hour <dtm>
```


- Like a natural join, but uses only some of the common variables:

```
flights2 %>%  
  left_join(planes, by = "tailnum")  
#> # A tibble: 336,776 x 16  
#>   year.x month   day hour origin dest tailnum carrier year.y type  
#>   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr> <int> <chr>  
#> 1  2013     1     1     5 EWR   IAH   N14228 UA      1999 Fixe~  
#> 2  2013     1     1     5 LGA   IAH   N24211 UA      1998 Fixe~  
#> 3  2013     1     1     5 JFK   MIA   N619AA AA       1990 Fixe~  
#> 4  2013     1     1     5 JFK   BQN   N804JB B6       2012 Fixe~  
#> 5  2013     1     1     6 LGA   ATL   N668DN DL       1991 Fixe~  
#> 6  2013     1     1     5 EWR   ORD   N39463 UA       2012 Fixe~  
#> 7  2013     1     1     6 EWR   FLL   N516JB B6       2000 Fixe~  
#> 8  2013     1     1     6 LGA   IAD   N829AS EV       1998 Fixe~  
#> 9  2013     1     1     6 JFK   MCO   N593JB B6       2004 Fixe~  
#> 10 2013     1     1     6 LGA   ORD   N3ALAA AA        NA <NA>  
#> # ... with 336,766 more rows, and 6 more variables:  
#> #   manufacturer <chr>, model <chr>, engines <int>, seats <int>,  
#> #   speed <int>, engine <chr>
```

Using a named character vector

- With `by = c("a" = "b")`, `left_join` matches variable `a` in table `x` to variable `b` in table `y`:

```
flights2 %>%  
  left_join(airports, c("dest" = "faa"))  
#> # A tibble: 336,776 x 15  
#>   year month   day hour origin dest tailnum carrier name      lat  
#>   <int> <int> <int> <dbl> <chr>  <chr> <chr>  <chr>  <chr>  <dbl>  
#> 1  2013     1     1     5 EWR   IAH   N14228 UA      George~ 30.0  
#> 2  2013     1     1     5 LGA   IAH   N24211 UA      George~ 30.0  
#> 3  2013     1     1     5 JFK   MIA   N619AA AA      Miami ~ 25.8  
#> 4  2013     1     1     5 JFK   BQN   N804JB B6      <NA>    NA  
#> 5  2013     1     1     6 LGA   ATL   N668DN DL      Hartsf~ 33.6  
#> 6  2013     1     1     5 EWR   ORD   N39463 UA      Chicag~ 42.0  
#> 7  2013     1     1     6 EWR   FLL   N516JB B6      Fort L~ 26.1  
#> 8  2013     1     1     6 LGA   IAD   N829AS EV      Washin~ 38.9  
#> 9  2013     1     1     6 JFK   MCO   N593JB B6      Orland~ 28.4  
#> 10 2013     1     1     6 LGA   ORD   N3ALAA AA      Chicag~ 42.0  
#> # ... with 336,766 more rows, and 5 more variables: lon <dbl>,  
#> #   alt <dbl>, tz <dbl>, dst <chr>, tzone <chr>
```

dplyr	SQL
<code>inner_join(x, y, by = "z")</code>	<code>SELECT * FROM x INNER JOIN y USING (z)</code>
<code>left_join(x, y, by = "z")</code>	<code>SELECT * FROM x LEFT OUTER JOIN y USING (z)</code>
<code>right_join(x, y, by = "z")</code>	<code>SELECT * FROM x RIGHT OUTER JOIN y USING (z)</code>
<code>full_join(x, y, by = "z")</code>	<code>SELECT * FROM x FULL OUTER JOIN y USING (z)</code>

■ Note that:

- ▶ “INNER” and “OUTER” are optional, and often omitted.
- ▶ Joining different variables between the tables uses a slightly different syntax in SQL.
 - E.g. `inner_join(x, y, by = c("a" = "b"))` vs `SELECT * FROM x INNER JOIN y ON x.a = y.b.`

- Similar to mutating joins, but affect the observations rather than the variables:
 - ▶ `semi_join(x, y)` **keeps** all observations in `x` that have a match in `y`.
 - Useful for matching filtered summary tables back to the original rows.
 - ▶ `anti_join(x, y)` **drops** all observations in `x` that have a match in `y`.
 - Useful for diagnosing join mismatches.

Flights that went to top destinations

```
top_dest <- flights %>%
  count(dest, sort = TRUE) %>%
  head(10)

flights %>%
  filter(dest %in% top_dest$dest) %>%
  print(n = 5)

#> # A tibble: 141,145 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>   <int>         <int>       <dbl>   <int>
#> 1  2013     1     1     542             540         2     923
#> 2  2013     1     1     554             600        -6     812
#> 3  2013     1     1     554             558        -4     740
#> 4  2013     1     1     555             600        -5     913
#> 5  2013     1     1     557             600        -3     838
#> # ... with 141,140 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dtm>
```


- How to extend to multiple variables?

- Only keeps rows in x having a match in y:


```
flights %>%  
  semi_join(top_dest)  
#> Joining, by = "dest"  
#> # A tibble: 141,145 x 19  
#>   year month   day dep_time sched_dep_time dep_delay arr_time  
#>   <int> <int> <int>   <int>         <int>         <dbl>   <int>  
#> 1  2013     1     1     542             540           2     923  
#> 2  2013     1     1     554             600          -6     812  
#> 3  2013     1     1     554             558          -4     740  
#> 4  2013     1     1     555             600          -5     913  
#> 5  2013     1     1     557             600          -3     838  
#> 6  2013     1     1     558             600          -2     753  
#> 7  2013     1     1     558             600          -2     924  
#> 8  2013     1     1     558             600          -2     923  
#> 9  2013     1     1     559             559           0     702  
#> 10 2013     1     1     600             600           0     851  
#> # ... with 141,135 more rows, and 12 more variables:  
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,  
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#> #   time_hour <dtm>
```

Visually understand the semi-join

■ One-to-many:




■ Many-to-many:



flights without a match in planes

```
flights %>%  
  anti_join(planes,  
            by = "tailnum") %>%  
  count(tailnum, sort = TRUE)  
#> # A tibble: 722 x 2  
#>   tailnum      n  
#>   <chr>   <int>  
#> 1 <NA>     2512  
#> 2 N725MQ    575  
#> 3 N722MQ    513  
#> 4 N723MQ    507  
#> 5 N713MQ    483  
#> 6 N735MQ    396  
#> 7 N0EGMQ    371  
#> 8 N534MQ    364  
#> 9 N542MQ    363  
#> 10 N531MQ   349  
#> # ... with 712 more rows
```



key	val_x
3	x3