

# DSFBA: Data Wrangling

*Data Science for Business Analytics*

Thibault Vatter

Department of Statistics, Columbia University

10/13/2021

# Outline

1 Tidy data

2 Filter


3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data



Most of the material (e.g., the picture above) is borrowed from

**R for data science**

# A grammar of data manipulation

- When working with data you must:
  - ▶ Figure out what you want to do.
  - ▶ Describe those tasks as a computer program.
  - ▶ Execute the program.
- The `dplyr` package makes this fast and easy with 5 verbs!
  - ▶ `filter()`: select observations based on their values.
  - ▶ `arrange()`: reorder the observations.
  - ▶ `select()`: select variables based on their names.
  - ▶ `mutate()`: add variables as functions of existing variables.
  - ▶ `summarize()`: collapse many values down to a single summary.
- Two important features:
  - ▶ Verbs can be used with `group_by()` to operate groupwise.
  - ▶ Verbs work similarly...
    1. First argument: a data frame.
    2. Other arguments: what to do with it using variable names.
    3. Result: a new data frame.

All 336,776 flights that departed from NYC in 2013 ([US BTS](#)):

```
nycflights13::flights
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>     <int>           <int>     <dbl>     <int>
#> 1 2013     1     1      517            515        2     830
#> 2 2013     1     1      533            529        4     850
#> 3 2013     1     1      542            540        2     923
#> 4 2013     1     1      544            545       -1    1004
#> 5 2013     1     1      554            600       -6     812
#> 6 2013     1     1      554            558       -4     740
#> 7 2013     1     1      555            600       -5     913
#> 8 2013     1     1      557            600       -3     709
#> 9 2013     1     1      557            600       -3     838
#> 10 2013    1     1      558            600       -2     753
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dttm>
```

# What is this code doing?

```
a1 <- group_by(flights, year, month, day)
a2 <- select(a1, arr_delay, dep_delay)
a3 <- summarize(a2,
                 arr = mean(arr_delay, na.rm = TRUE),
                 dep = mean(dep_delay, na.rm = TRUE))
filter(a3, arr > 30 | dep > 30)
#> # A tibble: 49 x 5
#> # Groups:   year, month [11]
#>   year month   day   arr   dep
#>   <int> <int> <int> <dbl> <dbl>
#> 1 2013     1     16  34.2  24.6
#> 2 2013     1     31  32.6  28.7
#> 3 2013     2     11  36.3  39.1
#> 4 2013     2     27  31.3  37.8
#> 5 2013     3      8  85.9  83.5
#> 6 2013     3     18  41.3  30.1
#> 7 2013     4     10  38.4  33.0
#> 8 2013     4     12  36.0  34.8
#> 9 2013     4     18  36.0  34.9
#> 10 2013    4     19  47.9  46.1
#> # ... with 39 more rows
```

# Same code (no unnecessary objects)

```
filter(summarize(select(group_by(flights, year, month, day),
                        arr_delay, dep_delay),
                 arr = mean(arr_delay, na.rm = TRUE),
                 dep = mean(dep_delay, na.rm = TRUE)),
       arr > 30 | dep > 30)
#> # A tibble: 49 x 5
#> # Groups:   year, month [11]
#>   year month   day   arr   dep
#>   <int> <int> <int> <dbl> <dbl>
#> 1 2013     1    16  34.2  24.6
#> 2 2013     1    31  32.6  28.7
#> 3 2013     2    11  36.3  39.1
#> 4 2013     2    27  31.3  37.8
#> 5 2013     3     8  85.9  83.5
#> 6 2013     3    18  41.3  30.1
#> 7 2013     4    10  38.4  33.0
#> 8 2013     4    12  36.0  34.8
#> 9 2013     4    18  36.0  34.9
#> 10 2013    4    19  47.9  46.1
#> # ... with 39 more rows
```

# ... Or use %>%

```
flights %>%
  group_by(year, month, day) %>%
  select(arr_delay, dep_delay) %>%
  summarize(arr = mean(arr_delay, na.rm = TRUE),
            dep = mean(dep_delay, na.rm = TRUE)) %>%
  filter(arr > 30 | dep > 30)
#> # A tibble: 49 x 5
#> # Groups:   year, month [11]
#>   year month   day   arr   dep
#>   <int> <int> <int> <dbl> <dbl>
#> 1 2013     1     1   34.2  24.6
#> 2 2013     1     31   32.6  28.7
#> 3 2013     2     11   36.3  39.1
#> 4 2013     2     27   31.3  37.8
#> 5 2013     3      8   85.9  83.5
#> 6 2013     3     18   41.3  30.1
#> 7 2013     4     10   38.4  33.0
#> 8 2013     4     12   36.0  34.8
#> 9 2013     4     18   36.0  34.9
#> 10 2013    4     19   47.9  46.1
#> # ... with 39 more rows
```

# Basic piping

- $x \%>% f$  is equivalent to  $f(x)$
- $x \%>% f(y)$  is equivalent to  $f(x, y)$
- $x \%>% f(y) \%>% g(z)$  is equivalent to  $g(f(x, y), z)$

```
x <- 1:10
y <- x + 1
z <- y + 1
f <- function(x, y) x + y

x %>% sum
#> [1] 55
x %>% f(y)
#> [1] 3 5 7 9 11 13 15 17 19 21
x %>% f(y) %>% f(z)
#> [1] 6 9 12 15 18 21 24 27 30 33
```

# The argument (“dot”) placeholder

- $x \%>% f(y, .)$  is equivalent to  $f(y, x)$
- $x \%>% f(y, z = .)$  is equivalent to  $f(y, z = x)$

```
x <- 1:10
y <- 2 * x
f <- function(z, y) y / z

x %>% f(y, .)
#> [1] 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5 0.5
x %>% f(y, z = .)
#> [1] 2 2 2 2 2 2 2 2 2 2
```

# Function composition

- Each of the three options has its own strengths and weaknesses:
  - ▶ Nesting,  $f(g(x))$ :
    - Concise, and well suited for short sequences.
    - Longer sequences harder to read (inside out & right to left).
    - Arguments can get spread out over long distances ([see Dagwood sandwich](#)).
  - ▶ Intermediate objects,  $y \leftarrow f(x); g(y)$ :
    - Requires you to name intermediate objects.
    - A strength when objects are important, but a weakness when values are truly intermediate.
  - ▶ Piping,  $x \%>% f() \%>% g()$ :
    - Allows to read code in straightforward left-to-right fashion.
    - Doesn't require to name intermediate objects.
    - Only for linear sequences of transformations of a single object.
- Most code use a combination of all three styles, but...
- **Piping is more common in data analysis code!**

# Outline

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

# Tidy data

---

*"Happy families are all alike; every unhappy family is unhappy in its own way."* — Leo Tolstoy

*"Tidy datasets are all alike, but every messy dataset is messy in its own way."* — Hadley Wickham

To learn more about the underlying theory, see the [Tidy Data paper](#).

# Which representation is “best”?

## ■ First representation?

```
table1
#> # A tibble: 6 x 4
#>   country     year  cases population
#>   <chr>      <int> <int>      <int>
#> 1 Afghanistan 1999    745  19987071
#> 2 Afghanistan 2000   2666  20595360
#> 3 Brazil       1999  37737 172006362
#> 4 Brazil       2000  80488 174504898
#> 5 China        1999 212258 1272915272
#> 6 China        2000 213766 1280428583
```

## ■ Second representation?

```
table2
#> # A tibble: 12 x 4
#>   country     year type     count
#>   <chr>      <int> <chr>    <int>
#> 1 Afghanistan 1999 cases     745
#> 2 Afghanistan 1999 population 19987071
#> 3 Afghanistan 2000 cases     2666
#> 4 Afghanistan 2000 population 20595360
#> 5 Brazil       1999 cases     37737
#> 6 Brazil       1999 population 172006362
#> 7 Brazil       2000 cases     80488
#> 8 Brazil       2000 population 174504898
#> 9 China        1999 cases     212258
#> 10 China       1999 population 1272915272
#> 11 China       2000 cases     213766
#> 12 China       2000 population 1280428583
```

## ■ Third representation?

```
table3
#> # A tibble: 6 x 3
#>   country     year rate
#>   <chr>      <int> <chr>
#> 1 Afghanistan 1999 745/19987071
#> 2 Afghanistan 2000 2666/20595360
#> 3 Brazil       1999 37737/172006362
#> 4 Brazil       2000 80488/174504898
#> 5 China        1999 212258/1272915272
#> 6 China        2000 213766/1280428583
```


## ■ Fourth representation?

```
table4a # cases
#> # A tibble: 3 x 3
#>   country `1999` `2000`
#>   <chr>    <int>   <int>
#> 1 Afghanistan    745    2666
#> 2 Brazil          37737   80488
#> 3 China           212258  213766
table4b # population
#> # A tibble: 3 x 3
#>   country      `1999`   `2000`
#>   <chr>        <int>    <int>
#> 1 Afghanistan  19987071 20595360
#> 2 Brazil       172006362 174504898
#> 3 China        1272915272 1280428583
```

# What makes a dataset tidy?

Three interrelated rules:

- Each variable must have its own column.
- Each observation must have its own row.
- Each value must have its own cell.



Because it's impossible to only satisfy two of the three:

- Put each dataset in a tibble.
- Put each variable in a column.

# Why ensure that your data is tidy?

- Why?
  - ▶ With consistent data structure, it's easier to learn the tools that work with it because they have an underlying uniformity.
  - ▶ Placing variables in columns allows R's vectorized nature to shine.
- Tidy data principles seem obvious, BUT:
  - ▶ Most people aren't familiar with them.
  - ▶ Data often organized to facilitate something different than analysis.
- Hence, you'll most likely need to do some tidying.

# The two steps of tidying

- Figure out what the variables and observations are.
  - Resolve one of two common problems:
    - ▶ One variable might be spread across multiple columns.
    - ▶ One observation might be scattered across multiple rows.
- ... To fix these problems, you'll need `pivot_longer()` and `pivot_wider()`.

# Longer with pivot\_wider()

table4a

```
#> # A tibble: 3 x 3
#>   country     `1999` `2000`
#>   <chr>       <int>  <int>
#> 1 Afghanistan    745    2666
#> 2 Brazil        37737   80488
#> 3 China         212258  213766
```

table4a %>%

```
pivot_longer(c(`1999`, `2000`),
             names_to = "year",
             values_to = "cases")
#> # A tibble: 6 x 3
#>   country     year   cases
#>   <chr>      <chr>  <int>
#> 1 Afghanistan 1999    745
#> 2 Afghanistan 2000   2666
#> 3 Brazil      1999   37737
#> 4 Brazil      2000   80488
#> 5 China       1999  212258
#> 6 China       2000  213766
```

country	year	cases	country	1999	2000
Afghanistan	1999	745	Afghanistan	745	2666
Afghanistan	2000	2666	Brazil	37737	80488
Brazil	1999	37737	China	212258	213766
Brazil	2000	80488			
China	1999	212258			
China	2000	213766			

table4

# Wider with pivot\_wider()

table2

```
#> # A tibble: 12 x 4
#>   country     year type     count
#>   <chr>      <int> <chr>    <int>
#> 1 Afghanistan 1999 cases     745
#> 2 Afghanistan 1999 population 19987071
#> 3 Afghanistan 2000 cases     2666
#> 4 Afghanistan 2000 population 20595360
#> 5 Brazil      1999 cases     37737
#> 6 Brazil      1999 population 172006362
#> 7 Brazil      2000 cases     80488
#> 8 Brazil      2000 population 174504898
#> 9 China       1999 cases     212258
#> 10 China      1999 population 1272915272
#> 11 China      2000 cases     213766
#> 12 China      2000 population 1280428583
```

table2 %>%

```
pivot_wider(names_from = type,
            values_from = count)
#> # A tibble: 6 x 4
#>   country     year cases population
#>   <chr>      <int> <int>      <int>
#> 1 Afghanistan 1999  745  19987071
#> 2 Afghanistan 2000  2666 20595360
#> 3 Brazil      1999  37737 172006362
#> 4 Brazil      2000  80488 174504898
#> 5 China       1999  212258 1272915272
#> 6 China       2000  213766 1280428583
```

country	year	key	value	country	year	cases	population
Afghanistan	1999	cases	745	Afghanistan	1999	745	19987071
Afghanistan	1999	population	19987071	Afghanistan	2000	2666	20595360
Afghanistan	2000	cases	2666	Brazil	1999	37737	172006362
Afghanistan	2000	population	20595360	Brazil	2000	80488	174504898
Brazil	1999	cases	37737	China	1999	212258	1272915272
Brazil	1999	population	172006362	China	2000	213766	1280428583
Brazil	2000	cases	80488				
Brazil	2000	population	174504898				
China	1999	cases	212258				
China	1999	population	1272915272				
China	2000	cases	213766				
China	2000	population	1280428583				

table2

# Separate a column with separate()

```
table3
#> # A tibble: 6 x 3
#>   country     year    rate
#>   <chr>      <int> <chr>
#> 1 Afghanistan 1999 745/19987071
#> 2 Afghanistan 2000 2666/20595360
#> 3 Brazil      1999 37737/172006362
#> 4 Brazil      2000 80488/174504898
#> 5 China       1999 212258/1272915272
#> 6 China       2000 213766/1280428583
```

```
table3 %>% separate(rate,
                      into = c("cases",
                               "population"))
#> # A tibble: 6 x 4
#>   country     year  cases population
#>   <chr>      <int> <chr>    <chr>
#> 1 Afghanistan 1999  745     19987071
#> 2 Afghanistan 2000  2666    20595360
#> 3 Brazil      1999  37737   172006362
#> 4 Brazil      2000  80488   174504898
#> 5 China       1999  212258  1272915272
#> 6 China       2000  213766  1280428583
```

country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20595360
Brazil	1999	37737	172006362
Brazil	2000	80488	174504898
China	1999	212258	1272915272
China	2000	213766	1280428583

table3

# separate() using convert = TRUE

```
table3 %>%  
  separate(rate, into = c("cases", "population"), convert = TRUE)  
#> # A tibble: 6 x 4  
#>   country      year  cases population  
#>   <chr>        <int> <int>     <int>  
#> 1 Afghanistan  1999    745  19987071  
#> 2 Afghanistan  2000   2666  20595360  
#> 3 Brazil       1999  37737  172006362  
#> 4 Brazil       2000  80488  174504898  
#> 5 China        1999 212258 1272915272  
#> 6 China        2000 213766 1280428583
```

# Unite two columns with unite()

table5

```
#> # A tibble: 6 x 4
#>   country  century year  rate
#>   <chr>     <chr>  <chr> <chr>
#> 1 Afghanistan 19      99    745/19987071
#> 2 Afghanistan 20      00    2666/20595360
#> 3 Brazil       19      99    37737/172006362
#> 4 Brazil       20      00    80488/174504898
#> 5 China        19      99    212258/1272915272
#> 6 China        20      00    213766/1280428583
```

table5 %>%

```
unite(new, century, year, sep = "")
#> # A tibble: 6 x 3
#>   country  new  rate
#>   <chr>    <chr> <chr>
#> 1 Afghanistan 1999 745/19987071
#> 2 Afghanistan 2000 2666/20595360
#> 3 Brazil      1999 37737/172006362
#> 4 Brazil      2000 80488/174504898
#> 5 China       1999 212258/1272915272
#> 6 China       2000 213766/1280428583
```



country	year	rate
Afghanistan	1999	745 / 19987071
Afghanistan	2000	2666 / 20595360
Brazil	1999	37737 / 172006362
Brazil	2000	80488 / 174504898
China	1999	212258 / 1272915272
China	2000	213766 / 1280428583

country	century	year	rate
Afghanistan	19	99	745 / 19987071
Afghanistan	20	0	2666 / 20595360
Brazil	19	99	37737 / 172006362
Brazil	20	0	80488 / 174504898
China	19	99	212258 / 1272915272
China	20	0	213766 / 1280428583

table6

# Missing values and tidy data

- A value can be missing in one of two possible ways:
  - ▶ **Explicitly**, i.e. flagged with NA.
  - ▶ **Implicitly**, i.e. simply not present in the data.

*"An explicit missing value is the presence of an absence; an implicit missing value is the absence of a presence." Hadley Wickham*

- Are there missing values in this dataset?

```
stocks <- tibble(  
  year = c(2015, 2015, 2015, 2015, 2016, 2016, 2016),  
  qtr = c(1, 2, 3, 4, 2, 3, 4),  
  return = c(1.88, 0.59, 0.35, NA, 0.92, 0.17, 2.66)  
)
```

# Implicit to explicit and conversely

- Implicit to explicit by pivoting:

```
stocks %>%  
  pivot_wider(names_from = year,  
             values_from = return)  
#> # A tibble: 4 x 3  
#>   qtr `2015` `2016`  
#>   <dbl> <dbl> <dbl>  
#> 1     1    1.88  NA  
#> 2     2    0.59  0.92  
#> 3     3    0.35  0.17  
#> 4     4    NA    2.66
```

- Implicit to explicit using complete:

```
stocks %>% complete(year, qtr)  
#> # A tibble: 8 x 3  
#>   year   qtr return  
#>   <dbl> <dbl> <dbl>  
#> 1  2015    1    1.88  
#> 2  2015    2    0.59  
#> 3  2015    3    0.35  
#> 4  2015    4    NA  
#> 5  2016    1    NA  
#> 6  2016    2    0.92  
#> 7  2016    3    0.17  
#> 8  2016    4    2.66
```

- Explicit to implicit via drop\_na().

# Fill in missing values with fill()

```
treatment <- tribble(  
  ~ person,                  ~ treatment, ~response,  
  "Derrick Whitmore", 1,      7,  
  NA,                      2,      10,  
  NA,                      3,      9,  
  "Katherine Burke", 1,      4  
)  
treatment %>%  
  fill(person)  
#> # A tibble: 4 x 3  
#>   person        treatment response  
#>   <chr>          <dbl>     <dbl>  
#> 1 Derrick Whitmore     1         7  
#> 2 Derrick Whitmore     2        10  
#> 3 Derrick Whitmore     3         9  
#> 4 Katherine Burke      1         4
```

# Outline

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

# Filter rows with filter()

```
filter(flights, month == 1, day == 1)
#> # A tibble: 842 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>     <int>           <int>    <dbl>     <int>
#> 1 2013     1     1      517            515       2     830
#> 2 2013     1     1      533            529       4     850
#> 3 2013     1     1      542            540       2     923
#> 4 2013     1     1      544            545      -1    1004
#> 5 2013     1     1      554            600      -6     812
#> 6 2013     1     1      554            558      -4     740
#> 7 2013     1     1      555            600      -5     913
#> 8 2013     1     1      557            600      -3     709
#> 9 2013     1     1      557            600      -3     838
#> 10 2013    1     1      558            600      -2     753
#> # ... with 832 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dttm>
```

# Comparisons

- The standard suite: `>`, `>=`, `<`, `<=`, `!=`, and `==`.
- Most common mistake:

```
filter(flights, month = 1)
```


- What happens in the following?

```
sqrt(2) ^ 2 == 2
#> [1] FALSE
1/49 * 49 == 1
#> [1] FALSE
near(sqrt(2) ^ 2, 2)
#> [1] TRUE
near(1 / 49 * 49, 1)
#> [1] TRUE
```


# Logical operators

Multiple arguments to `filter()` are combined with:


- `&` for “and”
- `|` for “or”
- `!` for “not”




$$y \& !x$$




$$x$$




$$x | y$$




$$x \& y$$



$$\text{xor}(x, y)$$



$$x \& !y$$



$$y$$

# What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

# What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

- Literally “finds all flights that departed in November or December”.
- ... filter(flights, month == 11 | 12) ?

# What is this code doing?

```
filter(flights, month == 11 | month == 12)
```

- Literally “finds all flights that departed in November or December”.
- ... filter(flights, month == 11 | 12) ?
- No, but a solution:

```
filter(flights, month %in% c(11, 12))
```

# De Morgan's law

- $!(x \ \& \ y)$  is the same as  $!x \ \mid \ !y$
- $!(x \ \mid \ y)$  is the same as  $!x \ \& \ !y$

```
all.equal(
  filter(flights, !(arr_delay > 120 | dep_delay > 120)),
  filter(flights, arr_delay <= 120, dep_delay <= 120)
)
#> [1] TRUE
```

# Missing values and filter()

```
df <- tibble(x = c(1, NA, 3))
filter(df, x > 1)
#> # A tibble: 1 x 1
#>   x
#>   <dbl>
#> 1     3
filter(df, is.na(x) | x > 1)
#> # A tibble: 2 x 1
#>   x
#>   <dbl>
#> 1    NA
#> 2     3
```

# Outline

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

# Arrange rows with arrange()

```
arrange(flights, year, month, day)
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>     <int>          <int>     <dbl>     <int>
#> 1 2013     1     1      517            515        2       830
#> 2 2013     1     1      533            529        4       850
#> 3 2013     1     1      542            540        2       923
#> 4 2013     1     1      544            545       -1      1004
#> 5 2013     1     1      554            600       -6       812
#> 6 2013     1     1      554            558       -4       740
#> 7 2013     1     1      555            600       -5       913
#> 8 2013     1     1      557            600       -3       709
#> 9 2013     1     1      557            600       -3       838
#> 10 2013    1     1      558            600       -2       753
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dttm>
```

# arrange() and desc()

```
arrange(flights, desc(arr_delay))
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>    <int>          <int>    <dbl>    <int>
#> 1 2013     1     9      641           900    1301    1242
#> 2 2013     6    15     1432          1935    1137    1607
#> 3 2013     1    10     1121          1635    1126    1239
#> 4 2013     9    20     1139          1845    1014    1457
#> 5 2013     7    22      845          1600    1005    1044
#> 6 2013     4    10     1100          1900     960    1342
#> 7 2013     3    17     2321          810     911     135
#> 8 2013     7    22     2257          759     898     121
#> 9 2013    12     5      756          1700     896    1058
#> 10 2013    5     3     1133          2055     878    1250
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dttm>
```

# arrange() and missing values

```
df <- tibble(x = c(5, NA, 2))
arrange(df, x)
#> # A tibble: 3 x 1
#>       x
#>   <dbl>
#> 1     2
#> 2     5
#> 3    NA
arrange(df, desc(x))
#> # A tibble: 3 x 1
#>       x
#>   <dbl>
#> 1     5
#> 2     2
#> 3    NA
```

# Outline

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

# Select columns with select()

```
select(flights, year, month, day)
#> # A tibble: 336,776 x 3
#>   year month   day
#>   <int> <int> <int>
#> 1 2013     1     1
#> 2 2013     1     1
#> 3 2013     1     1
#> 4 2013     1     1
#> 5 2013     1     1
#> 6 2013     1     1
#> 7 2013     1     1
#> 8 2013     1     1
#> 9 2013     1     1
#> 10 2013    1     1
#> # ... with 336,766 more rows
```

# All columns between year and day

```
select(flights, year:day)
#> # A tibble: 336,776 x 3
#>   year month   day
#>   <int> <int> <int>
#> 1 2013     1     1
#> 2 2013     1     1
#> 3 2013     1     1
#> 4 2013     1     1
#> 5 2013     1     1
#> 6 2013     1     1
#> 7 2013     1     1
#> 8 2013     1     1
#> 9 2013     1     1
#> 10 2013    1     1
#> # ... with 336,766 more rows
```

# All columns except from year to day

```
select(flights, -(year:day))
#> # A tibble: 336,776 x 16
#>   dep_time sched_dep_time dep_delay arr_time sched_arr_time
#>   <int>        <int>     <dbl>      <int>        <int>
#> 1 517          515       2         830        819
#> 2 533          529       4         850        830
#> 3 542          540       2         923        850
#> 4 544          545      -1        1004       1022
#> 5 554          600      -6        812        837
#> 6 554          558      -4        740        728
#> 7 555          600      -5        913        854
#> 8 557          600      -3        709        723
#> 9 557          600      -3        838        846
#> 10 558         600      -2        753        745
#> # ... with 336,766 more rows, and 11 more variables:
#> #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
#> #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#> #   hour <dbl>, minute <dbl>, time_hour <dttm>
```

# select() and everything()

```
select(flights, time_hour, air_time, everything())
#> # A tibble: 336,776 x 19
#>   time_hour           air_time year month day dep_time
#>   <dttm>             <dbl>  <int> <int> <int>    <int>
#> 1 2013-01-01 05:00:00     227  2013     1     1      517
#> 2 2013-01-01 05:00:00     227  2013     1     1      533
#> 3 2013-01-01 05:00:00     160  2013     1     1      542
#> 4 2013-01-01 05:00:00     183  2013     1     1      544
#> 5 2013-01-01 06:00:00     116  2013     1     1      554
#> 6 2013-01-01 05:00:00     150  2013     1     1      554
#> 7 2013-01-01 06:00:00     158  2013     1     1      555
#> 8 2013-01-01 06:00:00      53  2013     1     1      557
#> 9 2013-01-01 06:00:00     140  2013     1     1      557
#> 10 2013-01-01 06:00:00     138  2013     1     1      558
#> # ... with 336,766 more rows, and 13 more variables:
#> #   sched_dep_time <int>, dep_delay <dbl>, arr_time <int>,
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   distance <dbl>, hour <dbl>, minute <dbl>
```

# More on select

- Helper functions you can use within `select()`:
  - ▶ `starts_with("abc")`: matches names that begin with “abc”.
  - ▶ `ends_with("xyz")`: matches names that end with “xyz”.
  - ▶ `contains("ijk")`: matches names that contain “ijk”.
  - ▶ `matches("(.)\\1")`: selects variables that match a regular expression (this one matches any variables that contain repeated characters).
  - ▶ `num_range("x", 1:3)` matches x1, x2 and x3.
- `select()` can be used to rename variables, but it drops all of the variables not explicitly mentioned. Instead, use `rename()`
- See `?select` for more details.

# Outline

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

# Create a narrower dataset

```
(flights_sml <- select(flights,
  ends_with("delay"),
  distance,
  air_time))

#> # A tibble: 336,776 x 4
#>   dep_delay arr_delay distance air_time
#>   <dbl>     <dbl>     <dbl>     <dbl>
#> 1 2           11       1400      227
#> 2 4           20       1416      227
#> 3 2           33       1089      160
#> 4 -1          -18      1576      183
#> 5 -6          -25      762       116
#> 6 -4          12       719       150
#> 7 -5          19       1065      158
#> 8 -3          -14      229       53
#> 9 -3          -8       944       140
#> 10 -2          8       733       138
#> # ... with 336,766 more rows
```

# Add new variables with `mutate()`

```
mutate(flights_sml,
  gain = arr_delay - dep_delay,
  speed = distance / air_time * 60)
#> # A tibble: 336,776 x 6
#>   dep_delay arr_delay distance air_time   gain speed
#>   <dbl>     <dbl>    <dbl>     <dbl>   <dbl> <dbl>
#> 1      2       11     1400      227      9  370.
#> 2      4       20     1416      227     16  374.
#> 3      2       33     1089      160     31  408.
#> 4     -1      -18     1576      183     -17  517.
#> 5     -6      -25      762      116     -19  394.
#> 6     -4      -12      719      150      16  288.
#> 7     -5      -19     1065      158      24  404.
#> 8     -3      -14      229       53     -11  259.
#> 9     -3       -8      944      140      -5  405.
#> 10    -2        8      733      138      10  319.
#> # ... with 336,766 more rows
```

# Refer to columns just created

```
mutate(flights_sml,
  gain = arr_delay - dep_delay,
  hours = air_time / 60,
  gain_per_hour = gain / hours)
#> # A tibble: 336,776 x 7
#>   dep_delay arr_delay distance air_time   gain hours gain_per_hour
#>   <dbl>     <dbl>    <dbl>    <dbl>    <dbl> <dbl>    <dbl>
#> 1      2        11     1400     227      9 3.78     2.38
#> 2      4        20     1416     227     16 3.78     4.23
#> 3      2        33     1089     160     31 2.67    11.6 
#> 4     -1       -18     1576     183    -17 3.05   -5.57
#> 5     -6       -25      762     116    -19 1.93   -9.83
#> 6     -4       12      719     150     16 2.5      6.4 
#> 7     -5       19     1065     158     24 2.63    9.11
#> 8     -3       -14     229      53    -11 0.883   -12.5
#> 9     -3       -8      944     140     -5 2.33   -2.14
#> 10    -2        8      733     138     10 2.3      4.35
#> # ... with 336,766 more rows
```

# transmute()

```
transmute(flights,
  gain = arr_delay - dep_delay,
  hours = air_time / 60,
  gain_per_hour = gain / hours)
#> # A tibble: 336,776 x 3
#>   gain    hours gain_per_hour
#>   <dbl>   <dbl>      <dbl>
#> 1  9     3.78       2.38
#> 2  16    3.78       4.23
#> 3  31    2.67      11.6
#> 4 -17   3.05      -5.57
#> 5 -19   1.93      -9.83
#> 6  16    2.5        6.4
#> 7  24    2.63      9.11
#> 8 -11   0.883     -12.5
#> 9  -5   2.33      -2.14
#> 10 10    2.3        4.35
#> # ... with 336,766 more rows
```

# Useful creation functions I

Any vectorized function would work, but frequently useful are:

- Arithmetic operators: `+`, `-`, `*`, `/`, `^`.
  - ▶ Vectorized with “recycling rules” (e.g., `air_time / 60`).
  - ▶ Useful in conjunction with aggregate functions (e.g., `x / sum(x)` or `y - mean(y)`).
- Modular arithmetic: `%/%` (integer division) and `%%` (remainder), where  $x == y * (x \%/% y) + (x \% \% y)$ .
  - ▶ Allows you to break integers up into pieces (e.g., `hour = dep_time %/% 100` and `minute = dep_time %% 100`)
- Logs: `log()`, `log2()`, `log10()`.
  - ▶ Useful for data ranging across multiple orders of magnitude.
  - ▶ Convert multiplicative relationships to additive.

# Useful creation functions II

- Offsets: `lead()` and `lag()`:

- ▶ Refer to lead-/lagging values (e.g., compute running differences  $x - \text{lag}(x)$  or find values change  $x \neq \text{lag}(x)$ ).

```
x <- 1:10
lag(x)
#> [1] NA  1  2  3  4  5  6  7  8  9
lead(x)
#> [1]  2  3  4  5  6  7  8  9 10 NA
```

- Cumulative aggregates: `cumsum()`, `cumprod()`, `cummin()`, `cummax()`, `cummean()`.

```
cumsum(x)
#> [1]  1  3  6 10 15 21 28 36 45 55
cummean(x)
#> [1] 1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0 5.5
```

# Useful creation functions III

- Logical comparisons, <, <=, >, >=, !=
- Ranking functions: min\_rank(), row\_number(),  
dense\_rank(), percent\_rank(), cume\_dist(), ntile()

```
y <- c(1, 2, 2, NA, 3, 4)
min_rank(y)
#> [1] 1 2 2 NA 4 5
min_rank(desc(y))
#> [1] 5 3 3 NA 2 1
row_number(y)
#> [1] 1 2 3 NA 4 5
dense_rank(y)
#> [1] 1 2 2 NA 3 4
percent_rank(y)
#> [1] 0.00 0.25 0.25    NA 0.75 1.00
cume_dist(y)
#> [1] 0.2 0.6 0.6  NA 0.8 1.0
```

# Outline

1 Tidy data

2 Filter

3 Arrange

4 Select

5 Mutate

6 Summarize

7 Relational data

# Collapse values with summarize()

```
summarize(flights, delay = mean(dep_delay, na.rm = TRUE))  
#> # A tibble: 1 x 1  
#>   delay  
#>   <dbl>  
#> 1 12.6
```

# summarize() paired with group\_by()

```
flights %>%  
  group_by(year, month, day) %>%  
  summarize(delay = mean(dep_delay, na.rm = TRUE))  
#> `summarise()` has grouped output by 'year', 'month'. You can override using t  
#> # A tibble: 365 x 4  
#> # Groups: year, month [12]  
#>   year month   day delay  
#>   <int> <int> <int> <dbl>  
#> 1 2013     1     1  11.5  
#> 2 2013     1     2  13.9  
#> 3 2013     1     3  11.0  
#> 4 2013     1     4  8.95  
#> 5 2013     1     5  5.73  
#> 6 2013     1     6  7.15  
#> 7 2013     1     7  5.42  
#> 8 2013     1     8  2.55  
#> 9 2013     1     9  2.28  
#> 10 2013    1    10  2.84  
#> # ... with 355 more rows
```

- To suppress the summarize info

```
options(dplyr.summarise.inform = FALSE)
```

# An alternative to na.rm: pre-filter

```
not_cancelled <- flights %>%
  filter(!is.na(dep_delay))

not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(mean = mean(dep_delay))
#> # A tibble: 365 x 4
#> # Groups:   year, month [12]
#>   year month   day   mean
#>   <int> <int> <int> <dbl>
#> 1 2013     1     1  11.5
#> 2 2013     1     2  13.9
#> 3 2013     1     3  11.0
#> 4 2013     1     4  8.95
#> 5 2013     1     5  5.73
#> 6 2013     1     6  7.15
#> 7 2013     1     7  5.42
#> 8 2013     1     8  2.55
#> 9 2013     1     9  2.28
#> 10 2013    1    10  2.84
#> # ... with 355 more rows
```

# Useful summary functions I

- Measures of location: `mean()`, `median()`.
- Measures of spread: `sd()`, `IQR()`, `mad()`.
- Measures of rank: `min(x)`, `quantile(x, 0.25)`, `max(x)`.

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(first = min(dep_time), last = max(dep_time))
#> # A tibble: 365 x 5
#> # Groups:   year, month [12]
#>   year month   day first  last
#>   <int> <int> <int> <int> <int>
#> 1 2013     1     1   517  2356
#> 2 2013     1     2    42  2354
#> 3 2013     1     3    32  2349
#> 4 2013     1     4    25  2358
#> 5 2013     1     5    14  2357
#> 6 2013     1     6    16  2355
#> 7 2013     1     7    49  2359
#> 8 2013     1     8   454  2351
#> 9 2013     1     9     2  2252
#> 10 2013    1    10     3  2320
#> # ... with 355 more rows
```

# Useful summary functions II

- Measures of position: `first(x)`, `nth(x, 2)`, `last(x)`.

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(first_dep = first(dep_time), last_dep = last(dep_time))
#> # A tibble: 365 x 5
#> # Groups:   year, month [12]
#>   year month   day first_dep last_dep
#>   <int> <int> <int>     <int>    <int>
#> 1 2013     1     1      517    2356
#> 2 2013     1     2      42     2354
#> 3 2013     1     3      32     2349
#> 4 2013     1     4      25     2358
#> 5 2013     1     5      14     2357
#> 6 2013     1     6      16     2355
#> 7 2013     1     7      49     2359
#> 8 2013     1     8      454    2351
#> 9 2013     1     9       2     2252
#> 10 2013    1    10       3     2320
#> # ... with 355 more rows
```

# Useful summary functions III

- Counts: `n(x)`, `sum(!is.na(x))`, `n_distinct(x)`.

```
not_cancelled %>%
  group_by(dest) %>%
  summarize(carriers = n_distinct(carrier)) %>%
  arrange(desc(carriers))
#> # A tibble: 104 x 2
#>   dest    carriers
#>   <chr>    <int>
#> 1 ATL        7
#> 2 BOS        7
#> 3 CLT        7
#> 4 ORD        7
#> 5 TPA        7
#> 6 AUS        6
#> 7 DCA        6
#> 8 DTW        6
#> 9 IAD        6
#> 10 MSP       6
#> # ... with 94 more rows
```

# Useful summary functions IV

- A simple helper function for counts:

```
not_cancelled %>% count(dest)
#> # A tibble: 104 x 2
#>   dest      n
#>   <chr> <int>
#> 1 ABQ     254
#> 2 ACK     265
#> 3 ALB     419
#> 4 ANC      8
#> 5 ATL    16898
#> 6 AUS     2418
#> 7 AVL     263
#> 8 BDL     412
#> 9 BGR     360
#> 10 BHM    272
#> # ... with 94 more rows
```

# Useful summary functions V

- Counts with an optional weight variable:

```
not_cancelled %>% count(tailnum, wt = distance)
#> # A tibble: 4,037 x 2
#>   tailnum      n
#>   <chr>     <dbl>
#> 1 D942DN     3418
#> 2 NOEGMQ    240626
#> 3 N10156    110389
#> 4 N102UW    25722
#> 5 N103US    24619
#> 6 N104UW    25157
#> 7 N10575    141475
#> 8 N105UW    23618
#> 9 N107US    21677
#> 10 N108UW   32070
#> # ... with 4,027 more rows
```

# Useful summary functions VI

- Counts of logical values: e.g., `sum(x > 10)`.

```
not_cancelled %>%  
  group_by(year, month, day) %>%  
  summarize(n_early = sum(dep_time < 500))  
#> # A tibble: 365 x 4  
#>   year month   day n_early  
#>   <int> <int> <int>    <int>  
#> 1 2013     1     1      0  
#> 2 2013     1     2      3  
#> 3 2013     1     3      4  
#> 4 2013     1     4      3  
#> 5 2013     1     5      3  
#> 6 2013     1     6      2  
#> 7 2013     1     7      2  
#> 8 2013     1     8      1  
#> 9 2013     1     9      3  
#> 10 2013    1    10      3  
#> # ... with 355 more rows
```

# Useful summary functions VII

- Proportions of logical values: e.g., `mean(y == 0)`.

```
not_cancelled %>%
  group_by(year, month, day) %>%
  summarize(hour_perc = mean(arr_delay > 60, na.rm = TRUE))
#> # A tibble: 365 x 4
#>   year month   day hour_perc
#>   <int> <int> <int>     <dbl>
#> 1 2013     1     1  0.0722
#> 2 2013     1     2  0.0851
#> 3 2013     1     3  0.0567
#> 4 2013     1     4  0.0396
#> 5 2013     1     5  0.0349
#> 6 2013     1     6  0.0470
#> 7 2013     1     7  0.0333
#> 8 2013     1     8  0.0213
#> 9 2013     1     9  0.0202
#> 10 2013    1    10  0.0183
#> # ... with 355 more rows
```

# Grouping by multiple variables I

```
daily <- group_by(flights, year, month, day)
(per_day  <- summarize(daily, flights = n()))
#> # A tibble: 365 x 4
#> # Groups:   year, month [12]
#>   year month   day flights
#>   <int> <int> <int>   <int>
#> 1 2013     1     1     842
#> 2 2013     1     2     943
#> 3 2013     1     3     914
#> 4 2013     1     4     915
#> 5 2013     1     5     720
#> 6 2013     1     6     832
#> 7 2013     1     7     933
#> 8 2013     1     8     899
#> 9 2013     1     9     902
#> 10 2013    1    10     932
#> # ... with 355 more rows
```

# Grouping by multiple variables II

```
(per_month <- summarize(per_day, flights = sum(flights)))
#> # A tibble: 12 x 3
#> # Groups:   year [1]
#>   year month flights
#>   <int> <int>    <int>
#> 1 2013     1    27004
#> 2 2013     2    24951
#> 3 2013     3    28834
#> 4 2013     4    28330
#> 5 2013     5    28796
#> 6 2013     6    28243
#> 7 2013     7    29425
#> 8 2013     8    29327
#> 9 2013     9    27574
#> 10 2013    10    28889
#> 11 2013    11    27268
#> 12 2013    12    28135
(per_year <- summarize(per_month, flights = sum(flights)))
#> # A tibble: 1 x 2
#>   year flights
#>   <int>    <int>
#> 1 2013    336776
```

# Ungrouping

```
daily %>%
  ungroup() %>%                      # no longer grouped by date
  summarize(flights = n())   # all flights
#> # A tibble: 1 x 1
#>   flights
#>   <int>
#> 1 336776
```

# Grouped filters

```
(popular_dests <- flights %>%
  group_by(dest) %>%
  filter(n() > 365))
#> # A tibble: 332,577 x 19
#> # Groups: dest [77]
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>     <int>          <int>    <dbl>     <int>
#> 1 2013     1     1      517            515        2     830
#> 2 2013     1     1      533            529        4     850
#> 3 2013     1     1      542            540        2     923
#> 4 2013     1     1      544            545       -1    1004
#> 5 2013     1     1      554            600       -6     812
#> 6 2013     1     1      554            558       -4     740
#> 7 2013     1     1      555            600       -5     913
#> 8 2013     1     1      557            600       -3     709
#> 9 2013     1     1      557            600       -3     838
#> 10 2013    1     1      558            600       -2     753
#> # ... with 332,567 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dttm>
```

# Grouped mutates

```
popular_dests %>%  
  filter(arr_delay > 0) %>%  
  mutate(prop_delay = arr_delay / sum(arr_delay)) %>%  
  select(year:day, dest, arr_delay, prop_delay)  
#> # A tibble: 131,106 x 6  
#> # Groups: dest [77]  
#>   year month day dest arr_delay prop_delay  
#>   <int> <int> <int> <chr>     <dbl>      <dbl>  
#> 1 2013     1     1 IAH        11 0.000111  
#> 2 2013     1     1 IAH        20 0.000201  
#> 3 2013     1     1 MIA        33 0.000235  
#> 4 2013     1     1 ORD        12 0.0000424  
#> 5 2013     1     1 FLL        19 0.0000938  
#> 6 2013     1     1 ORD         8 0.0000283  
#> 7 2013     1     1 LAX         7 0.0000344  
#> 8 2013     1     1 DFW        31 0.000282  
#> 9 2013     1     1 ATL        12 0.0000400  
#> 10 2013    1     1 DTW        16 0.000116  
#> # ... with 131,096 more rows
```

# Outline

1 Tidy data

2 Filter

3 Arrange

4 Select


5 Mutate

6 Summarize

7 Relational data

- Until now: analysis of a single table of data.
- Typically: multiple tables of data to be combined.
  - ▶ Called **relational data**:
    - Because relations, not just the individual datasets, are important.
    - Relations are always defined for a pair of tables.
    - Relations of three or more tables are built from the relations between pairs.

- Common place to find relational data.
- Oracle, MySQL, Microsoft SQL Server, PostgreSQL, IBM DB2, Microsoft Access, SQLite, and others.



- All 336,776 flights that departed from NYC in 2013 ([US BTS](#)):

```
flights
#> # A tibble: 336,776 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>     <int>           <int>     <dbl>     <int>
#> 1 2013     1     1      517            515        2     830
#> 2 2013     1     1      533            529        4     850
#> 3 2013     1     1      542            540        2     923
#> 4 2013     1     1      544            545       -1    1004
#> 5 2013     1     1      554            600       -6     812
#> 6 2013     1     1      554            558       -4     740
#> 7 2013     1     1      555            600       -5     913
#> 8 2013     1     1      557            600       -3     709
#> 9 2013     1     1      557            600       -3     838
#> 10 2013    1     1      558            600       -2     753
#> # ... with 336,766 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dttm>
```

## airlines

```
#> # A tibble: 16 x 2
#>   carrier name
#>   <chr>   <chr>
#> 1 9E      Endeavor Air Inc.
#> 2 AA      American Airlines Inc.
#> 3 AS      Alaska Airlines Inc.
#> 4 B6      JetBlue Airways
#> 5 DL      Delta Air Lines Inc.
#> 6 EV      ExpressJet Airlines Inc.
#> 7 F9      Frontier Airlines Inc.
#> 8 FL      AirTran Airways Corporation
#> 9 HA      Hawaiian Airlines Inc.
#> 10 MQ     Envoy Air
#> 11 OO     SkyWest Airlines Inc.
#> 12 UA     United Air Lines Inc.
#> 13 US     US Airways Inc.
#> 14 VX     Virgin America
#> 15 WN     Southwest Airlines Co.
#> 16 YV     Mesa Airlines Inc.
```

# nycflights13::airports


```
airports
#> # A tibble: 1,458 x 8
#>   faa     name          lat    lon    alt    tz dst  tzone
#>   <chr>  <chr>        <dbl>  <dbl>  <dbl>  <dbl> <chr> <chr>
#> 1 04G  Lansdowne Airport  41.1  -80.6  1044  -5 A America/Ne-
#> 2 06A  Moton Field Muni~  32.5  -85.7   264  -6 A America/Ch-
#> 3 06C  Schaumburg Regio~  42.0  -88.1   801  -6 A America/Ch-
#> 4 06N  Randall Airport    41.4  -74.4   523  -5 A America/Ne-
#> 5 09J  Jekyll Island Ai~  31.1  -81.4    11  -5 A America/Ne-
#> 6 0A9  Elizabethton Mun~  36.4  -82.2  1593  -5 A America/Ne-
#> 7 0G6  Williams County ~  41.5  -84.5   730  -5 A America/Ne-
#> 8 0G7  Finger Lakes Reg~  42.9  -76.8   492  -5 A America/Ne-
#> 9 0P2  Shoestring Aviat~  39.8  -76.6  1000  -5 U America/Ne-
#> 10 0S9 Jefferson County~ 48.1  -123.    108  -8 A America/Lo-
#> # ... with 1,448 more rows
```

## planes


```
#> # A tibble: 3,322 x 9
#>   tailnum  year type  manufacturer model engines seats speed engine
#>   <chr>    <int> <chr>       <chr>     <chr>    <int> <int> <int> <chr>
#> 1 N10156  2004 Fixed~ EMBRAER    EMB~-    2      55    NA Turbo~
#> 2 N102UW  1998 Fixed~ AIRBUS INDU~ A320~    2     182    NA Turbo~
#> 3 N103US  1999 Fixed~ AIRBUS INDU~ A320~    2     182    NA Turbo~
#> 4 N104UW  1999 Fixed~ AIRBUS INDU~ A320~    2     182    NA Turbo~
#> 5 N10575  2002 Fixed~ EMBRAER    EMB~~    2      55    NA Turbo~
#> 6 N105UW  1999 Fixed~ AIRBUS INDU~ A320~    2     182    NA Turbo~
#> 7 N107US  1999 Fixed~ AIRBUS INDU~ A320~    2     182    NA Turbo~
#> 8 N108UW  1999 Fixed~ AIRBUS INDU~ A320~    2     182    NA Turbo~
#> 9 N109UW  1999 Fixed~ AIRBUS INDU~ A320~    2     182    NA Turbo~
#> 10 N110UW 1999 Fixed~ AIRBUS INDU~ A320~   2     182    NA Turbo~
#> # ... with 3,312 more rows
```

## weather

```
#> # A tibble: 26,115 x 15
#>   origin year month day hour temp  dewp humid wind_dir
#>   <chr>  <int> <int> <int> <dbl> <dbl> <dbl> <dbl>
#> 1 EWR    2013     1     1     1  39.0  26.1  59.4   270
#> 2 EWR    2013     1     1     2  39.0  27.0  61.6   250
#> 3 EWR    2013     1     1     3  39.0  28.0  64.4   240
#> 4 EWR    2013     1     1     4  39.9  28.0  62.2   250
#> 5 EWR    2013     1     1     5  39.0  28.0  64.4   260
#> 6 EWR    2013     1     1     6  37.9  28.0  67.2   240
#> 7 EWR    2013     1     1     7  39.0  28.0  64.4   240
#> 8 EWR    2013     1     1     8  39.9  28.0  62.2   250
#> 9 EWR    2013     1     1     9  39.9  28.0  62.2   260
#> 10 EWR   2013     1     1    10  41    28.0  59.6   260
#> # ... with 26,105 more rows, and 6 more variables: wind_speed <dbl>,
#> #   wind_gust <dbl>, precip <dbl>, pressure <dbl>, visib <dbl>,
#> #   time_hour <dttm>
```




# Exercise 1




- Imagine you wanted to draw (approximately) the route each plane flies from its origin to its destination.
  - ▶ What variables would you need?
  - ▶ What tables would you need to combine?

# Exercise 2



- I forgot to draw the relationship between `weather` and `airports`.
  - ▶ What is the relationship and how should it appear in the diagram?

# Exercise 3



- **weather** only contains information for the origin (NYC) airports.
  - ▶ If it contained weather records for all airports in the USA, what additional relation would it define with **flights**?

# Keys

- Keys:
  - ▶ Variables used to connect pair of tables.
  - ▶ Uniquely identifies an observation.
  - ▶ Can be:
    - A single variable (e.g., `tailnum` for planes).
    - Multiple variables (e.g., `year`, `month`, `day`, `hour`, and `origin` for weather).
- Two types of **keys**:
  - ▶ **Primary**: uniquely identifies an observation **in its own table**.
    - E.g., `planes$tailnum`.
  - ▶ **Foreign**: uniquely identifies an observation **in another table**.
    - E.g., `flights$tailnum`.
- Note that:
  - ▶ A variable can be both a primary key *and* a foreign key.
  - ▶ A primary key and the corresponding foreign key in another table form a **relation**.
  - ▶ Relations are typically one-to-many (e.g., flights and planes).

# Is a given key primary?

```
planes %>%
  count(tailnum) %>%
  filter(n > 1)
#> # A tibble: 0 x 2
#> # ... with 2 variables: tailnum <chr>, n <int>

weather %>%
  count(year, month, day, hour, origin) %>%
  filter(n > 1)
#> # A tibble: 3 x 6
#>   year month   day hour origin     n
#>   <int> <int> <int> <int> <chr> <int>
#> 1 2013    11     3     1 EWR      2
#> 2 2013    11     3     1 JFK      2
#> 3 2013    11     3     1 LGA      2
```

# No explicit primary key?

```
flights %>%
  count(year, month, day, flight) %>%
  filter(n > 1)
#> # A tibble: 29,768 x 5
#>   year month   day flight     n
#>   <int> <int> <int> <int> <int>
#> 1 2013     1     1     1     2
#> 2 2013     1     1     3     2
#> 3 2013     1     1     4     2
#> 4 2013     1     1    11     3
#> 5 2013     1     1    15     2
#> 6 2013     1     1    21     2
#> 7 2013     1     1    27     4
#> 8 2013     1     1    31     2
#> 9 2013     1     1    32     2
#> 10 2013    1     1    35     2
#> # ... with 29,758 more rows
```

- Solution: add one with `mutate()` and `row_number()`.
- This is called a **surrogate key**.

- Two families of verbs to work with relational data:
  - ▶ **Mutating joins**
    - Add new variables to one data frame from matching observations in another.
  - ▶ **Filtering joins**
    - Filter observations from one data frame based on whether or not they match an observation in the other table.

# Create a narrower dataset

```
flights2 <- flights %>%
  select(year:day, hour, origin, dest, tailnum, carrier)

flights2
#> # A tibble: 336,776 x 8
#>   year month   day hour origin dest tailnum carrier
#>   <int> <int> <int> <dbl> <chr>  <chr>  <chr>  <chr>
#> 1 2013     1     1     5   EWR    IAH    N14228  UA
#> 2 2013     1     1     5   LGA    IAH    N24211  UA
#> 3 2013     1     1     5   JFK    MIA    N619AA  AA
#> 4 2013     1     1     5   JFK    BQN    N804JB  B6
#> 5 2013     1     1     6   LGA    ATL    N668DN  DL
#> 6 2013     1     1     5   EWR    ORD    N39463  UA
#> 7 2013     1     1     6   EWR    FLL    N516JB  B6
#> 8 2013     1     1     6   LGA    IAD    N829AS  EV
#> 9 2013     1     1     6   JFK    MCO    N593JB  B6
#> 10 2013    1     1     6   LGA    ORD    N3ALAA  AA
#> # ... with 336,766 more rows
```

# A simple example


```
flights2 %>%
  select(-origin, -dest) %>%
  left_join(airlines, by = "carrier")
#> # A tibble: 336,776 x 7
#>   year month   day hour tailnum carrier name
#>   <int> <int> <int> <dbl> <chr>   <chr>
#> 1 2013     1     1     5 N14228  UA     United Air Lines Inc.
#> 2 2013     1     1     5 N24211  UA     United Air Lines Inc.
#> 3 2013     1     1     5 N619AA  AA     American Airlines Inc.
#> 4 2013     1     1     5 N804JB  B6     JetBlue Airways
#> 5 2013     1     1     6 N668DN  DL     Delta Air Lines Inc.
#> 6 2013     1     1     5 N39463  UA     United Air Lines Inc.
#> 7 2013     1     1     6 N516JB  B6     JetBlue Airways
#> 8 2013     1     1     6 N829AS  EV     ExpressJet Airlines Inc.
#> 9 2013     1     1     6 N593JB  B6     JetBlue Airways
#> 10 2013    1     1     6 N3ALAA  AA     American Airlines Inc.
#> # ... with 336,766 more rows
```

# Why mutating join?


```
flights2 %>%
  select(-origin, -dest) %>%
  mutate(name = airlines$name[match(carrier, airlines$carrier)])
#> # A tibble: 336,776 x 7
#>   year month   day hour tailnum carrier name
#>   <int> <int> <int> <dbl> <chr>   <chr>   <chr>
#> 1 2013     1     1     5 N14228  UA      United Air Lines Inc.
#> 2 2013     1     1     5 N24211  UA      United Air Lines Inc.
#> 3 2013     1     1     5 N619AA  AA      American Airlines Inc.
#> 4 2013     1     1     5 N804JB  B6      JetBlue Airways
#> 5 2013     1     1     6 N668DN  DL      Delta Air Lines Inc.
#> 6 2013     1     1     5 N39463  UA      United Air Lines Inc.
#> 7 2013     1     1     6 N516JB  B6      JetBlue Airways
#> 8 2013     1     1     6 N829AS  EV      ExpressJet Airlines Inc.
#> 9 2013     1     1     6 N593JB  B6      JetBlue Airways
#> 10 2013    1     1     6 N3ALAA  AA      American Airlines Inc.
#> # ... with 336,766 more rows
```

# Understanding mutating joins

```
x <- tribble(~key, ~val_x,
             1, "x1",
             2, "x2",
             3, "x3")
y <- tribble(~key, ~val_y,
             1, "y1",
             2, "y2",
             4, "y3")
```



# Inner join




```
x %>%
  inner_join(y, by = "key")
#> # A tibble: 2 x 3
#>   key  val_x val_y
#>   <dbl> <chr> <chr>
#> 1     1 x1    y1
#> 2     2 x2    y2
```


# Outer joins

- **Outer joins** keep observations that appear in at least one of the tables:
  - ▶ **Left join:** keeps all observations in x.
  - ▶ **Right join:** keeps all observations in y.
  - ▶ **Full join:** keeps all observations in x and y
- They work by adding to each table an additional “virtual” observation which
  - ▶ has a key that always matches (if no other key matches),
  - ▶ and a value filled with NA.


# Outer joins II




# A Venn diagram for joins




inner\_join(x, y)



full\_join(x, y)



left\_join(x, y)



right\_join(x, y)

- Two possibilities:
  - ▶ One table has duplicate keys.
    - Useful to add in additional information as there is typically a one-to-many relationship.
  - ▶ Both tables have duplicate keys.
    - Usually an error because in neither table do the keys uniquely identify an observation.
    - When you join duplicated keys, you get all possible combinations (i.e., the Cartesian product).

# One table has duplicate keys


- Only x has duplicated keys:

```
x <- tribble(~key, ~val_x,
  1, "x1",
  2, "x2",
  2, "x3",
  1, "x4")
```

```
y <- tribble(~key, ~val_y,
  1, "y1",
  2, "y2")
```

- The join adds val\_y to the matching rows:

```
left_join(x, y, by = "key")
#> # A tibble: 4 x 3
#>   key  val_x val_y
#>   <dbl> <chr> <chr>
#> 1     1 x1    y1
#> 2     2 x2    y2
#> 3     2 x3    y2
#> 4     1 x4    y1
```



	val_x	key	val_y
x1	1	y1	
x2	2	y2	
x3	2	y2	
x4	1	y1	

# Both tables have duplicate keys


- Both x and y have duplicated keys:

```
x <- tribble(~key, ~val_x,
             1, "x1",
             2, "x2",
             2, "x3",
             3, "x4")
```

```
y <- tribble(~key, ~val_y,
             1, "y1",
             2, "y2",
             2, "y3",
             3, "y4")
```

- The joint creates all combinations:

```
left_join(x, y, by = "key")
#> # A tibble: 6 x 3
#>   key  val_x val_y
#>   <dbl> <chr> <chr>
#> 1     1 x1    y1
#> 2     2 x2    y2
#> 3     2 x2    y3
#> 4     2 x3    y2
#> 5     2 x3    y3
#> 6     3 x4    y4
```



# Defining the key columns

- Default uses all variables that appear in both tables.
- Called a **natural join**.

```
flights2 %>%
  left_join(weather)
#> Joining, by = c("year", "month", "day", "hour", "origin")
#> # A tibble: 336,776 x 18
#>   year month   day hour origin dest tailnum carrier temp dewp
#>   <int> <int> <int> <dbl> <chr> <chr> <chr> <chr> <dbl> <dbl>
#> 1 2013     1     1     5 EWR   IAH N14228 UA    39.0 28.0
#> 2 2013     1     1     5 LGA   IAH N24211 UA    39.9 25.0
#> 3 2013     1     1     5 JFK   MIA N619AA AA    39.0 27.0
#> 4 2013     1     1     5 JFK   BQN N804JB B6    39.0 27.0
#> 5 2013     1     1     6 LGA   ATL N668DN DL    39.9 25.0
#> 6 2013     1     1     5 EWR   ORD N39463 UA    39.0 28.0
#> 7 2013     1     1     6 EWR   FLL N516JB B6    37.9 28.0
#> 8 2013     1     1     6 LGA   IAD N829AS EV    39.9 25.0
#> 9 2013     1     1     6 JFK   MCO N593JB B6    37.9 27.0
#> 10 2013    1     1     6 LGA   ORD N3ALAA AA    39.9 25.0
#> # ... with 336,766 more rows, and 8 more variables: humid <dbl>,
#> #   wind_dir <dbl>, wind_speed <dbl>, wind_gust <dbl>, precip <dbl>,
#> #   pressure <dbl>, visib <dbl>, time_hour <dttm>
```

# Using a character vector

- Like a natural join, but uses only some of the common variables:

```
flights2 %>%
  left_join(planes, by = "tailnum")
#> # A tibble: 336,776 x 16
#>   year.x month   day hour origin dest tailnum carrier year.y type
#>   <int> <int> <int> <dbl> <chr>  <chr>  <chr>  <chr>  <chr>  <int> <chr>
#> 1 2013     1     1     5 EWR    IAH    N14228  UA     1999 Fixe-
#> 2 2013     1     1     5 LGA    IAH    N24211  UA     1998 Fixe-
#> 3 2013     1     1     5 JFK    MIA    N619AA AA     1990 Fixe-
#> 4 2013     1     1     5 JFK    BQN    N804JB B6     2012 Fixe-
#> 5 2013     1     1     6 LGA    ATL    N668DN DL     1991 Fixe-
#> 6 2013     1     1     5 EWR    ORD    N39463  UA     2012 Fixe-
#> 7 2013     1     1     6 EWR    FLL    N516JB B6     2000 Fixe-
#> 8 2013     1     1     6 LGA    IAD    N829AS EV     1998 Fixe-
#> 9 2013     1     1     6 JFK    MCO    N593JB B6     2004 Fixe-
#> 10 2013    1     1     6 LGA   ORD    N3ALAA AA      NA <NA>
#> # ... with 336,766 more rows, and 6 more variables:
#> #   manufacturer <chr>, model <chr>, engines <int>, seats <int>,
#> #   speed <int>, engine <chr>
```

# Using a named character vector

- With by = c("a" = "b"), left\_join matches variable a in table x to variable b in table y:

```
flights2 %>%
  left_join(airports, c("dest" = "faa"))
#> # A tibble: 336,776 x 15
#>   year month   day hour origin dest tailnum carrier name      lat
#>   <int> <int> <int> <dbl> <chr>  <chr>  <chr>  <chr>  <chr>  <dbl>
#> 1 2013     1     1     5 EWR    IAH    N14228  UA     George~ 30.0
#> 2 2013     1     1     5 LGA    IAH    N24211  UA     George~ 30.0
#> 3 2013     1     1     5 JFK    MIA    N619AA  AA     Miami ~ 25.8
#> 4 2013     1     1     5 JFK    BQN    N804JB  B6     <NA>    NA
#> 5 2013     1     1     6 LGA    ATL    N668DN  DL     Hartsf~ 33.6
#> 6 2013     1     1     5 EWR    ORD    N39463  UA     Chicag~ 42.0
#> 7 2013     1     1     6 EWR    FLL    N516JB  B6     Fort L~ 26.1
#> 8 2013     1     1     6 LGA    IAD    N829AS  EV     Washin~ 38.9
#> 9 2013     1     1     6 JFK    MCO    N593JB  B6     Orland~ 28.4
#> 10 2013    1     1     6 LGA    ORD    N3ALAA  AA     Chicag~ 42.0
#> # ... with 336,766 more rows, and 5 more variables: lon <dbl>,
#> #   alt <dbl>, tz <dbl>, dst <chr>, tzone <chr>
```

# SQL is the inspiration

---

dplyr	SQL
inner_join(x, y, by = "z")	SELECT * FROM x INNER JOIN y USING (z)
left_join(x, y, by = "z")	SELECT * FROM x LEFT OUTER JOIN y USING (z)
right_join(x, y, by = "z")	SELECT * FROM x RIGHT OUTER JOIN y USING (z)
full_join(x, y, by = "z")	SELECT * FROM x FULL OUTER JOIN y USING (z)

---

- Note that:

- ▶ “INNER” and “OUTER” are optional, and often omitted.
- ▶ Joining different variables between the tables uses a slightly different syntax in SQL.
  - E.g. `inner_join(x, y, by = c("a" = "b"))` vs `SELECT * FROM x INNER JOIN y ON x.a = y.b.`

# Filtering joins

- Similar to mutating joins, but affect the observations rather than the variables:
  - ▶ `semi_join(x, y)` **keeps** all observations in `x` that have a match in `y`.
    - Useful for matching filtered summary tables back to the original rows.
  - ▶ `anti_join(x, y)` **drops** all observations in `x` that have a match in `y`.
    - Useful for diagnosing join mismatches.

# Flights that went to top destinations

```
top_dest <- flights %>%
  count(dest, sort = TRUE) %>%
  head(10)

flights %>%
  filter(dest %in% top_dest$dest) %>%
  print(n = 5)
#> # A tibble: 141,145 x 19
#>   year month   day dep_time sched_dep_time dep_delay arr_time
#>   <int> <int> <int>     <int>           <int>     <dbl>     <int>
#> 1 2013     1     1      542            540        2     923
#> 2 2013     1     1      554            600       -6     812
#> 3 2013     1     1      554            558       -4     740
#> 4 2013     1     1      555            600       -5     913
#> 5 2013     1     1      557            600       -3     838
#> # ... with 141,140 more rows, and 12 more variables:
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
#> #   time_hour <dttm>
```

- How to extend to multiple variables?


# Semi-join

- Only keeps rows in x having a match in y:


```
flights %>%  
  semi_join(top_dest)  
#> Joining, by = "dest"  
#> # A tibble: 141,145 x 19  
#>   year month   day dep_time sched_dep_time dep_delay arr_time  
#>   <int> <int> <int>     <int>           <int>     <dbl>     <int>  
#> 1 2013     1     1      542            540        2     923  
#> 2 2013     1     1      554            600       -6     812  
#> 3 2013     1     1      554            558       -4     740  
#> 4 2013     1     1      555            600       -5     913  
#> 5 2013     1     1      557            600       -3     838  
#> 6 2013     1     1      558            600       -2     753  
#> 7 2013     1     1      558            600       -2     924  
#> 8 2013     1     1      558            600       -2     923  
#> 9 2013     1     1      559            559        0     702  
#> 10 2013    1     1      600            600        0     851  
#> # ... with 141,135 more rows, and 12 more variables:  
#> #   sched_arr_time <int>, arr_delay <dbl>, carrier <chr>,  
#> #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
#> #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,  
#> #   time_hour <dttm>
```

# Visually understand the semi-join

- One-to-many:



- Many-to-many:



# flights without a match in planes

```
flights %>%  
  anti_join(planes,  
            by = "tailnum") %>%  
  count(tailnum, sort = TRUE)  
#> # A tibble: 722 x 2  
#>   tailnum      n  
#>   <chr>     <int>  
#> 1 N725MQ     2512  
#> 2 N722MQ      575  
#> 3 N723MQ      513  
#> 4 N723MQ      507  
#> 5 N713MQ      483  
#> 6 N735MQ      396  
#> 7 NOEGMQ      371  
#> 8 N534MQ      364  
#> 9 N542MQ      363  
#> 10 N531MQ     349  
#> # ... with 712 more rows
```

