

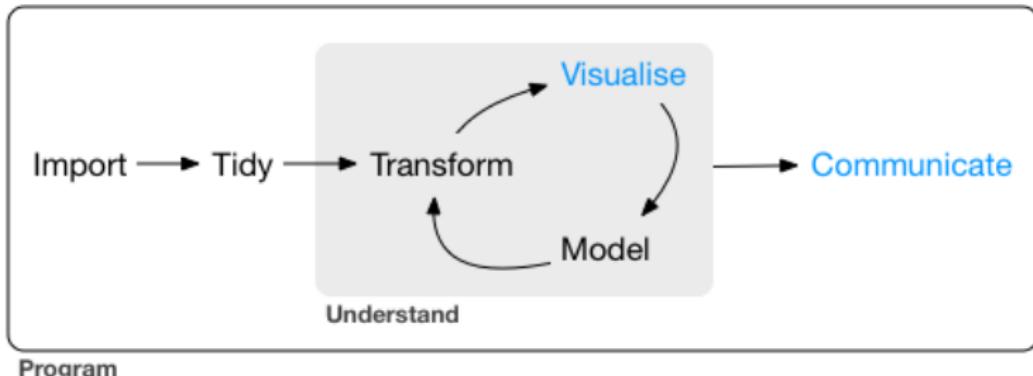
DSFBA: Visualization

Data Science for Business Analytics

Thibault Vatter

Department of Statistics, Columbia University

10/27/2021



Most of the material (e.g., the picture above) is borrowed from

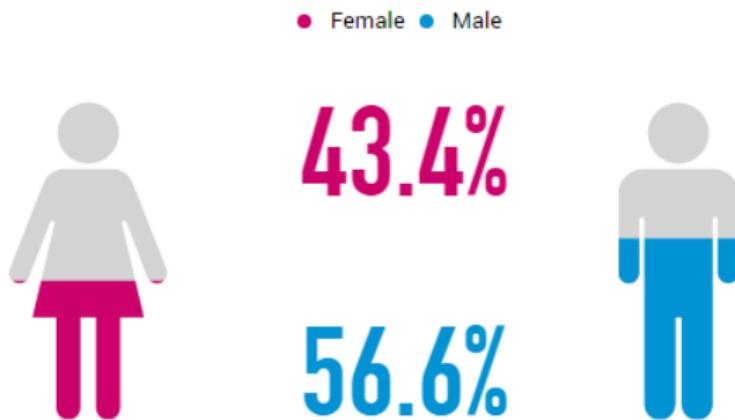
R for data science

- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems and the layered grammar of graphics

- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems and the layered grammar of graphics

Data content

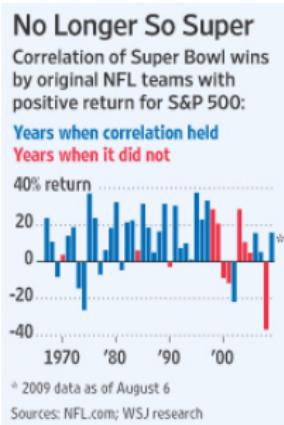
- Makes no sense to use graphs for very small amounts of data.
- The human brain is capable of grasping a few values.



source: talkwalker.com

Data relevance

- Graphs are only as good as the data they display.
- No creativity can produce a good graph from poor data.



- Leinweber (author of *Nerds on Wall Street*):
 - ▶ The S&P500 could be “predicted” at 75% by the butter production in Bangladesh.
 - ▶ ... Or 99% when adding cheese production in the USA, and the population of sheep.

Complexity

- Graphs shouldn't be more complex than the data they portray.
- Unnecessary complexity can be introduced by irrelevant
 - ▶ decoration
 - ▶ color
 - ▶ 3d effects
- ... Collectively known as "chartjunk" !

Distribution of All TFBS Regions

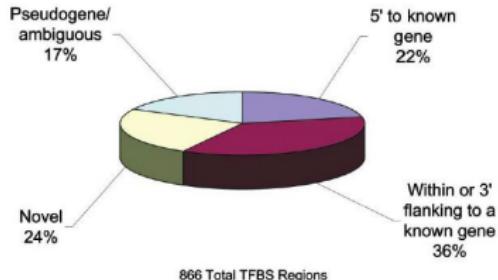
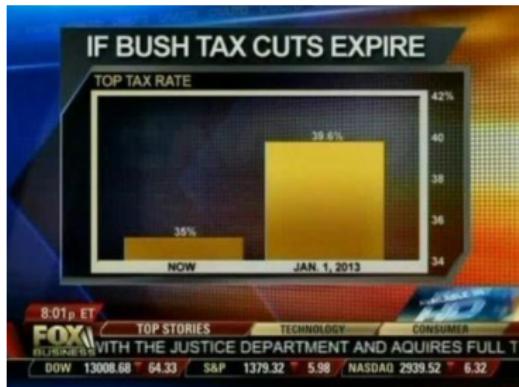


Figure 1. Classification of TFBS Regions
TFBS regions for Sp1, cMyc, and p53 were classified based upon proximity to annotations (RefSeq, Sanger hand-curated annotations, GenBank full-length mRNAs, and Ensembl predicted genes). The proximity was calculated from the center of each TFBS region. TFBS regions were classified as follows: within 5 kb of the 5' most exon of a gene, within 5 kb of the 3' terminal exon, or within a gene, novel or outside of any annotation, and pseudogene/ambiguous (TFBS overlapping or flanking pseudogene annotations, limited to chromosome 22, or TFBS regions falling into more than one of the above categories).

source: Cawley S, et al. (2004), Cell 116:499-509, Figure 1

Distortion

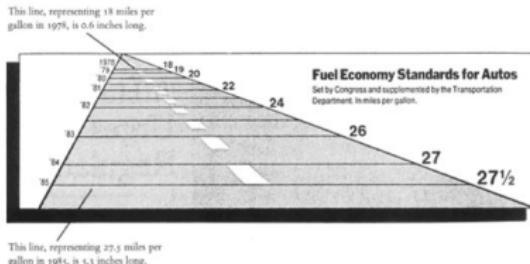
- Graphs shouldn't be distorted pictures of the portrayed values:
 - ▶ Can be either deliberate or accidental.
 - ▶ Useful to know how to produce truth bending graphs.
 - ▶ Misleading often used as a synonym of distorted.
 - ▶ See https://en.wikipedia.org/wiki/Misleading_graph



source: [statisticshowto.com/misleading-graphs/](http://statisticsshowto.com/misleading-graphs/)

More on distortion

- Common sources of distortion:
 - ▶ 3 dimensional “effects”.
 - ▶ linear scaling when using area or volume to represent values.
- The “lie factor”:
 - ▶ Measure of the amount of distortion in a graph.
 - lie factor = $\frac{\text{size of effect shown in graphic}}{\text{size of effect shown in data}}$
 - Don't take this too seriously, defined by Ed Tufte of Yale.
 - ▶ If > 1 , the graph is exaggerating the effect.



$$\text{lie factor} = \frac{\frac{5.3 - 0.6}{0.6}}{\frac{27.5 - 18}{18}} = 14.8!$$



Drawing good graphs

- The three main rules:
 - ▶ If the “story” is simple, keep it simple.
 - ▶ If the “story” is complex, make it look simple.
 - ▶ Tell the truth – do not distort the data.
- Specifically:
 - ▶ There should be a high data to chart ratio.
 - ▶ Use the appropriate graph for the appropriate purpose.
 - Most graphs presented in Excel are POOR CHOICES!
 - In particular, never use a pie chart!
 - ▶ Make sure that the graph is complete:
 - All axes must be labeled.
 - The units should be indicated.
 - There should be a title.
 - A legend can provide needed additional information (e.g., for colors or line types).

A grammar of graphics

"A grammar of graphics is a tool that enables us to concisely describe the components of a graphic. Such a grammar allows us to move beyond named graphics (e.g., the "scatterplot") and gain insight into the deep structure that underlies statistical graphics." — Hadley Wickham

- ggplot2 is an R implementation of the concept:
 - ▶ A coherent system for describing and creating graphs.
 - ▶ Based on [The Grammar of Graphics](#).
 - ▶ Learn one system and apply it in many places.
 - ▶ The equivalent of dplyr for graphs.
- To learn more, read [The Layered Grammar of Graphics](#).
- Implementations exist in other languages (e.g., Python)

The mpg data frame

- Data from the US EPA on 38 models of car:

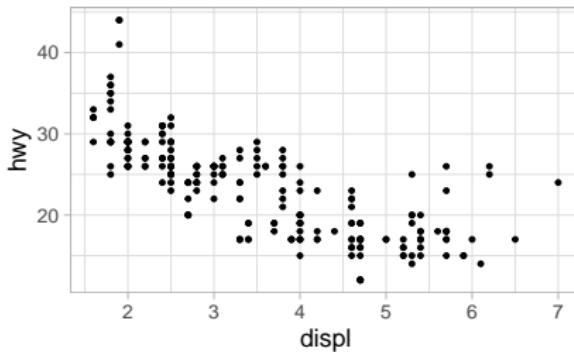
```
mpg %>% print(n = 5)
#> # A tibble: 234 x 11
#>   manufacturer model displ year cyl trans drv cty hwy fl
#>   <chr>        <chr>  <dbl> <int> <int> <chr> <chr> <int> <int> <chr>
#> 1 audi         a4      1.8  1999    4 auto(~ f       18    29 p
#> 2 audi         a4      1.8  1999    4 manua~ f       21    29 p
#> 3 audi         a4      2     2008    4 manua~ f       20    31 p
#> 4 audi         a4      2     2008    4 auto(~ f       21    30 p
#> 5 audi         a4      2.8  1999    6 auto(~ f       16    26 p
#> # ... with 229 more rows, and 1 more variable: class <chr>
```

- Among the variables in `mpg` are:
 - ▶ `displ`, a car's engine size, in litres.
 - ▶ `hwy`, a car's fuel efficiency on the highway (in miles per gallon).
- A few questions
 - ▶ Do cars with big engines use more fuel ?
 - ▶ What does the relationship between engine size and fuel efficiency look like? Positive? Negative? Linear? Nonlinear?

Creating a plot

- A simple scatterplot:

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy))
```



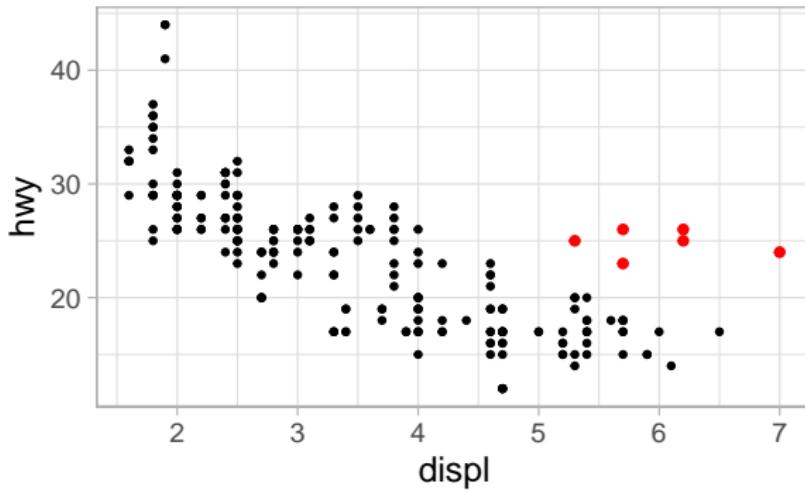
- A graphing template:

```
ggplot(data = <DATA>) +  
  <GEOM_FUNCTION>(mapping = aes(<MAPPINGS>))
```

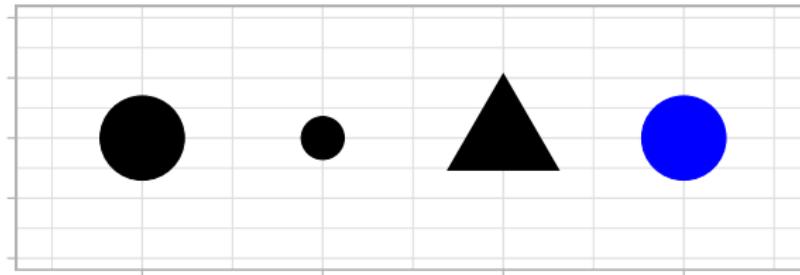
- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems and the layered grammar of graphics

Aesthetic mappings

"The greatest value of a picture is when it forces us to notice what we never expected to see." — John Tukey

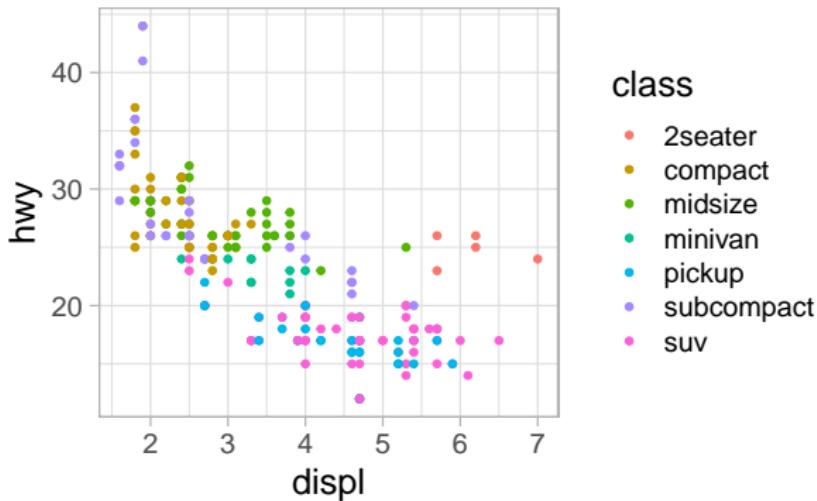


- How to add more variables to a two dimensional plot?
- By mapping them to an **aesthetic**:
 - ▶ A visual property of the objects in your plot.
 - ▶ Include the size, the shape, or the color of the points.
- We use the words
 - ▶ “**value**” to describe data,
 - ▶ and “**level**” to describe aesthetic properties.



Adding classes to your plot

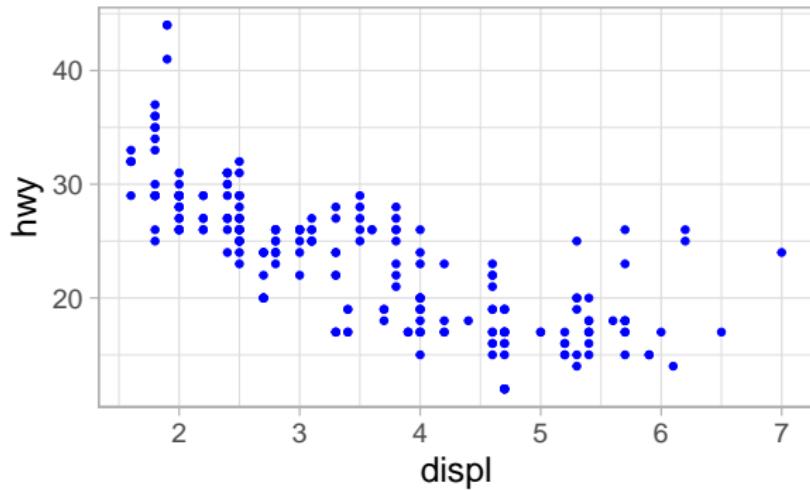
```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, color = class))
```



- If you prefer British English, use `color` instead of `color`.

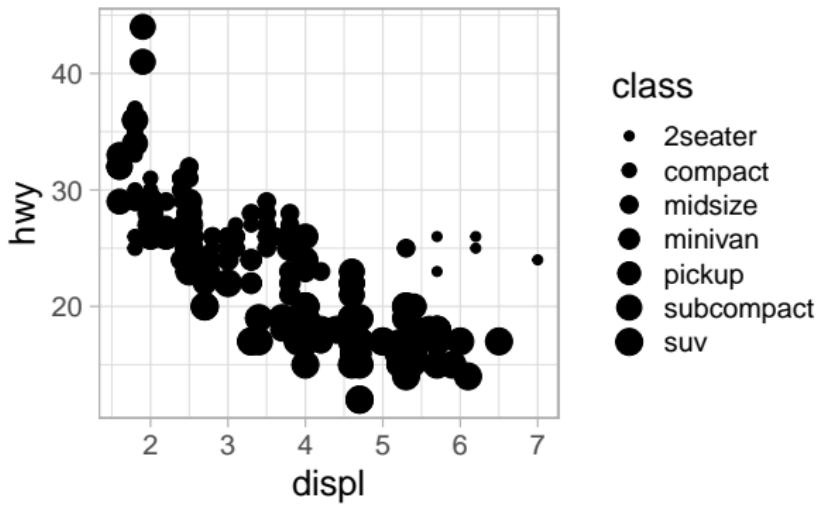
Set the aesthetics manually

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy), color = "blue")
```



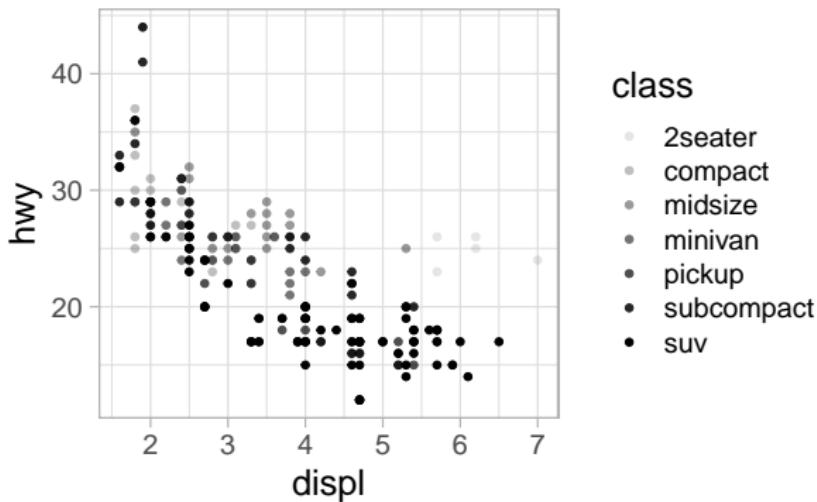
The size aesthetic

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, size = class))  
#> Warning: Using size for a discrete variable is not advised.
```



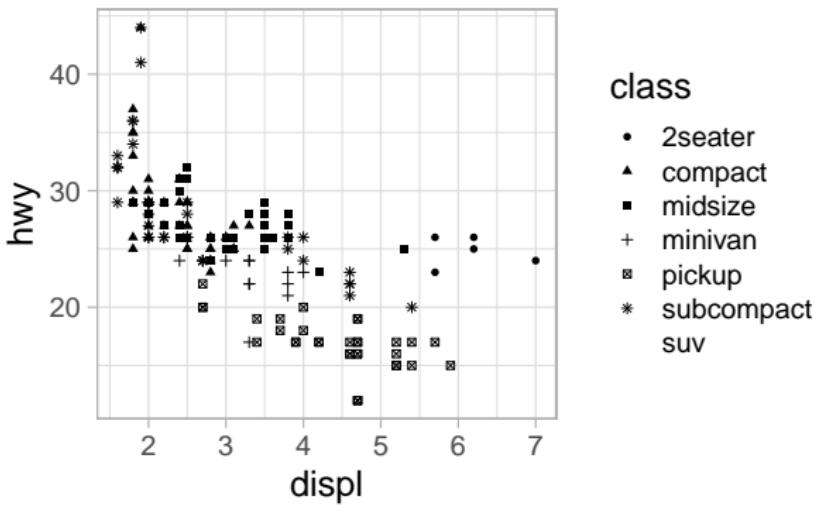
The alpha aesthetic

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, alpha = class))  
#> Warning: Using alpha for a discrete variable is not advised.
```



The shape aesthetic

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy, shape = class))  
#> Warning: The shape palette can deal with a maximum of 6 discrete  
#> values because more than 6 becomes difficult to discriminate;  
#> you have 7. Consider specifying shapes manually if you must  
#> have them.  
#> Warning: Removed 62 rows containing missing values (geom_point).
```



A remark on aesthetic values

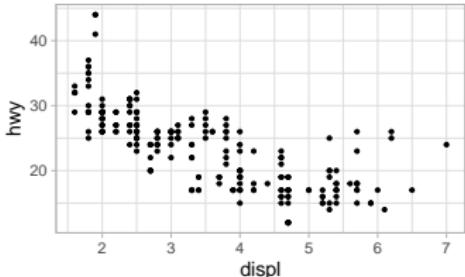
- Need values that make sense for that aesthetic:
 - ▶ The name of a color as a character string.
 - ▶ The size of a point in mm.
 - ▶ The shape of a point as a number.

□ 0	\times 4	\oplus 10	■ 15	■ 22
○ 1	∇ 6	\bigcirclearrowleft 11	● 16	● 21
△ 2	\boxtimes 7	\boxplus 12	▲ 17	▲ 24
◇ 5	$*$ 8	\otimes 13	◆ 18	◆ 23
+ 3	\diamondsuit 9	\boxminus 14	● 19	● 20

Figure 1: The hollow shapes (0–14) have a border determined by ‘color’; the solid shapes (15–18) are filled with ‘color’; the filled shapes (21–24) have a border of ‘color’ and are filled with ‘fill’.

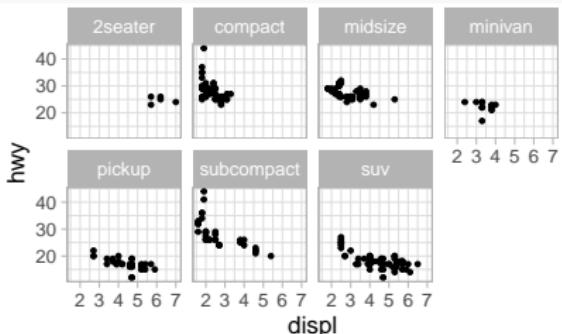
Facetting

```
(p <- ggplot(data = mpg) + geom_point(mapping = aes(x = displ, y = hwy)))
```



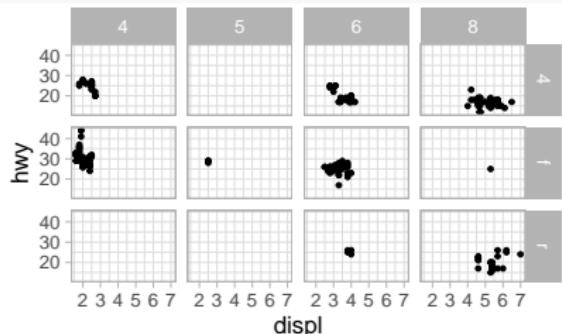
■ Facet wrap

```
p + facet_wrap(~ class, nrow = 2)
```



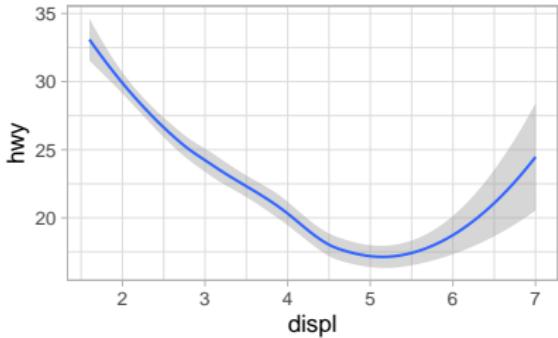
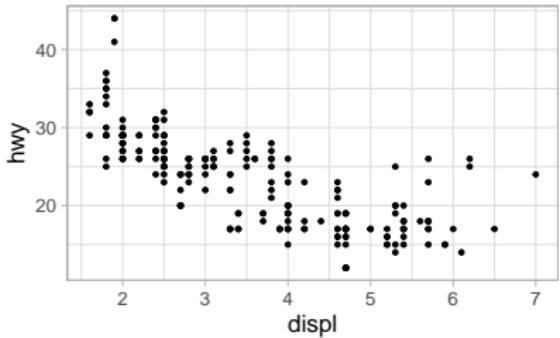
■ Facet grid

```
p + facet_grid(drv ~ cyl)
```



- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems and the layered grammar of graphics

How are these two plots similar?



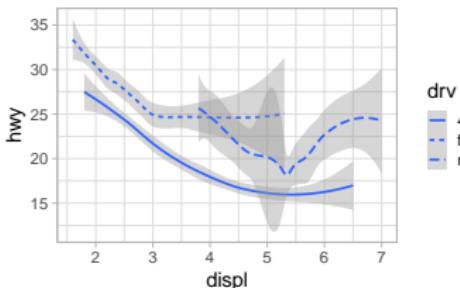
■ A geom:

- ▶ The object that a plot uses to represent data.
- ▶ Plots often described by the geom type:
 - Bar charts use bar geoms.
 - Line charts use line geoms.
 - Boxplots use boxplot geoms.
- ▶ An exception:
 - Scatterplots use the point geom.

Geometric objects

- Every **geom** function takes a `mapping` argument.
- But **not every aesthetic works with every geom**:
 - ▶ `shape` exists for `geom_point` but not for `geom_line`,
 - ▶ and conversely for `linetype`.

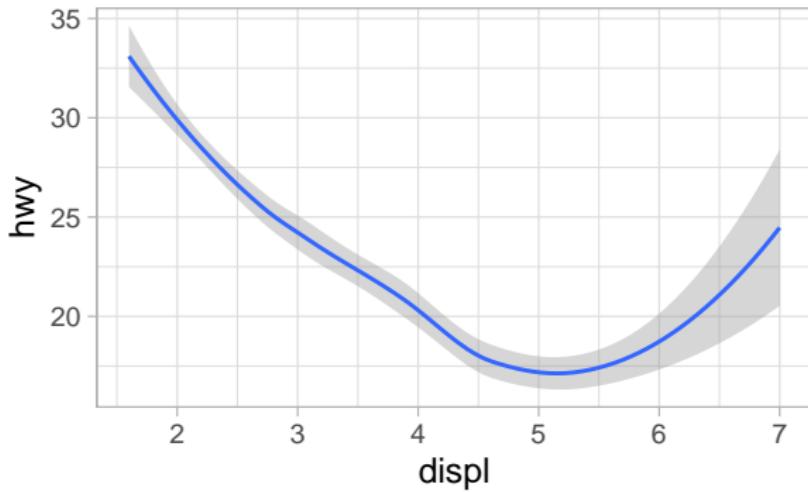
```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy, linetype = drv))
```



- Additionally
 - ▶ `ggplot2` provides over 30 geoms.
 - ▶ `extension packages` provide even more.
 - ▶ Use RStudio's `data visualization cheatsheet`.
 - ▶ To learn more about any single geom, use help: `?geom_smooth`.

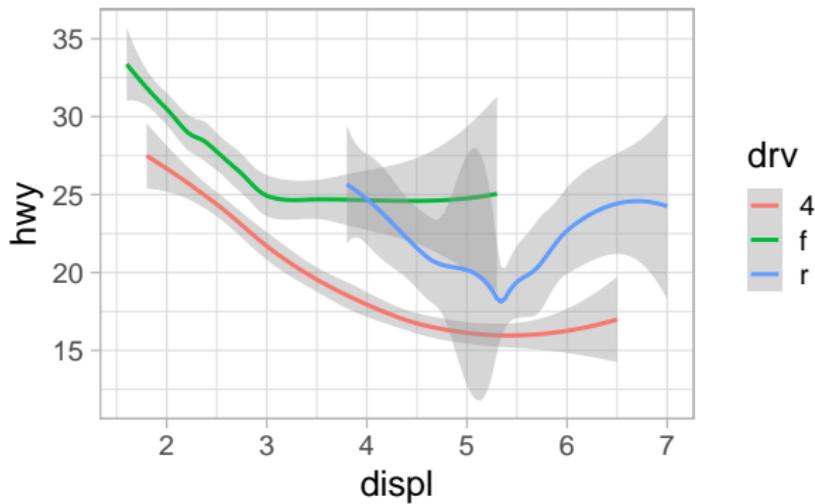
Geoms and legends

```
ggplot(data = mpg) + geom_smooth(mapping = aes(x = displ, y = hwy))
```



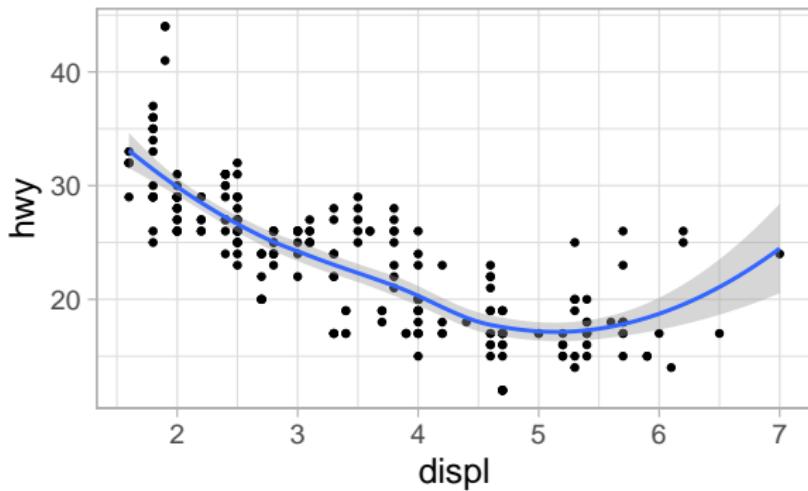
Geoms and legends

```
ggplot(data = mpg) +  
  geom_smooth(mapping = aes(x = displ, y = hwy, color = drv))
```



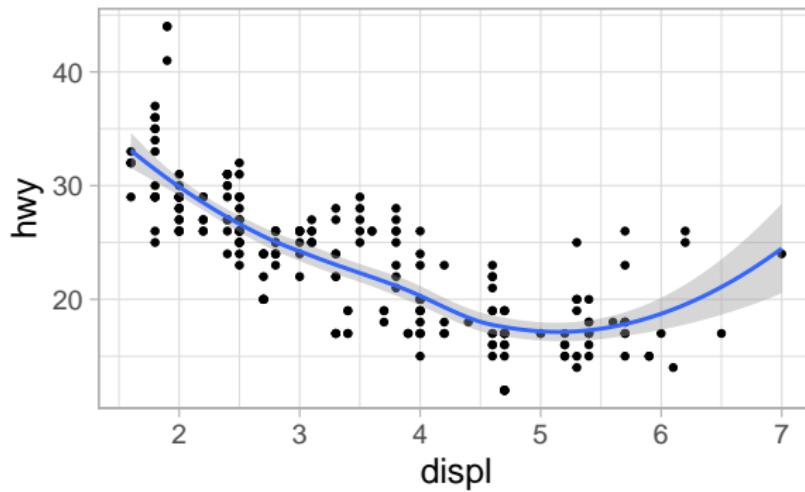
Multiple geoms in the same plot

```
ggplot(data = mpg) +  
  geom_point(mapping = aes(x = displ, y = hwy)) +  
  geom_smooth(mapping = aes(x = displ, y = hwy))
```



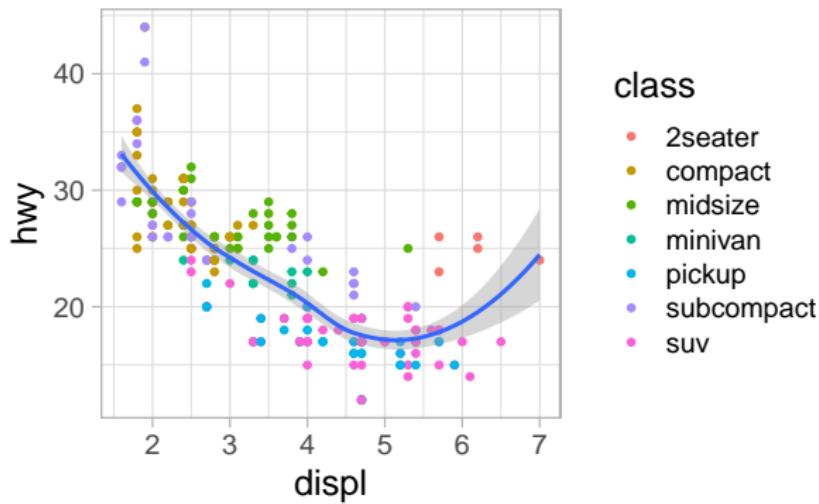
A better way

```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_point() +  
  geom_smooth()
```



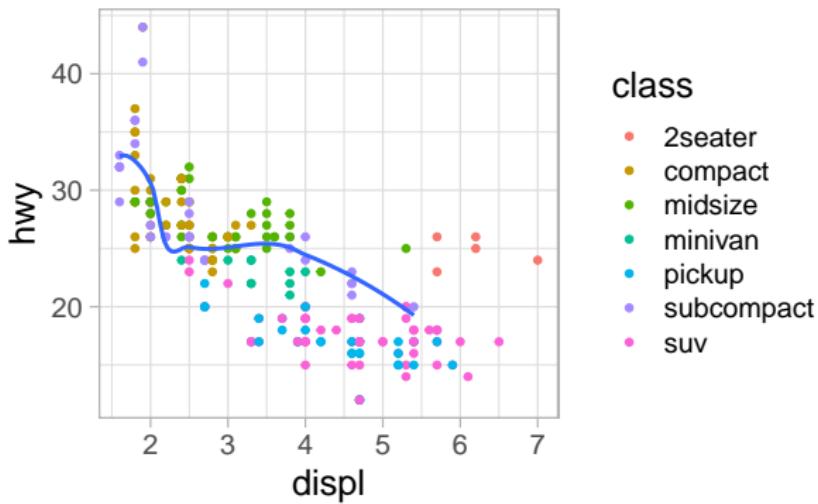
Local vs global mappings

```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_point(mapping = aes(color = class)) +  
  geom_smooth()
```



Layer dependent data

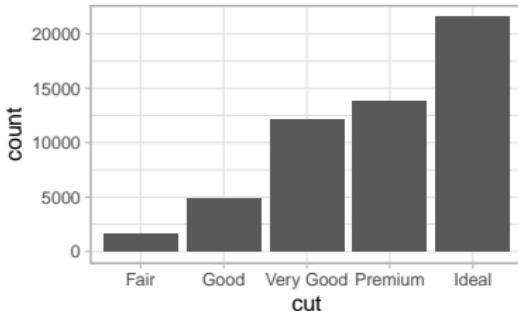
```
ggplot(data = mpg, mapping = aes(x = displ, y = hwy)) +  
  geom_point(mapping = aes(color = class)) +  
  geom_smooth(data = filter(mpg, class == "subcompact"), se = FALSE)
```



Beyond scatterplots

- Other graphs, like bar charts, calculate new values to plot.
 - ▶ Bar charts, histograms, and frequency polygons:
 - Bin data.
 - Plot bin counts (number of points falling in each bin).
 - ▶ Smoothers:
 - Fit a model to your data.
 - Plot predictions from the model.
 - ▶ Boxplots:
 - Compute a robust summary of the distribution.
 - Display a specially formatted box.

```
ggplot(data = diamonds) + geom_bar(mapping = aes(x = cut))
```



Statistical transformations

■ A stat:

- ▶ The algorithm used to calculate new values for a graph.
- ▶ Short for statistical transformation.

1. `geom_bar()` begins with the `diamonds` data set

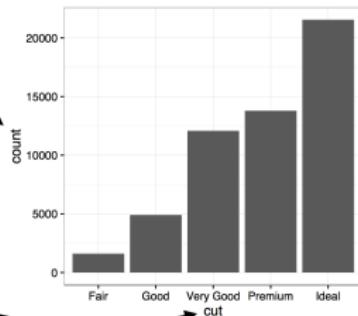
carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.98	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...

2. `geom_bar()` transforms the data with the "count" stat, which returns a data set of cut values and counts.

cut	count	prop
Fair	1610	1
Good	4906	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1

`stat_count()`

3. `geom_bar()` uses the transformed data to build the plot. cut is mapped to the x axis, count is mapped to the y axis.



- `ggplot2` provides over 20 stats.
- Each stat is a function, get help as usual, e.g. `?stat_bin`.
- Use RStudio's data visualization cheatsheet for a complete list.

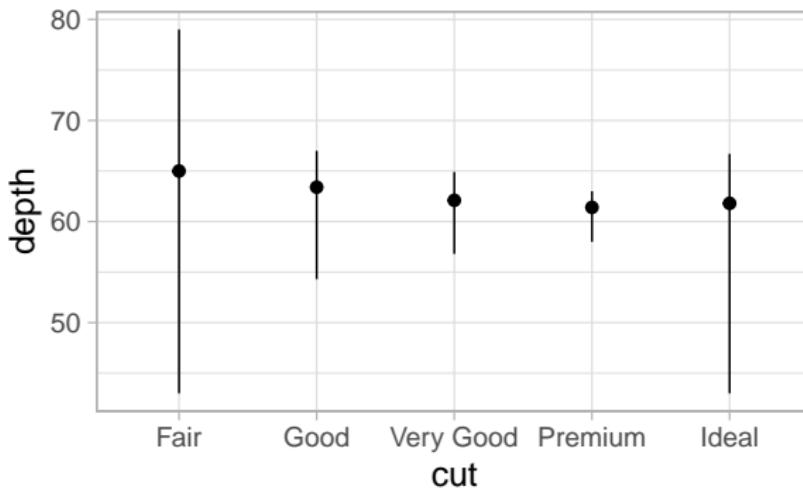
- Every geom has a default stat and conversely.
 - ▶ ?geom_bar shows that the default value for stat is “count”.
 - ▶ Means that geom_bar() uses stat_count().
 - ▶ ?stat_count has a section called “Computed variables” with two new variables: count and prop.
- You can generally use geoms and stats interchangeably!

```
ggplot(data = diamonds) +  
  stat_count(mapping = aes(x = cut))
```

- Typically, use geoms without worrying about the stat.
- Three reasons to use a stat explicitly:
 - ▶ To override the default stat.
 - ▶ To override the default mapping from transformed variables to aesthetics.
 - ▶ To draw greater attention to the stat in your code.

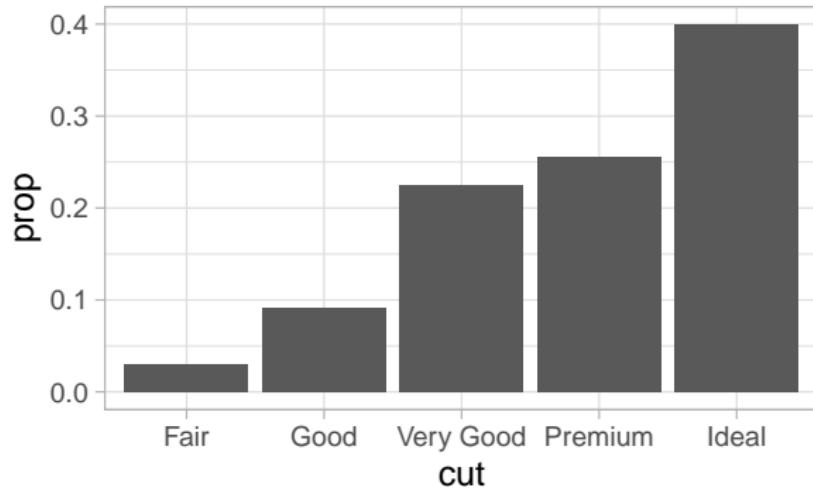
Use a stat explicitly !

```
ggplot(data = diamonds) +  
  stat_summary(  
    mapping = aes(x = cut, y = depth),  
    fun.min = min,  
    fun.max = max,  
    fun = median  
  )
```



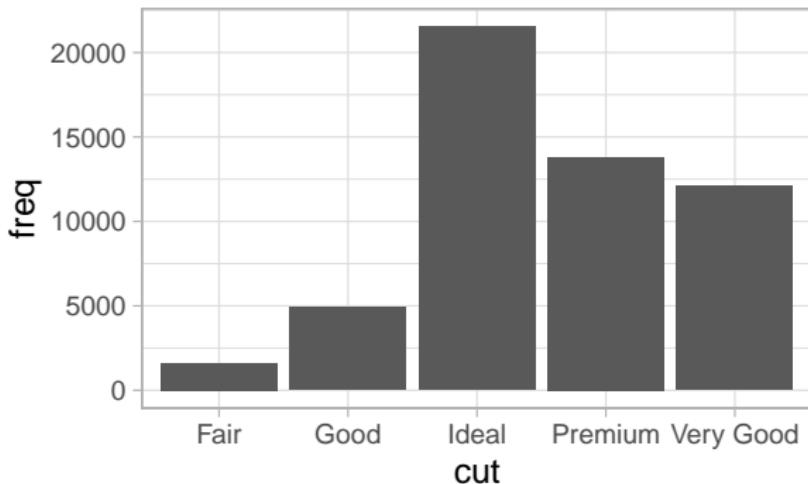
Use a stat explicitly II

```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, y = ..prop.., group = 1))
```



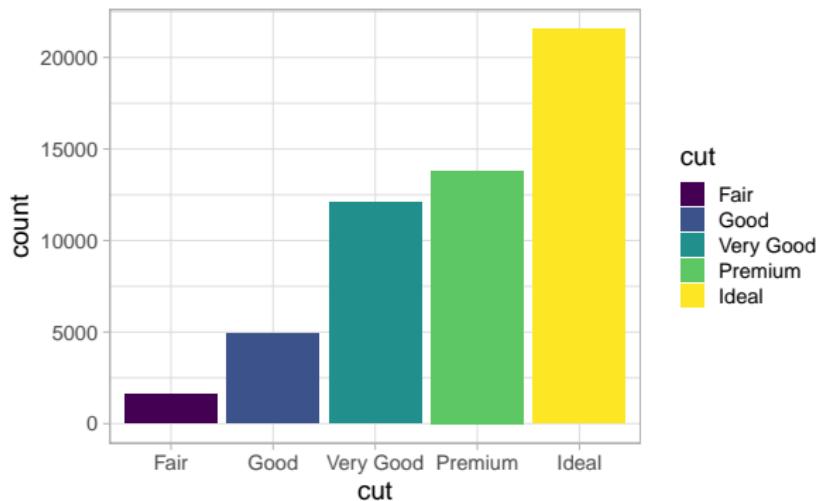
Use a stat explicitly III

```
demo <- tribble(~cut,          ~freq,
                 "Fair",        1610,
                 "Good",        4906,
                 "Very Good",  12082,
                 "Premium",    13791,
                 "Ideal",       21551)
ggplot(data = demo) +
  geom_bar(mapping = aes(x = cut, y = freq), stat = "identity")
```



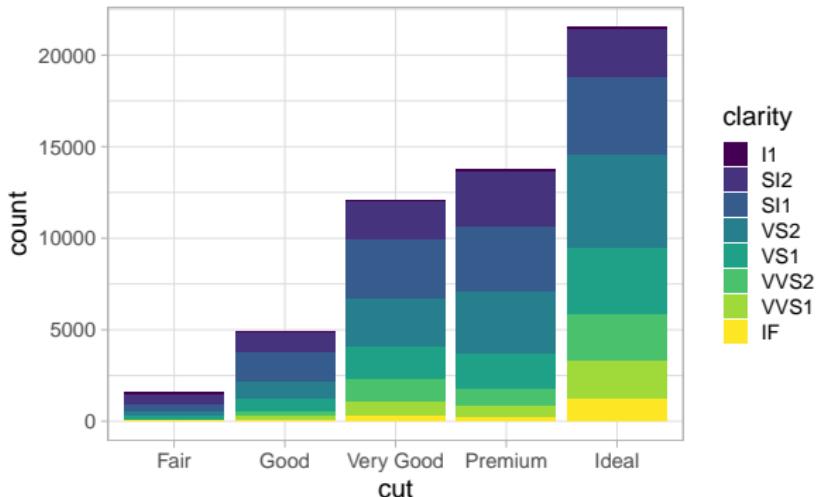
The fill aesthetic

```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = cut))
```



Fill and position adjustments

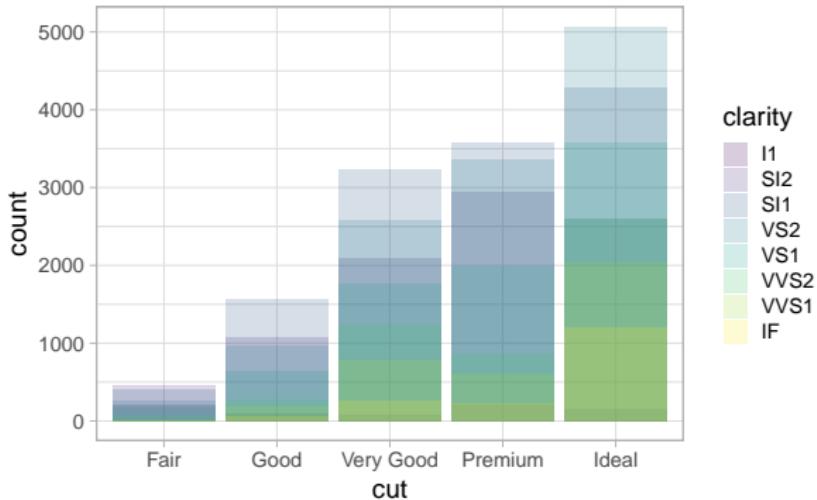
```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = clarity))
```



- Automatically stacked by the **position adjustment**.
- ?position_stack to learn more.

Fill with position = "identity"

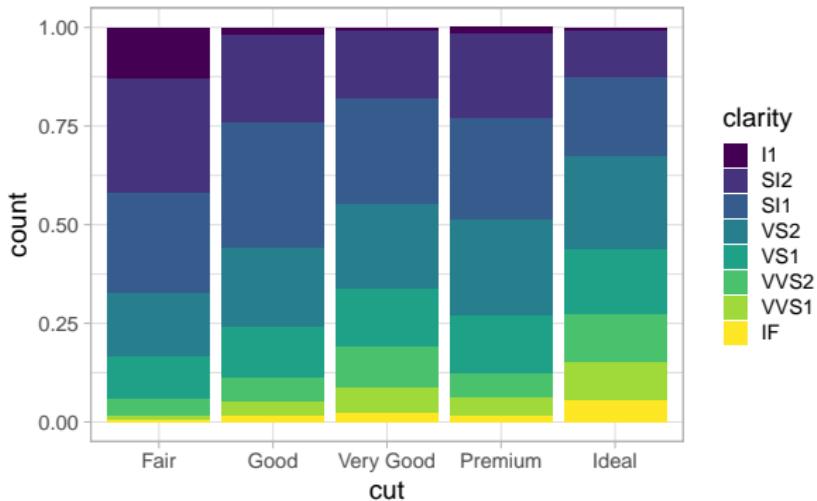
```
ggplot(data = diamonds, mapping = aes(x = cut, fill = clarity)) +  
  geom_bar(alpha = 1/5, position = "identity")
```



- Not very useful for bars because of overlap.
- ?position_identity to learn more.

Fill with position = "fill"

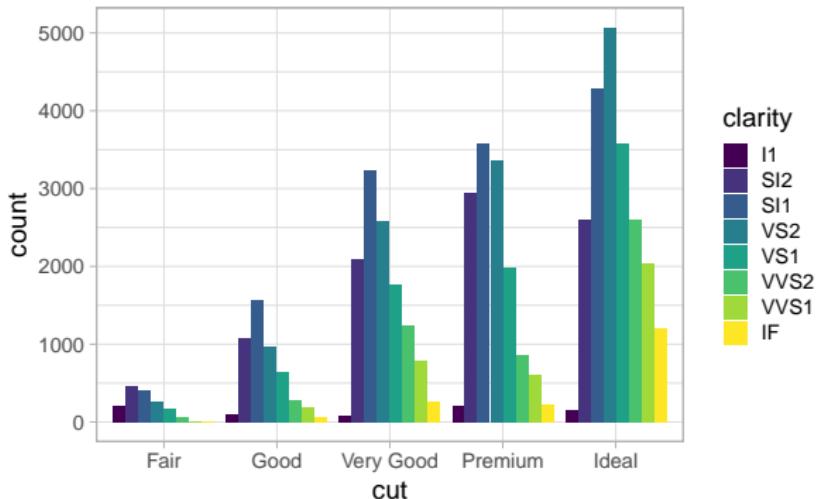
```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = clarity), position = "fill")
```



- Makes it easier to compare proportions across groups.
- ?position_fill to learn more.

Fill with position = "dodge"

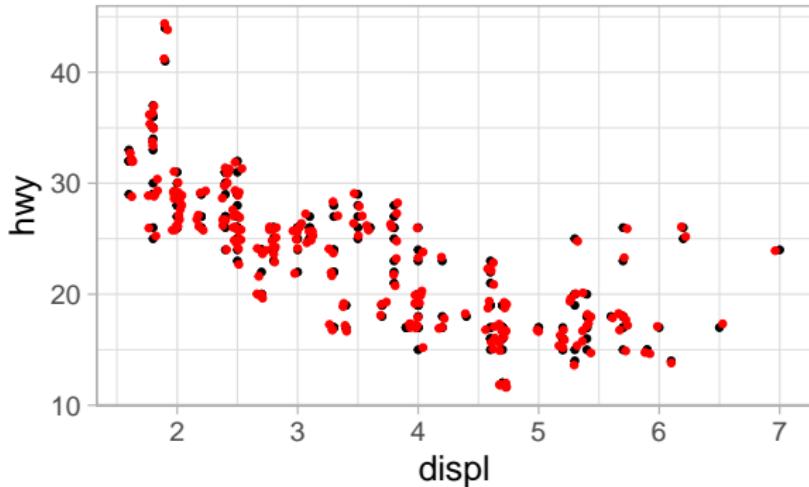
```
ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = clarity), position = "dodge")
```



- Makes it easier to compare individual values.
- ?position_dodge to learn more.

position = "jitter"

```
ggplot(data = mpg, aes(x = displ, y = hwy)) +  
  geom_point() +  
  geom_point(position = "jitter", color = "red")
```



- Graph less/**more** accurate/**revealing** at small/**large** scales.
- ?position_jitter to learn more.

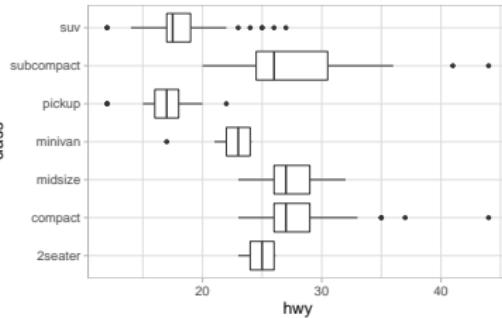
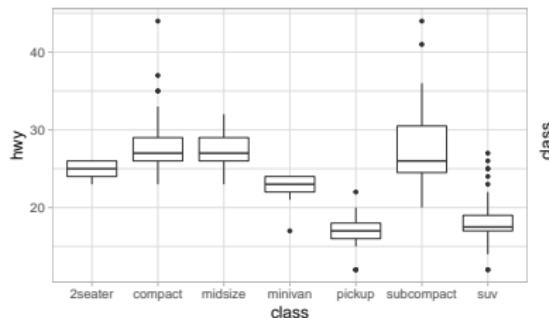
- 1 From bad graphs to the grammar of graphics
- 2 Aesthetics and facetting
- 3 Geometric objects and statistical transformations
- 4 Coordinate systems and the layered grammar of graphics

Coordinate systems

- The most complicated part of `ggplot2`.
- Default: the Cartesian coordinate system.
- Other systems occasionally helpful:
 - ▶ `coord_flip()` switches the x and y axes.
 - ▶ `coord_quickmap()` sets the aspect ratio correctly for maps.
 - ▶ `coord_polar()` uses polar coordinates.

coord_flip()

```
ggplot(data = mpg, mapping = aes(x = class, y = hwy)) +
  geom_boxplot()
ggplot(data = mpg, mapping = aes(x = class, y = hwy)) +
  geom_boxplot() +
  coord_flip()
```

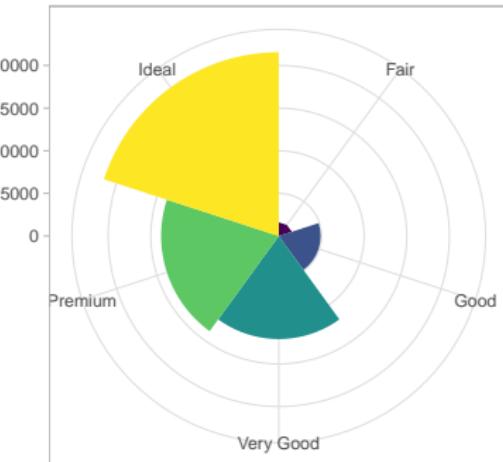
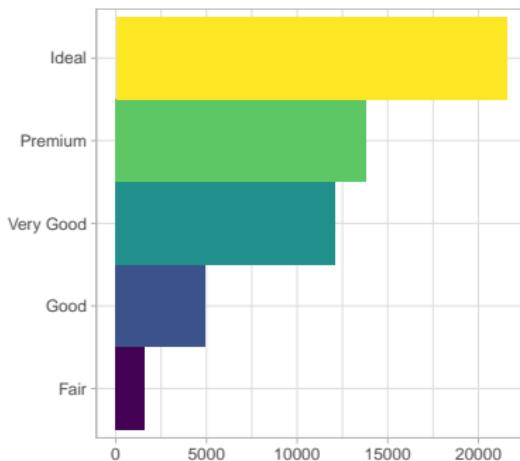


- Useful for:

- ▶ horizontal boxplots,
- ▶ and long labels.

coord_polar()

```
bar <- ggplot(data = diamonds) +  
  geom_bar(mapping = aes(x = cut, fill = cut),  
            show.legend = FALSE, width = 1) +  
  theme(aspect.ratio = 1) + labs(x = NULL, y = NULL)  
bar + coord_flip()  
bar + coord_polar()
```



The layered grammar of graphics

```
ggplot(data = <DATA>) +  
<GEOM_FUNCTION>(  
  mapping = aes(<MAPPINGS>),  
  stat = <STAT>,  
  position = <POSITION>  
) +  
<COORDINATE_FUNCTION> +  
<FACET_FUNCTION>
```

- A formal system for building plots,
- Uniquely describes *any* plot as a combination of
 - ▶ a dataset,
 - ▶ a geom,
 - ▶ a set of mappings,
 - ▶ a stat,
 - ▶ a position adjustment,
 - ▶ a coordinate system,
 - ▶ and a faceting scheme.

Example

1. Begin with the **iamonds** data set

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...

stat_count()

2. Compute counts for each cut value with **stat_count()**.

cut	count	prop
Fair	1610	1
Good	4906	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1

Example

3. Represent each observation with a bar.
4. Map the **fill** of each bar to the `..count..` variable.

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...

stat_count()



cut	count	prop
Fair	1610	1
Good	4906	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1

Example

carat	cut	color	clarity	depth	table	price	x	y	z
0.23	Ideal	E	SI2	61.5	55	326	3.95	3.98	2.43
0.21	Premium	E	SI1	59.8	61	326	3.89	3.84	2.31
0.23	Good	E	VS1	56.9	65	327	4.05	4.07	2.31
0.29	Premium	I	VS2	62.4	58	334	4.20	4.23	2.63
0.31	Good	J	SI2	63.3	58	335	4.34	4.35	2.75
...

stat_count()

cut	count	prop
Fair	1610	1
Good	4908	1
Very Good	12082	1
Premium	13791	1
Ideal	21551	1

5. Place geom in a cartesian coordinate system.

6. Map the y values to ..count.. and the x values to cut.

