

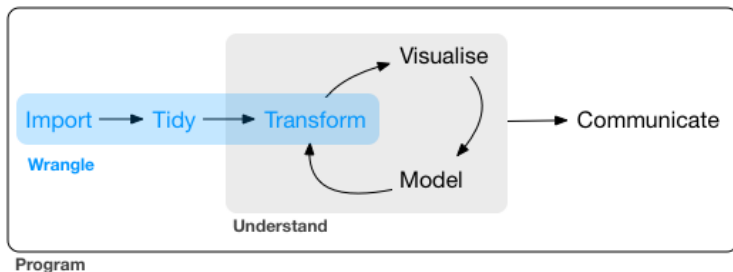
DSFBA: Wrangling

Data Science for Business Analytics

Thibault Vatter

Department of Statistics, Columbia University

11/10/2021



Most of the material (e.g., the picture above) is borrowed from

R for data science

1 Dates and times

2 Factors

3 Strings

1 Dates and times

2 Factors

3 Strings

- Does every year have 365 days?
- Does every day have 24 hours?
- Does every minute have 60 seconds?

- Two types of date/time data:
 - ▶ A **date**.
 - Tibbles print this as `<date>`.
 - ▶ A **date-time** is a date plus a time.
 - Uniquely identifies an instant in time (typically to the nearest second).
 - Tibbles print this as `<dtm>`.
 - Elsewhere in R, `POSIXct`.
- Use the **simplest possible data type satisfying your needs!**

- The **lubridate** package:
 - ▶ Makes it easier to work with dates and times in R.
 - ▶ Not part of core tidyverse because only needed when working with dates/times.

```
library(lubridate)
today()
#> [1] "2021-11-10"
now()
#> [1] "2021-11-10 09:07:40 EST"
```

- Other (usual) ways to create a date/time:
 - ▶ From a string.
 - ▶ From individual date-time components.
 - ▶ From an existing date/time object.

```
as_datetime(today())
#> [1] "2021-11-10 UTC"
as_date(now())
#> [1] "2021-11-10"
```

```
ymd("2017-01-31")
#> [1] "2017-01-31"
mdy("January 31st, 2017")
#> [1] "2017-01-31"
dmy("31-Jan-2017")
#> [1] "2017-01-31"

ymd_hms("2017-01-31 20:11:59")
#> [1] "2017-01-31 20:11:59 UTC"
mdy_hm("01/31/2017 08:01")
#> [1] "2017-01-31 08:01:00 UTC"
```

■ Additionally:

```
ymd(20170131)
#> [1] "2017-01-31"
ymd(20170131, tz = "UTC")
#> [1] "2017-01-31 UTC"
```


From individual components

```
flights %>%
  select(year:day, hour, minute, dep_time) %>%
  mutate(departure = make_datetime(year, month, day, hour, minute))
#> # A tibble: 336,776 x 7
#>   year month   day hour minute dep_time departure
#>   <int> <int> <int> <dbl> <dbl>   <int> <dtm>
#> 1  2013     1     1     5     15     517 2013-01-01 05:15:00
#> 2  2013     1     1     5     29     533 2013-01-01 05:29:00
#> 3  2013     1     1     5     40     542 2013-01-01 05:40:00
#> 4  2013     1     1     5     45     544 2013-01-01 05:45:00
#> 5  2013     1     1     6     0     554 2013-01-01 06:00:00
#> 6  2013     1     1     5     58     554 2013-01-01 05:58:00
#> 7  2013     1     1     6     0     555 2013-01-01 06:00:00
#> 8  2013     1     1     6     0     557 2013-01-01 06:00:00
#> 9  2013     1     1     6     0     557 2013-01-01 06:00:00
#> 10 2013     1     1     6     0     558 2013-01-01 06:00:00
#> # ... with 336,766 more rows
```


- For `dep_time` and others such as `arr_time`:

```
flights_dt <- flights %>%
  mutate(dep_time = make_datetime(
    year, month, day, dep_time %/% 100, dep_time %% 100))
```

■ Rounding:

- ▶ `floor_date()` rounds down.
- ▶ `round_date()` rounds to.
- ▶ `ceiling_date()` rounds up.

```
flights_dt %>%  
  filter(!is.na(dep_time)) %>%  
  count(week = floor_date(dep_time, "week")) %>%  
  ggplot(aes(week, n)) +  
  geom_line()
```



■ Getting the components:

```
datetime <- ymd_hms("2016-07-08 12:34:56")  
map_dbl(list(year, month, mday, yday, wday), function(f) f(datetime))  
#> [1] 2016    7    8   190    6
```

■ Setting the components:

```
year(datetime) <- 2020  
datetime  
#> [1] "2020-07-08 12:34:56 UTC"  
month(datetime) <- 01  
datetime  
#> [1] "2020-01-08 12:34:56 UTC"  
hour(datetime) <- hour(datetime) + 1  
datetime  
#> [1] "2020-01-08 13:34:56 UTC"
```

■ Alternatively:

```
update(datetime, year = 2019)  
#> [1] "2019-01-08 13:34:56 UTC"
```

- Goal: to do arithmetic (i.e., subtraction, addition, and division) with dates/times.
- Three classes that represent time spans:
 - ▶ **Durations** (number of seconds).
 - ▶ **Periods** (human units like weeks and months).
 - ▶ **Intervals** (a starting and ending point).

- A **duration** always record a time span in seconds.
- Larger units created at the standard rate.
 - ▶ E.g., 60s/mn, 60mn/h, 24h/d, 7d/w, 365d/y.

```
dseconds(15)
#> [1] "15s"
dminutes(10)
#> [1] "600s (~10 minutes)"
dhours(c(12, 24))
#> [1] "43200s (~12 hours)" "86400s (~1 days)"
ddays(0:5)
#> [1] "0s" "86400s (~1 days)" "172800s (~2 days)"
#> [4] "259200s (~3 days)" "345600s (~4 days)" "432000s (~5 days)"
dweeks(3)
#> [1] "1814400s (~3 weeks)"
dyears(1)
#> [1] "31557600s (~1 years)"
```

■ Add and multiply durations:

```
2 * dyears(1)
#> [1] "63115200s (~2 years)"
dyears(1) + dweeks(12) + dhours(15)
#> [1] "38869200s (~1.23 years)"
```

■ Add and subtract durations to and from dates/datetimes:

```
tomorrow <- today() + ddays(1)
last_year <- today() - dyears(1)
```

■ What happens here?

```
one_pm <- ymd_hms("2016-03-12 13:00:00", tz = "America/New_York")
one_pm
#> [1] "2016-03-12 13:00:00 EST"
one_pm + ddays(1)
#> [1] "2016-03-13 14:00:00 EDT"
```

- Work with “human” times, like days (no fixed length in secs):

```
one_pm
#> [1] "2016-03-12 13:00:00 EST"
one_pm + days(1)
#> [1] "2016-03-13 13:00:00 EDT"
seconds(15)
#> [1] "15S"
minutes(10)
#> [1] "10M 0S"
hours(c(12, 24))
#> [1] "12H 0M 0S" "24H 0M 0S"
days(7)
#> [1] "7d 0H 0M 0S"
months(1:3)
#> [1] "1m 0d 0H 0M 0S" "2m 0d 0H 0M 0S" "3m 0d 0H 0M 0S"
weeks(3)
#> [1] "21d 0H 0M 0S"
years(1)
#> [1] "1y 0m 0d 0H 0M 0S"
```

■ Add and multiply periods:

```
10 * (months(6) + days(1))  
#> [1] "60m 10d 0H 0M 0S"  
days(50) + hours(25) + minutes(2)  
#> [1] "50d 25H 2M 0S"
```

■ Add periods to dates/datetimes:

```
# A leap year  
ymd("2016-01-01") + dyears(1)  
#> [1] "2016-12-31 06:00:00 UTC"  
ymd("2016-01-01") + years(1)  
#> [1] "2017-01-01"  
  
# Daylight Savings Time  
one_pm + ddays(1)  
#> [1] "2016-03-13 14:00:00 EDT"  
one_pm + days(1)  
#> [1] "2016-03-13 13:00:00 EDT"
```


- What should the following code return?

```
years(1) / days(1)
```

- A duration with a starting point:

```
next_year <- today() + years(1)
(today() %--% next_year) / ddays(1)
#> [1] 365
```

	date			date time			duration			period			interval			number		
date	-						-	+		-	+					-	+	
date time				-			-	+		-	+					-	+	
duration	-	+		-	+		-	+	/							-	+	×
period	-	+		-	+					-	+					-	+	×
interval									/			/						
number	-	+		-	+		-	+	×	-	+	×	-	+	×	-	+	×

- Pick the simplest data structure that solves your problem:
 - ▶ If you only care about physical time, use a duration.
 - ▶ If you need to add human times, use a period.
 - ▶ If you need to figure out how long a span is in human units, use an interval.

```
Sys.timezone()
#> [1] "America/New_York"
length(OlsonNames())
#> [1] 608
head(OlsonNames())
#> [1] "Africa/Abidjan"      "Africa/Accra"        "Africa/Addis_Ababa"
#> [4] "Africa/Algiers"     "Africa/Asmara"       "Africa/Asmera"
```

■ Same instant, different place:

```
(x1 <- ymd_hms("2015-06-01 12:00:00", tz = "America/New_York"))  
#> [1] "2015-06-01 12:00:00 EDT"  
(x2 <- ymd_hms("2015-06-01 18:00:00", tz = "Europe/Copenhagen"))  
#> [1] "2015-06-01 18:00:00 CEST"  
(x3 <- ymd_hms("2015-06-02 04:00:00", tz = "Pacific/Auckland"))  
#> [1] "2015-06-02 04:00:00 NZST"  
x1 - x2  
#> Time difference of 0 secs  
x1 - x3  
#> Time difference of 0 secs
```

■ Note the behavior of 'c()':

```
x4 <- c(x1, x2, x3)  
x4  
#> [1] "2015-06-01 12:00:00 EDT" "2015-06-01 12:00:00 EDT"  
#> [3] "2015-06-01 12:00:00 EDT"
```

■ Keep the instant in time:

```
x4a <- with_tz(x4, tzone = "Australia/Lord_Howe")
x4a
#> [1] "2015-06-02 02:30:00 +1030" "2015-06-02 02:30:00 +1030"
#> [3] "2015-06-02 02:30:00 +1030"
x4a - x4
#> Time differences in secs
#> [1] 0 0 0
```

■ Change the instant in time:

```
x4b <- force_tz(x4, tzone = "Australia/Lord_Howe")
x4b
#> [1] "2015-06-01 12:00:00 +1030" "2015-06-01 12:00:00 +1030"
#> [3] "2015-06-01 12:00:00 +1030"
x4b - x4
#> Time differences in hours
#> [1] -14.5 -14.5 -14.5
```

1 Dates and times

2 Factors

3 Strings

- Factors are:
 - ▶ Used to work with categorical variables (i.e., that have a fixed and known set of possible values).
 - ▶ Useful to display character vectors in a non-alphabetical order.
- The **forcats** package:
 - ▶ Range of helpers for working with factors.

```
library(forcats)
```

- Imagine that you have a variable that records month:

```
x1 <- c("Dec", "Apr", "Jan", "Mar")
```

- Using a string to record this variable has two problems:
 - ▶ Twelve possible months and nothing saving you from typos.
 - ▶ It doesn't sort in a useful way.

```
sort(x1)  
#> [1] "Apr" "Dec" "Jan" "Mar"
```


- Start by creating a list of the valid **levels**:

```
month_levels <- c("Jan", "Feb", "Mar", "Apr", "May", "Jun",  
                  "Jul", "Aug", "Sep", "Oct", "Nov", "Dec")
```

- Then create a factor:

```
y1 <- factor(x1, levels = month_levels)  
y1  
#> [1] Dec Apr Jan Mar  
#> Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec  
sort(y1)  
#> [1] Jan Mar Apr Dec  
#> Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec  
factor(x1) ## without levels  
#> [1] Dec Apr Jan Mar  
#> Levels: Apr Dec Jan Mar
```

■ Notice:

```
x2 <- c("Dec", "Apr", "Jam", "Mar")
y2 <- factor(x2, levels = month_levels)
y2
#> [1] Dec Apr <NA> Mar
#> Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

■ Other ordering:

```
factor(x1, levels = unique(x1))
#> [1] Dec Apr Jan Mar
#> Levels: Dec Apr Jan Mar
factor(x1) %>%
  fct_inorder()
#> [1] Dec Apr Jan Mar
#> Levels: Dec Apr Jan Mar
```

■ Sample from the General Social Survey:


```
gss_cat
#> # A tibble: 21,483 x 9
#>   year marital    age race rincome partyid relig denom tuhours
#>   <int> <fct>    <int> <fct> <fct>    <fct> <fct> <fct>    <int>
#> 1  2000 Never m~    26 White $8000 t~ Ind,nea~ Prote~ South~    12
#> 2  2000 Divorced    48 White $8000 t~ Not str~ Prote~ Bapti~    NA
#> 3  2000 Widowed     67 White Not app~ Indepen~ Prote~ No de~     2
#> 4  2000 Never m~    39 White Not app~ Ind,nea~ Ortho~ Not a~     4
#> 5  2000 Divorced    25 White Not app~ Not str~ None   Not a~     1
#> 6  2000 Married     25 White $20000 ~ Strong ~ Prote~ South~    NA
#> 7  2000 Never m~    36 White $25000 ~ Not str~ Chris~ Not a~     3
#> 8  2000 Divorced    44 White $7000 t~ Ind,nea~ Prote~ Luthe~    NA
#> 9  2000 Married     44 White $25000 ~ Not str~ Prote~ Other     0
#> 10 2000 Married     47 White $25000 ~ Strong ~ Prote~ South~     3
#> # ... with 21,473 more rows
```

■ More info with ?gss_cat.

See levels of a factor from a tibble

■ A barplot:

```
ggplot(gss_cat, aes(race)) +  
  geom_bar()
```




■ Or a count:

```
gss_cat %>%  
  count(race)  
  
#> # A tibble: 3 x 2  
#>   race      n  
#>   <fct> <int>  
#> 1 Other  1959  
#> 2 Black  3129  
#> 3 White 16395
```


What's wrong here?

```
relig_summary <- gss_cat %>%  
  group_by(relig) %>%  
  summarize(age = mean(age, na.rm = TRUE),  
            tvhours = mean(tvhours, na.rm = TRUE),  
            n = n())  
  
ggplot(relig_summary, aes(tvhours, relig)) +  
  geom_point()
```




Modifying factor order

```
relig_summary %>%  
  mutate(relig = fct_reorder(relig, tvhours)) %>%  
  ggplot(aes(tvhours, relig)) +  
  geom_point()
```



Modify factor order II

```
gss_cat %>%  
  mutate(marital = marital %>% fct_infreq() %>% fct_rev()) %>%  
  ggplot(aes(marital)) +  
  geom_bar()
```



- More powerful than changing the orders of the levels is changing their values:
 - ▶ To clarify labels for publication.
 - ▶ To collapse levels for high-level displays.
- What's wrong here?

```
gss_cat %>%  
  count(partyid)  
#> # A tibble: 10 x 2  
#>   partyid      n  
#>   <fct>    <int>  
#> 1 No answer      154  
#> 2 Don't know       1  
#> 3 Other party     393  
#> 4 Strong republican 2314  
#> 5 Not str republican 3032  
#> 6 Ind,near rep     1791  
#> 7 Independent     4119  
#> 8 Ind,near dem     2499  
#> 9 Not str democrat 3690  
#> 10 Strong democrat 3490
```


Modifying factor levels II

```
gss_cat %>%
  mutate(partyid = fct_recode(partyid,
    "Republican, strong" = "Strong republican",
    "Republican, weak" = "Not str republican",
    "Independent, near rep" = "Ind,near rep",
    "Independent, near dem" = "Ind,near dem",
    "Democrat, weak" = "Not str democrat",
    "Democrat, strong" = "Strong democrat")) %>%
  count(partyid)

#> # A tibble: 10 x 2
#>   partyid      n
#>   <fct>    <int>
#> 1 No answer    154
#> 2 Don't know     1
#> 3 Other party   393
#> 4 Republican, strong 2314
#> 5 Republican, weak  3032
#> 6 Independent, near rep 1791
#> 7 Independent    4119
#> 8 Independent, near dem 2499
#> 9 Democrat, weak  3690
#> 10 Democrat, strong 3490
```

Collapsing factors

```
gss_cat %>%
  mutate(partyid = fct_recode(partyid,
    "Republican, strong"      = "Strong republican",
    "Republican, weak"       = "Not str republican",
    "Independent, near rep"  = "Ind,near rep",
    "Independent, near dem"  = "Ind,near dem",
    "Democrat, weak"        = "Not str democrat",
    "Democrat, strong"      = "Strong democrat",
    "Other"                 = "No answer",
    "Other"                 = "Don't know",
    "Other"                 = "Other party" )) %>%
  count(partyid)

#> # A tibble: 8 x 2
#>   partyid          n
#>   <fct>        <int>
#> 1 Other          548
#> 2 Republican, strong 2314
#> 3 Republican, weak  3032
#> 4 Independent, near rep 1791
#> 5 Independent      4119
#> 6 Independent, near dem 2499
#> 7 Democrat, weak    3690
#> 8 Democrat, strong  3490
```

```
gss_cat %>%
  mutate(partyid = fct_collapse(partyid,
    other = c("No answer", "Don't know", "Other party"),
    rep = c("Strong republican", "Not str republican"),
    ind = c("Ind,near rep", "Independent", "Ind,near dem"),
    dem = c("Not str democrat", "Strong democrat")
  )) %>%
  count(partyid)
#> # A tibble: 4 x 2
#>   partyid     n
#>   <fct>   <int>
#> 1 other     548
#> 2 rep     5346
#> 3 ind     8409
#> 4 dem     7180
```

Collapsing factor III

```
gss_cat %>%  
  mutate(relig = fct_lump(relig)) %>%  
  count(relig)  
  
#> # A tibble: 2 x 2  
#>   relig      n  
#>   <fct>    <int>  
#> 1 Protestant 10846  
#> 2 Other      10637  
  
gss_cat %>%  
  mutate(relig = fct_lump(relig, n = 3)) %>%  
  count(relig, sort = TRUE)  
  
#> # A tibble: 4 x 2  
#>   relig      n  
#>   <fct>    <int>  
#> 1 Protestant 10846  
#> 2 Catholic    5124  
#> 3 None        3523  
#> 4 Other        1990
```

1 Dates and times

2 Factors

3 Strings

```
library(stringr) # package for string manipulation

# To create strings
string1 <- "This is a string"
string2 <- 'To get a "quote" inside a string, use single quotes'
```

■ Backslash as escape character:

```
double_quote <- "\" # or '"'
single_quote <- '\'' # or '"'
```

■ The printed representation is not the string itself:

```
x <- c("\"", "\\")
x
#> [1] "\" "\\\"
writeLines(x)
#> "
#> \
```

- Special characters:
 - ▶ Use `"\n"`, for newline, or `"\t"`, for tab.
 - ▶ Complete list by requesting help on `"(?''', or ?''")`
- Other useful things:

```
(x <- "\u00b5") # Non-English characters
#> [1] "µ"
c("one", "two", "three") # Character vectors
#> [1] "one" "two" "three"
str_length(c("a", "R for data science", NA)) # String length
#> [1] 1 18 NA
```

- `stringr` autocomplete:

<pre>> str_c > str_conv > str_count > str_detect > str_dup > str_extract > str_extract_all > str_</pre>	<pre>{stringr} {stringr} {stringr} {stringr} {stringr} {stringr} {stringr}</pre>	<p><code>str_c(..., sep = "", collapse = NULL)</code></p> <p>To understand how <code>str_c</code> works, you need to imagine that you are building up a matrix of strings. Each input argument forms a column, and is expanded to the length of the longest argument, using the usual recycling rules. The <code>sep</code> string is inserted between each column. If <code>collapse</code> is <code>NULL</code> each row is collapsed into a single string. If non-<code>NULL</code> that string is inserted at the end of each row, and the entire matrix collapsed to a single string.</p> <p>Press F1 for additional help</p>
---	--	--

■ Combining strings:

```
str_c("x", "y")  
#> [1] "xy"  
str_c("x", "y", "z")  
#> [1] "xyz"  
str_c("x", "y", sep = ", ")  
#> [1] "x, y"
```

■ Missing values:

```
x <- c("abc", NA)  
str_c("|-", x, "-|")  
#> [1] "/-abc-/" NA  
str_c("|-", str_replace_na(x), "-|")  
#> [1] "/-abc-/" "/-NA-/"
```

■ Recycling:

```
str_c("prefix-", c("a", "b", "c"), "-suffix")  
#> [1] "prefix-a-suffix" "prefix-b-suffix" "prefix-c-suffix"
```

■ Collapsing a vector of strings:

```
str_c(c("x", "y", "z"), collapse = ", ")  
#> [1] "x, y, z"
```



```
x <- c("Apple", "Banana", "Pear")
str_sub(x, 1, 3)
#> [1] "App" "Ban" "Pea"
str_sub(x, -3, -1)
#> [1] "ple" "ana" "ear"
str_sub("a", 1, 5)
#> [1] "a"
str_sub(x, 1, 1) <- str_to_lower(str_sub(x, 1, 1))
x
#> [1] "apple" "banana" "pear"
```

- See also `str_to_upper()` or `str_to_title()`.

```
# Turkish has two i's: with and without a dot, and it  
# has a different rule for capitalising them:  
str_to_upper(c("i", "ı"))  
#> [1] "I" "I"  
str_to_upper(c("i", "ı"), locale = "tr")  
#> [1] "İ" "I"
```

■ The locale:

- ▶ An ISO 639 language code, which is a **two or three letter abbreviation**
- ▶ If blank, R uses the current locale, as provided by your operating system.

Some people, when confronted with a problem, think “I know, I’ll use regular expressions.” Now, they have two problems. — Jamie Zawinski

- A language that allows you to describe patterns in strings.
- Allows you for instance to:
 - ▶ Determine which strings match a pattern.
 - ▶ Find the positions of matches.
 - ▶ Extract the content of matches.
 - ▶ Replace matches with new values.
 - ▶ Split a string based on a match.
- Read the chapter on strings from the book!

- The simplest patterns match exact strings:

```
x <- c("apple", "banana", "pear")  
str_view(x, "an")
```

apple

ban**a**na

pear

- Next step is ., which matches any character (except a newline):

```
str_view(x, ".a.")
```

apple

ban**a**na

pear**a**

- If "." matches any character, how to match the character "."?

- If “.” matches any character, how to match the character “.”?
 - ▶ Need to use an “escape” (like string, a backslash \).
 - ▶ So to match an ., need the regexp \.
 - ▶ But \ is also an escape symbol in strings.
 - ▶ So to create the regexp \., use the string "\\.”.

```
# To create the regexp, we need \\  
dot <- "\\.”  
# But the expression itself only contains one:  
writeLines(dot)  
#> \.  
# And this tells R to look for an explicit .  
str_view(c("abc", "a.c", "bef"), "a\\.c")
```

abc

a.c

bef

- If \ is an escape character, how do you match a literal \?
 - ▶ Need to escape it, i.e. create the regexp \\.
 - ▶ To create that regexp with a string, which also needs to escape \, need to write "\\\\"
 - ▶ I.e., need four backslashes to match one!

```
x <- "a\\b"  
writeLines(x)  
#> a\b  
str_view(x, "\\")
```

a\b

- By default, regexps match any part of a string.
- Often useful to *anchor* the regexp:
 - ▶ `^` to match the start of the string.
 - ▶ `$` to match the end of the string.

```
x <- c("apple", "banana", "pear")
```

```
str_view(x, "^a")
```

a
apple
banana
pear

```
str_view(x, "a$")
```

apple
bananaa
pear

- To remember, [Evan Misshula's mnemonic](#): if you begin with power (`^`), you end up with money (`$`).

- To force a regexp to only match a complete string, anchor it with both `^` and `$`:

```
x <- c("apple pie", "apple", "apple cake")
```

```
str_view(x, "apple")
```

```
apple pie  
apple  
apple cake
```

```
str_view(x, "^apple$")
```

```
apple pie  
apple  
apple cake
```


- Some special patterns match more than one character:
 - ▶ Already seen `.` (matches any character apart from a newline).
 - ▶ Two other useful tools:
 - `\d`: matches any digit.
 - `\s`: matches any whitespace (e.g. space, tab, newline).
 - ▶ To create a regexp containing `\d` or `\s`:
 - Need to escape the `\` for the string.
 - So type `"\\d"` or `"\\s"`.
- The other two tools are:
 - ▶ **Character classes**
 - `[abc]`: matches a, b, or c.
 - `[^abc]`: matches anything except a, b, or c.
 - ▶ **Alternatives**
 - `abc|d..f`: matches either "abc", or "deaf".

- Can be used as an alternative to backslash escapes.

```
str_view(c("abc",  
          "a.c",  
          "a*c",  
          "a c"),  
        "a[.]c")
```

abc

a.c

a*c

a c

```
str_view(c("abc",  
          "a.c",  
          "a*c",  
          "a c"),  
        "a.[*]c")
```

abc

a.c

a*c

a c

```
str_view(c("abc",  
          "a.c",  
          "a*c",  
          "a c"),  
        "a[ ]")
```

abc

a.c

a*c

a c

- Used to pick between one or more alternative patterns.
- Works for most regex metacharacters: \$. | ? * + () [{.
- But some have special meaning even inside a character class.
 - ▶ Must be handled with backslash escapes:] \ ^ and -.

- Note that the precedence for `|` is low:
 - ▶ `abc|xyz`: matches `abc` or `xyz`, not `abxyz` or `abxyz`.
- Same as mathematical expressions: if it gets confusing, use parentheses.

```
str_view(c("grey", "gray"), "gr(e|a)y")
```

grey

gray

- To control how many times a pattern matches:

- ▶?: 0 or 1.
- ▶+: 1 or more.
- ▶*: 0 or more.

```
# 1888 is the longest year in Roman numerals
```

```
x <- "MDCCCLXXXVIII"
```

```
str_view(x, "CC?")
```

MDCCCLXXXVIII

```
str_view(x, "CC+")
```

MDCCCLXXXVIII

```
str_view(x, 'C[LCX]+')
```

MDCCCLXXXVIII

- The precedence of these operators is high:
 - ▶ colou?r: matches either US or British spellings.
 - ▶ Most uses will need parentheses, like bana(na)+.

■ To specify the number of matches precisely:

- ▶ `{n}`: exactly n.
- ▶ `{n,}`: n or more.
- ▶ `{,m}`: at most m.
- ▶ `{n,m}`: between n and m.

```
str_view(x, "C{2}")
```

MDCCCLXXXVIII

```
str_view(x, "C{2,}")
```

MDCCCLXXXVIII

```
str_view(x, "C{2,3}")
```

MDCCCLXXXVIII

- Earlier: parentheses as a way to disambiguate complex expressions.
- But parentheses also create a *numbered* capturing group.
- A capturing group stores *the part of the string* matched by the part of the regexp inside the parentheses.
- Refer to the same text as previously matched by a capturing group with *backreferences*, like `\1`, `\2` etc.

```
str_view(fruit, "(..)\1", match = TRUE)
```

```
banana  
coconut  
cucumber  
jujube  
papaya  
salal berry
```

- Cool applications in chapter 14.4!